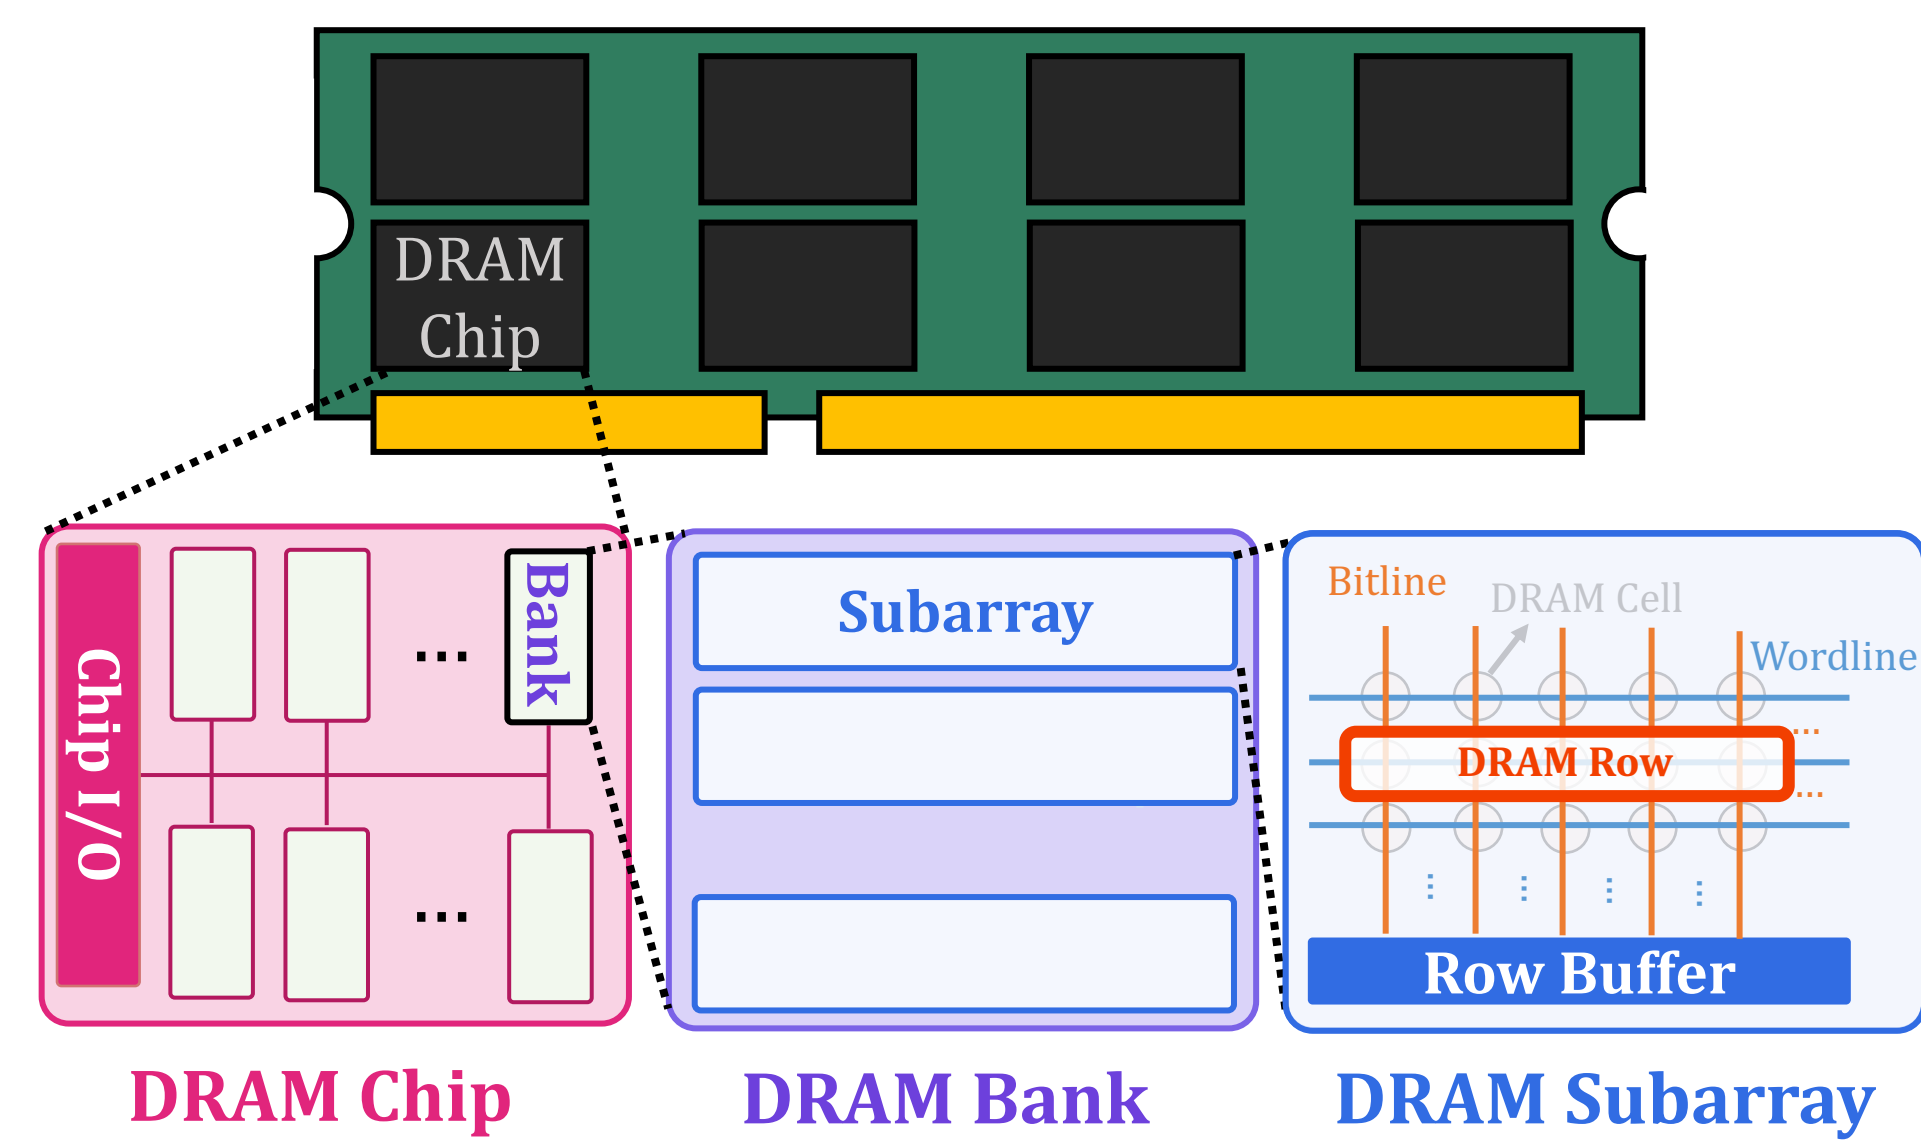
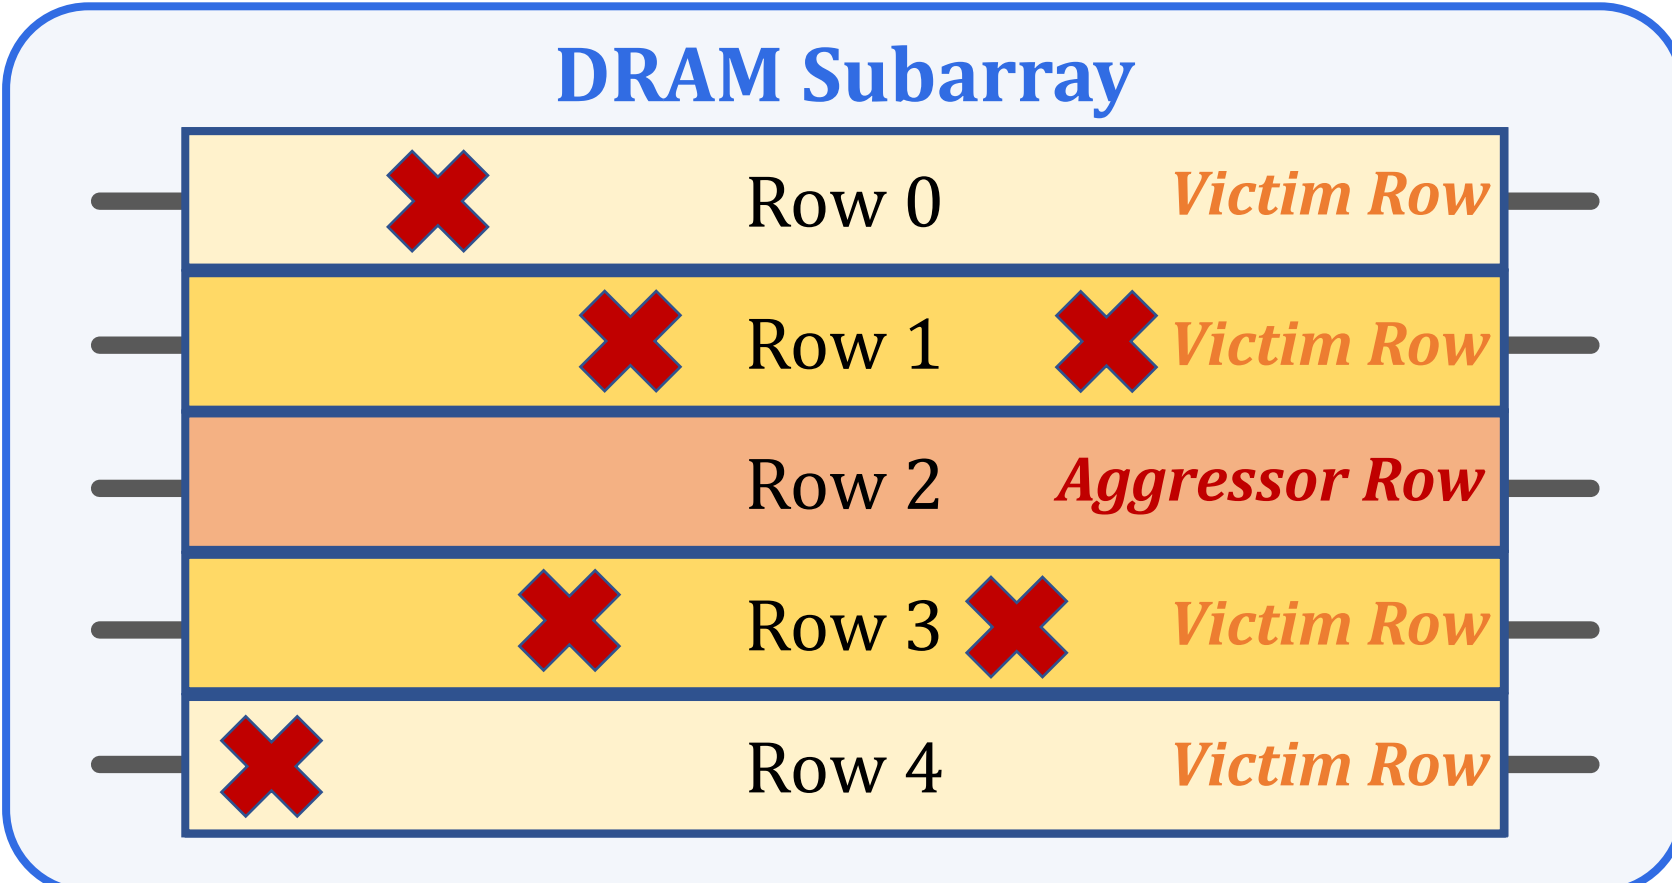


Nisa Bostanci Ataberk Olgun Giray Yaglikci Geraldo de Oliveira Yahya Tugrul
Haocong Luo Ismail Yuksel Konstantinos Kanellopoulos Mohammad Sadrosadati Onur Mutlu

1: DRAM Organization




2: Read Disturbance Vulnerabilities



- Repeatedly **activating** and **precharging** a DRAM row causes **RowHammer bitflips** in nearby cells
- The minimum number of activations that causes a bitflip is called the **RowHammer threshold**

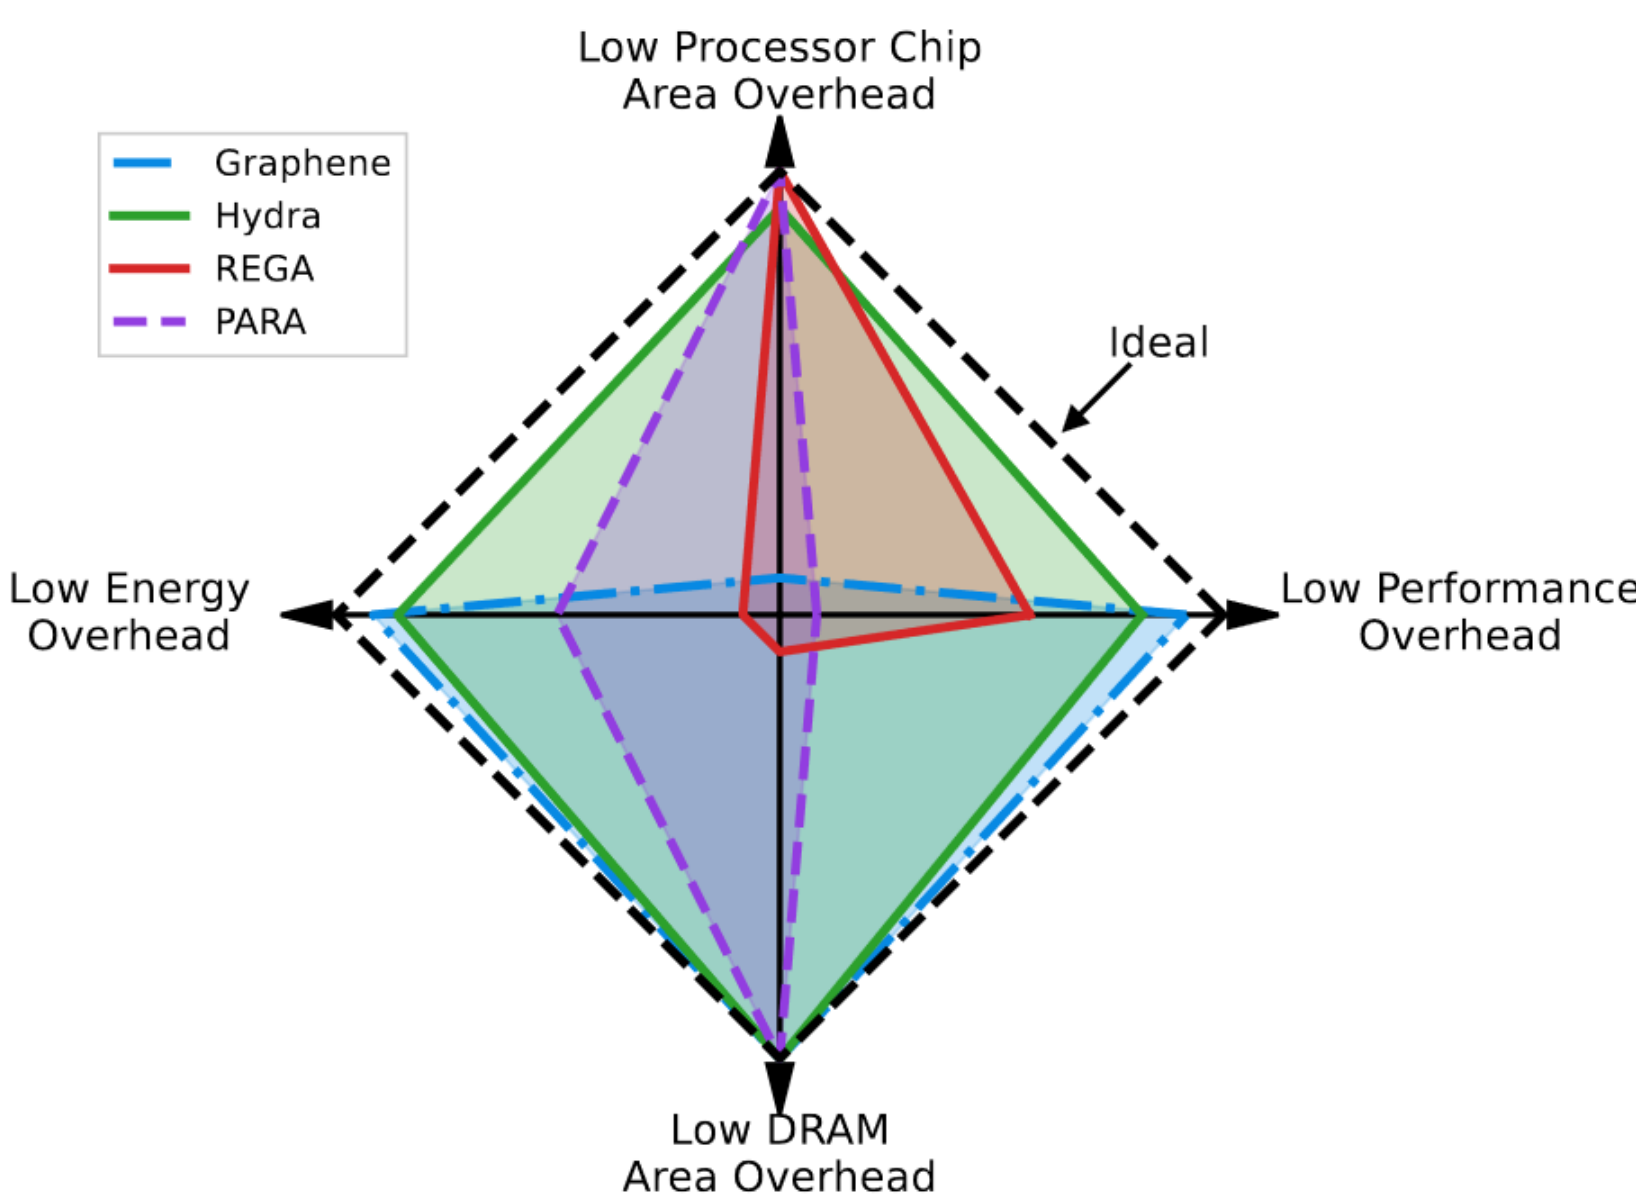
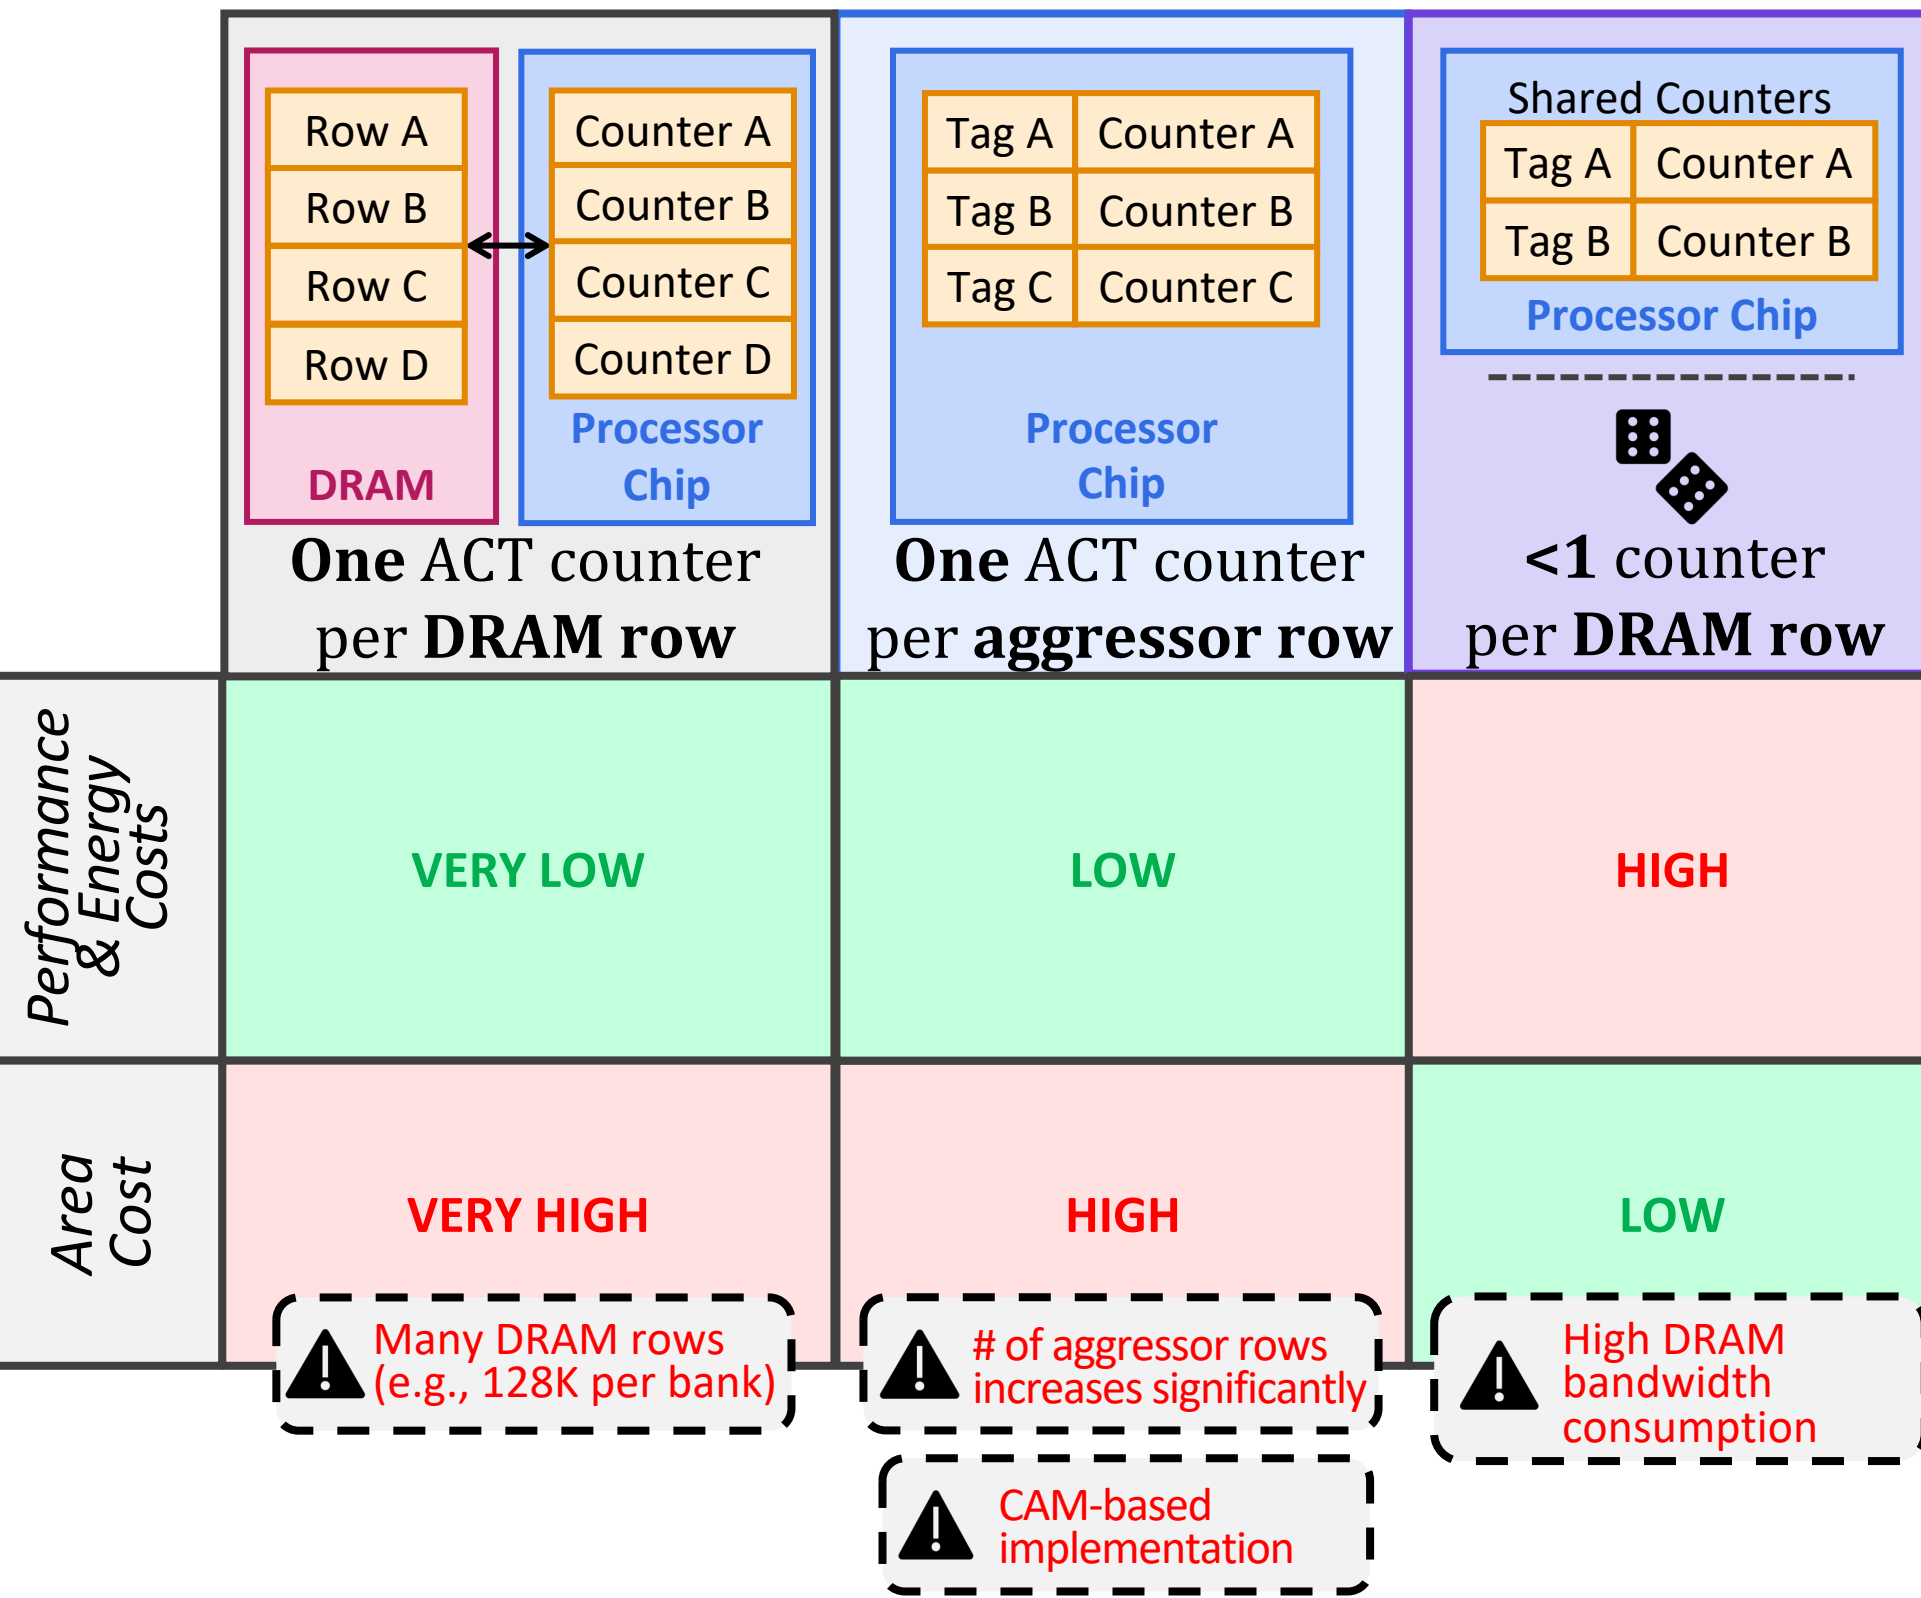
DRAM chips are more vulnerable to read disturbance today

Read disturbance bitflips occur at much lower activation counts
(more than two orders of magnitude decrease in less than a decade):



Mitigation techniques against read disturbance attacks need to be **effective** and **efficient** for highly vulnerable systems

3: Limitations of Existing Mitigations



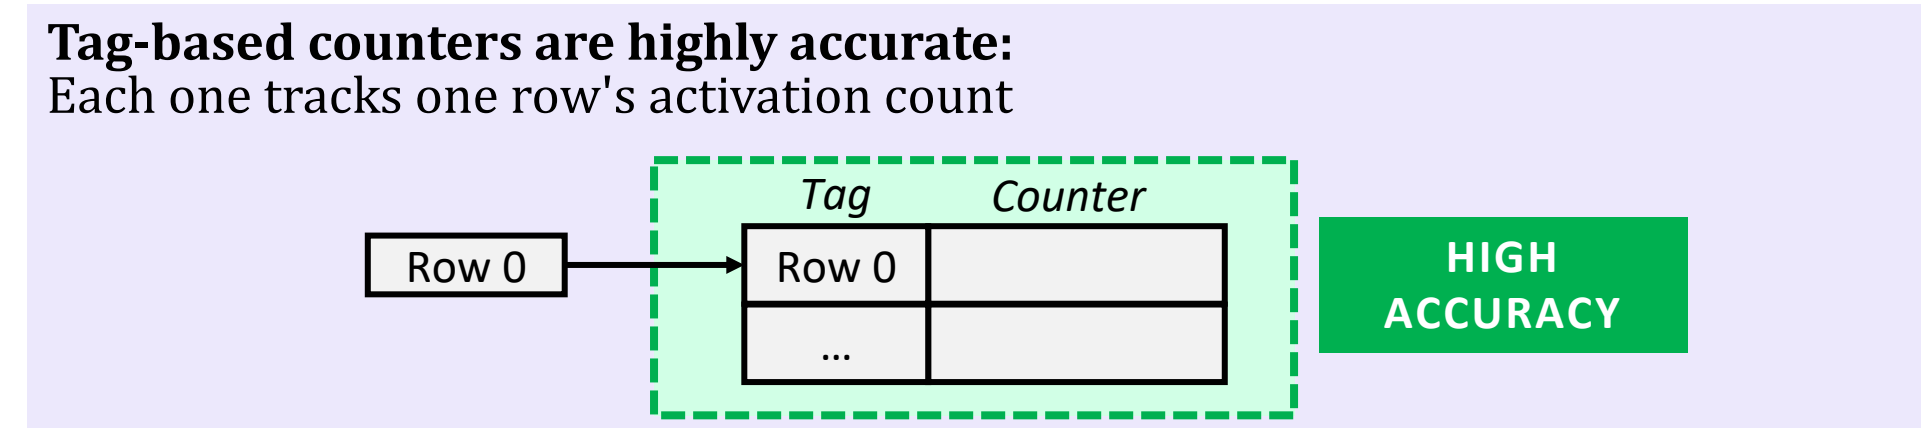
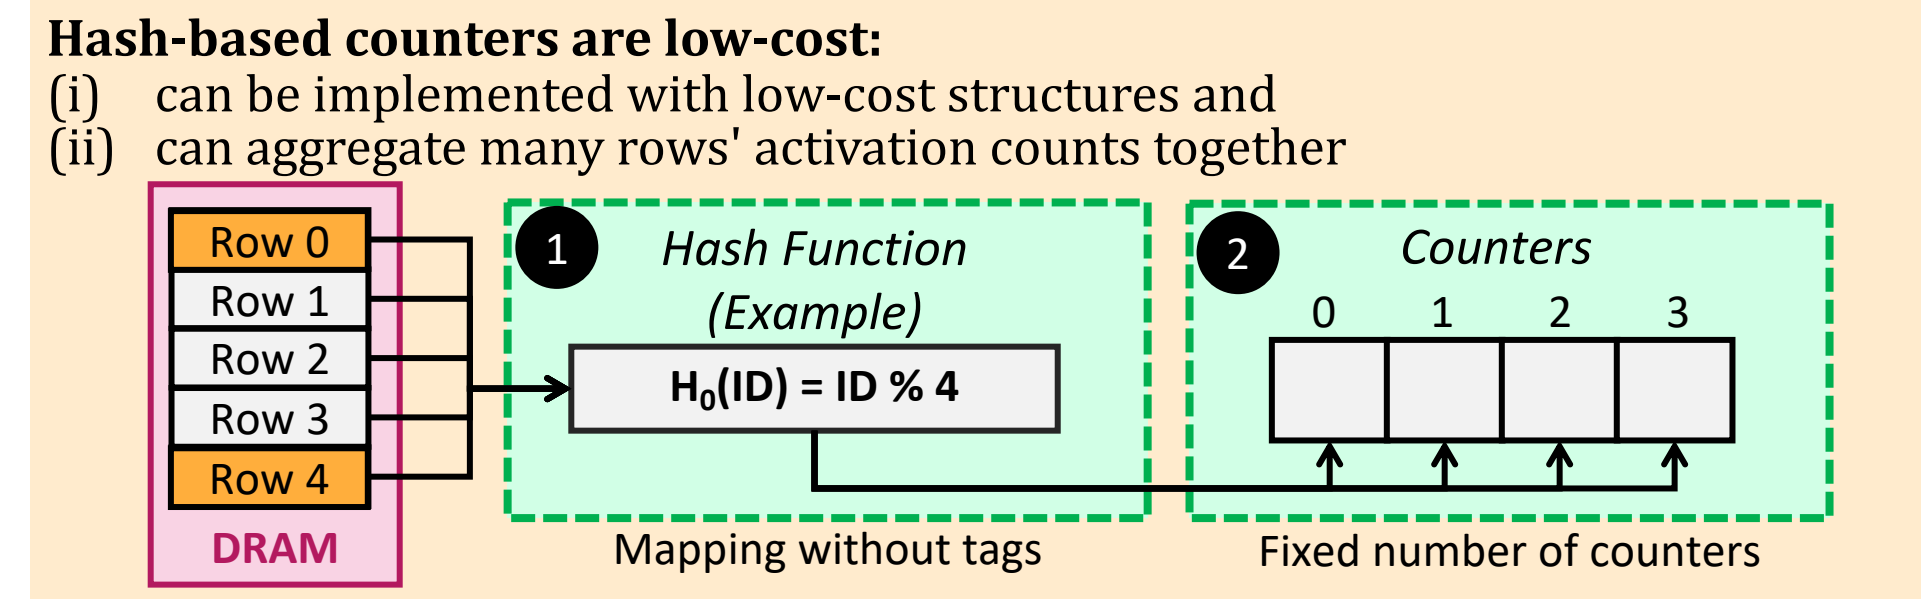
No existing mitigation technique prevents RowHammer bitflips at low area, performance and energy costs

4: Goal

Prevent RowHammer bitflips with low area, performance, and energy overheads in highly RowHammer-vulnerable DRAM-based systems (e.g., a RowHammer threshold of 125)

5: CoMeT's Key Observation and Key Idea

Key Observation



Counter Table (CT):
- Maps each DRAM row to a group of low-cost hash-based counters by employing the Count-Min Sketch technique
- Triggers a preventive refreshes when the aggressor's counter group reaches an activation threshold
Tracks DRAM row activations at low area cost

Recent Aggressor Table (RAT):
Allocates highly accurate per-DRAM-row counters for *only* a small set of DRAM rows that are activated many times
Reduces performance penalties by increasing tracking accuracy

Key Idea


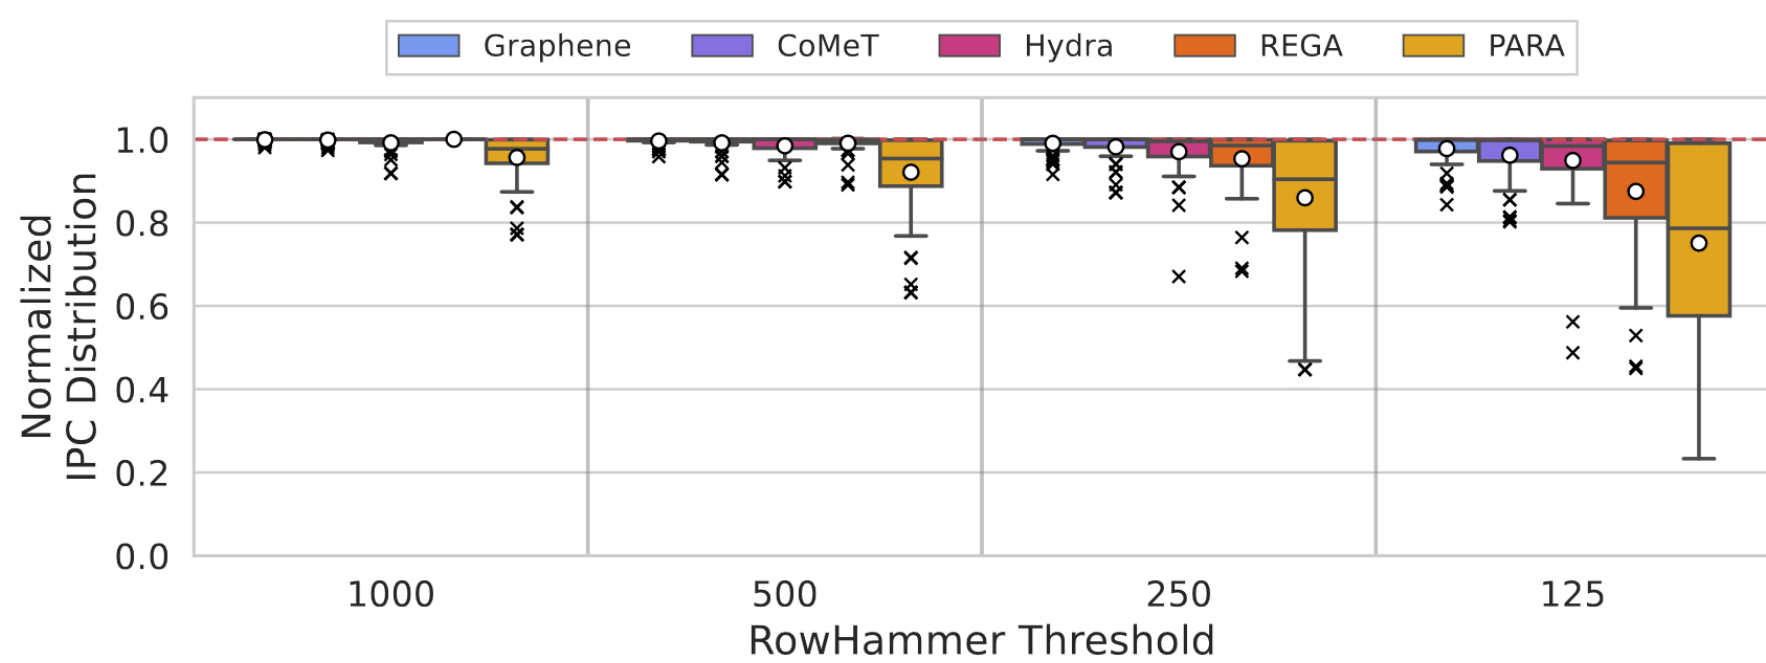
1 Track most DRAM rows' activations with low area cost

2 Track a small set of hot rows accurately

6: CoMeT's Evaluation and Conclusion

| | |
|------------------|--|
| Processor | 1 or 8 cores, 3.6GHz clock frequency, 4-wide issue, 128-entry instruction window |
| DRAM | DDR4, 1 channel, 2 rank/channel, 4 bank groups, 4 banks/bank group, 128K rows/bank |
| Memory Ctrl. | 64-entry read and write requests queues, Scheduling policy: FR-FCFS [137, 138] with a column cap of 16 [139] |
| Last-Level Cache | 8 MiB (single-core), 16 MiB (8-core) |

Cycle-level simulations using Ramulator Workloads:
62 1- & 8-core workloads
Four different very low nRH values: 1000, 500, 250, 125
Four state-of-the-art mitigation mechanisms: Graphene, Hydra, PARA, REGA



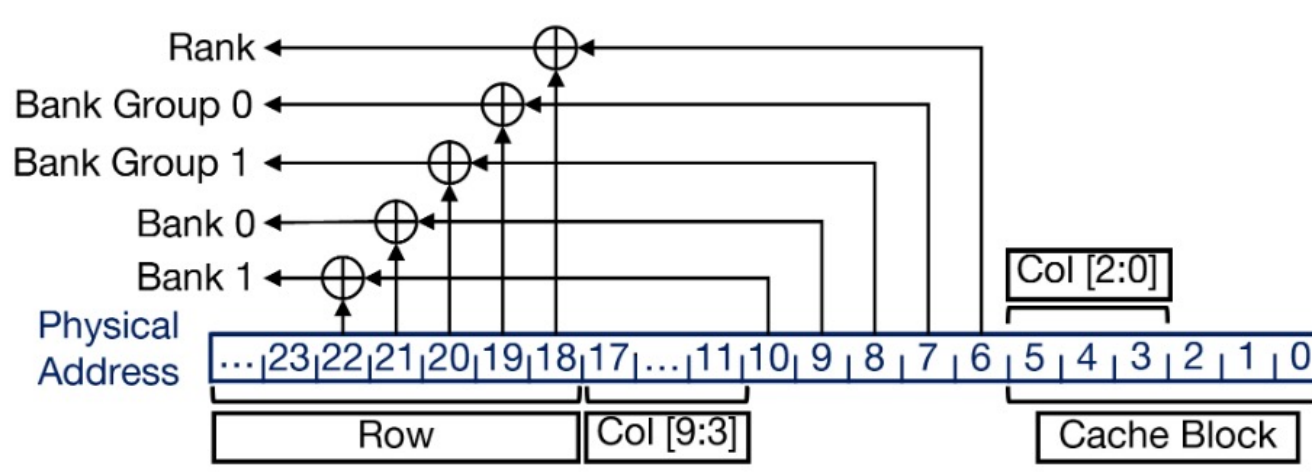
CoMeT achieves a good trade-off between area, performance and energy costs
- incurs significantly less area overhead (74.2x)
- outperforms the state-of-the-art (by up to 39.1%)

Open-sourced: <https://github.com/CMU-SAFARI/CoMeT>




7: ABACuS's Key Observation and Key Idea

- Address mappings distribute **consecutive** cache blocks to **different** banks (but to the **same row ID**)
- Workloads tend to access cache blocks in close proximity at around the same time

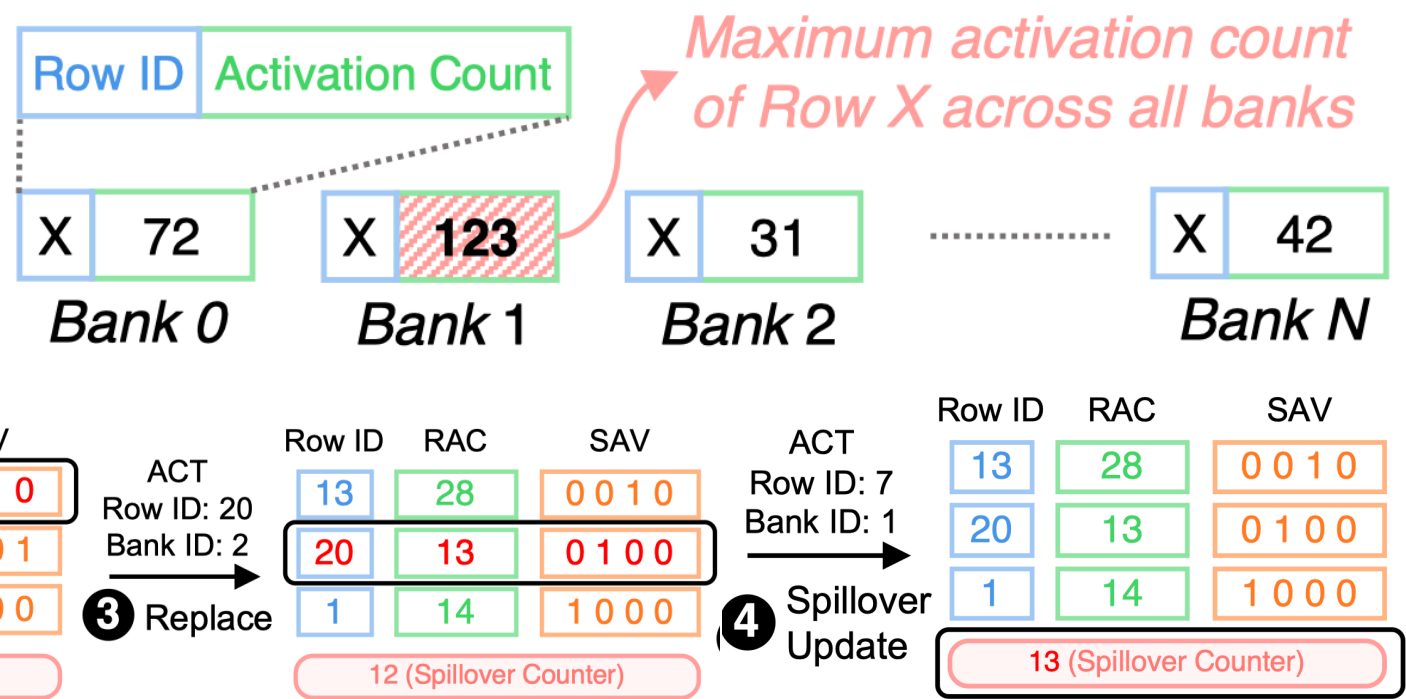


Many workloads access the same row ID in different banks at around the same time
Sibling rows: Rows with the same ID across all banks



Key Idea: Sibling rows can share one hardware counter
Reduce the number of counters by a factor of the number of banks in the chip

Key Mechanism: ABACuS
Track the **maximum** (worst) **activation count** of sibling rows using **one counter**

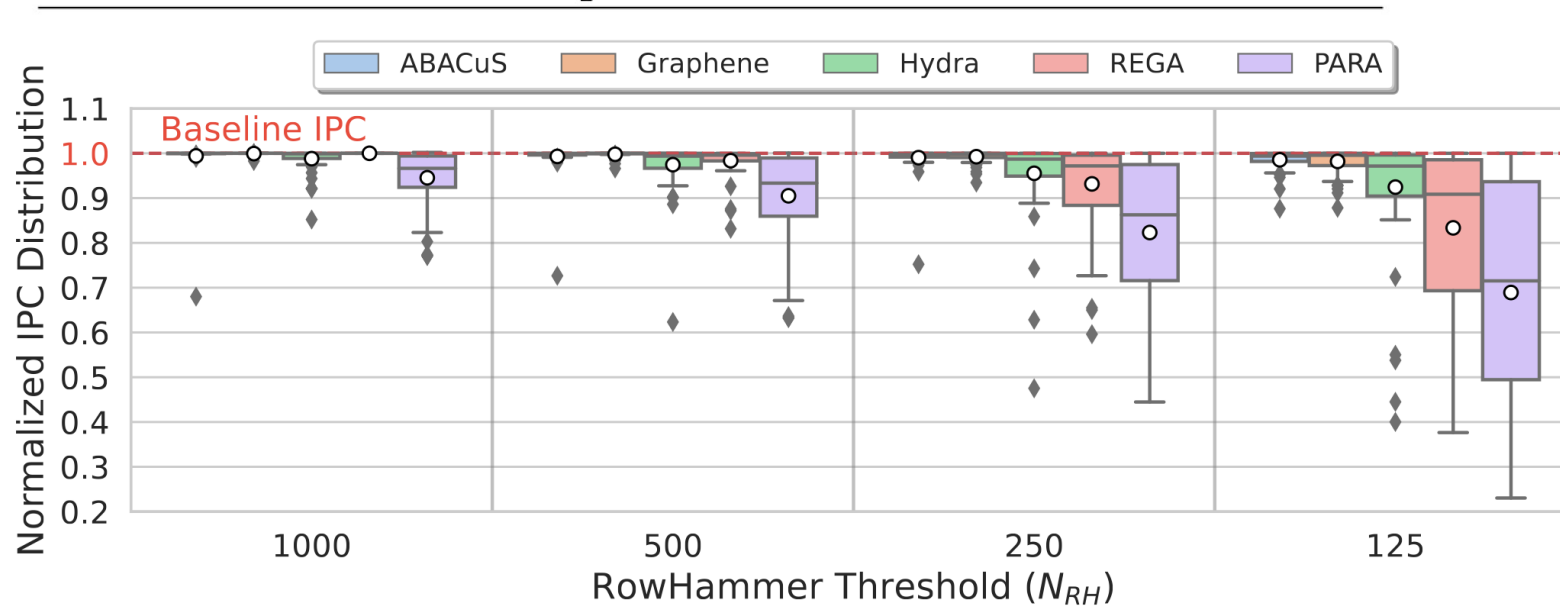


| Row ID | *RAC | *SAV | ACT | Row ID | RAC | SAV | ACT | Row ID | RAC | SAV | ACT | Row ID | RAC | SAV | |
|------------------------|------|---------|--------------------------|------------------------|-----|---------|--------------------------|------------------------|-----|---------|--------------------------|------------------------|-----|---------|--------------------|
| 13 | 27 | 0 0 0 1 | Row ID: 13 Bank ID: 1 | 13 | 27 | 0 0 1 1 | Row ID: 13 Bank ID: 1 | 13 | 28 | 0 0 1 0 | Row ID: 13 Bank ID: 2 | 13 | 28 | 0 0 1 0 | |
| 9 | 12 | 0 1 0 1 | | 9 | 12 | 0 1 0 1 | | 9 | 12 | 0 1 0 1 | | 9 | 12 | 0 1 0 1 | |
| 1 | 14 | 1 0 0 0 | | 1 | 14 | 1 0 0 0 | | 1 | 14 | 1 0 0 0 | | 1 | 14 | 1 0 0 0 | |
| 12 (Spillover Counter) | | | 1 Update | 12 (Spillover Counter) | | | 2 Update | 12 (Spillover Counter) | | | 3 Replace | 12 (Spillover Counter) | | | 4 Spillover Update |
| | | | | | | | | | | | | 13 (Spillover Counter) | | | |

8: ABACuS's Evaluation and Conclusion

| | |
|------------------|--|
| Processor | 1 or 8 cores, 3.6GHz clock frequency, 4-wide issue, 128-entry instruction window |
| DRAM | DDR4, 1 channel, 2 rank/channel, 4 bank groups, 4 banks/bank group, 128K rows/bank, 3200 MT/s |
| Memory Ctrl. | 64-entry read and write requests queues, Scheduling policy: FR-FCFS [181, 182] with a column cap of 16 [183], Address mapping: MOP [166, 168] 45 ns tRC, 7.9 µs tREFI, 64 ms tREFW 64 ms ABACuS reset period |
| Last-Level Cache | 2 MiB per core |

Cycle-level simulations using Ramulator Workloads:
62 1- & 8-core workloads
Four different very low nRH values: 1000, 500, 250, 125
Four state-of-the-art mitigation mechanisms: Graphene, Hydra, PARA, REGA



Lower overhead than all evaluated state-of-the-art mechanisms

Open-sourced: <https://github.com/CMU-SAFARI/ABACuS>

| Mitigation Mechanism | SRAM KB | CAM KB | mm² | % CPU | % DRAM | Access Energy (pJ) | Static Power (mW) |
|------------------------------|---------|--------|------|-------|--------|--------------------|-------------------|
| ABACuS | 10.63 | 8.30 | 0.04 | 0.02 | - | 25.98 | 12.22 |
| Row ID Table | - | 5.64 | 0.01 | <0.01 | - | 12.85 | 6.61 |
| Row Activation Counter Table | - | 2.66 | 0.02 | <0.01 | - | 11.13 | 4.66 |
| Sibling Activation Vector | 10.63 | - | 0.01 | <0.01 | - | 2.00 | 0.95 |
| Graphene [102] | - | 286.51 | 0.81 | 0.35 | - | 873.38 | 187.98 |
| Hydra [106] | 61.56 | - | 0.10 | 0.04 | - | 43.07 | 24.17 |
| REGA [177] | - | - | - | - | 2.06 | - | - |

