Enabling Efficient and Scalable DRAM Read Disturbance Mitigation via New Experimental Insights into Modern DRAM Chips

Abdullah Giray Yağlıkçı

Aug 14, 2024

Rambus

SAFARI

ETH zürich

Brief Self Introduction

- Abdullah Giray Yaglikci
  - Researcher @ SAFARI Research Group since August 2016
    - ETH Zurich (Feb 2018 – ongoing)
    - Intel Labs (Aug 2017 – Feb 2018)
    - Carnegie Mellon University (Aug 2016 – Aug 2017)
  - Defended my PhD thesis, advised by Onur Mutlu, in April 2024
  - https://agyaglikci.github.io/
  - agirayyaglikci@gmail.com (Best way to reach me)
  - https://safari.ethz.ch

- Research interests:
  - Computer architecture, hardware security
  - Memory and storage systems
  - Hardware security, safety, reliability, performance, availability, fairness, energy efficiency
  - Hardware/software cooperation
  - ...

Papers in This Talk

- Lois Orosa, Abdullah Giray Yaglikci, Haocong Luo, Ataberk Olgun, Jisung Park, Hasan Hassan, Minesh Patel, Jeremie S. Kim, and Onur Mutlu, **"A Deeper Look into RowHammer's Sensitivities: Experimental Analysis of Real DRAM Chips and Implications on Future Attacks and Defenses"** *in MICRO, 2021*. [Slides (pptx) (pdf)] [Short Talk Slides (pptx) (pdf)] [Lightning Talk Slides (pptx) (pdf)] [Talk Video (21 mins)] [Lightning Talk]

- A. Giray Yağlıkçı, Haocong Luo, Geraldo F. de Oliviera, Ataberk Olgun, Minesh Patel, Jisung Park, Hasan Hassan, Jeremie S. Kim, Lois Orosa, and Onur Mutlu, **"Understanding RowHammer Under Reduced Wordline Voltage: An Experimental Study Using Real DRAM Devices"** *in DSN, 2022.* [Slides (pptx) (pdf)] [Lightning Talk Slides (pptx) (pdf)] [arXiv version] [Talk Video (34 mins)] [Lightning Talk Video (2 mins)]

- Abdullah Giray Yaglikci, Geraldo Francisco de Oliveira, Yahya Can Tugrul, Ismail Yuksel, Ataberk Olgun, Haocong Luo, and Onur Mutlu, **"Spatial Variation-Aware Read Disturbance Defenses: Experimental Analysis of Real DRAM Chips and Implications on Future Solutions"** *in HPCA, 2024*. [Slides (pptx) (pdf)] [arXiv version]

- A. Giray Yaglıkcı, Ataberk Olgun, Minesh Patel, Haocong Luo, Hasan Hassan, Lois Orosa, Oguz Ergin, and Onur Mutlu, **"HiRA: Hidden Row Activation for Reducing Refresh Latency of Off-the-Shelf DRAM Chips"** *in MICRO, 2022*. [Slides (pptx) (pdf)][Longer Lecture Slides (pptx) (pdf)] [Lecture Video (36 minutes)]

- A. Giray Yaglikci, Minesh Patel, Jeremie S. Kim, Roknoddin Azizi, Ataberk Olgun, Lois Orosa, Hasan Hassan, Jisung Park, Konstantinos Kanellopoulos, Taha Shahroodi, Saugata Ghose, and Onur Mutlu, **"BlockHammer: Preventing RowHammer at Low Cost by Blacklisting Rapidly-Accessed DRAM Rows,"** in *HPCA*, 2021. [Slides (pptx) (pdf)] [Short Talk Slides (pptx) (pdf)] [Talk Video (22 minutes)] [Short Talk Video (7 minutes)] [Intel Hardware Security Academic Awards Short Talk Slides (pptx) (pdf)] [Intel Hardware Security Academic Awards Short Talk Video (2 minutes)] [BlockHammer Source Code] ***Intel Hardware Security Academic Award Finalist (one of 4 finalists out of 34 nominations)***

- Oğuzhan Canpolat, A. Giray Yağlıkçı, Geraldo F. Oliveira, Ataberk Olgun, Oğuz Ergin, and Onur Mutlu, **"Understanding the Security Benefits and Overheads of Emerging Industry Solutions to DRAM Read Disturbance"** ***DRAMSec****, held with* ***ISCA***, 2024. [Slides (pptx) (pdf)] [arXiv version] [Source Code]

- Oğuzhan Canpolat and A. Giray Yağlıkçı and Ataberk Olgun and İsmail Emir Yüksel and Yahya Can Tuğrul and Konstantinos Kanellopoulos and Oğuz Ergin and Onur Mutlu **"Leveraging Adversarial Detection to Enable Scalable and Low Overhead RowHammer Mitigations,"** arXiv [cs.CR 2404.13477], 2024.

Memory Isolation

- *A memory access should* ***not have unintended side effects*** *on data stored in* ***other addresses***
- A fundamental property for **robustness** (safety, security, and reliability)

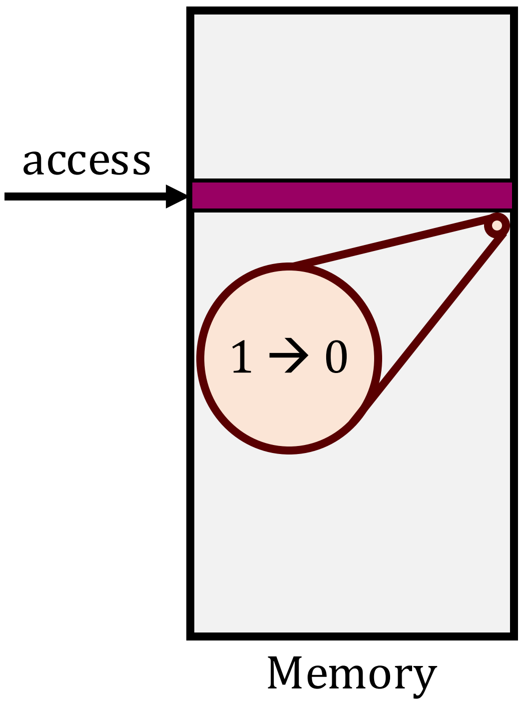

Data loss or **corruption**

Compromise **application correctness**

Leak private information

Take over a computer

Memory isolation is **difficult in modern memory chips**

Read Disturbance in Modern Memory Chips

- Prevalent memory technology:
  Dynamic Random Access Memory (DRAM)

- DRAM stores **data** in the form of **electrical charge** on a capacitor

- DRAM **leaks charge** over time and needs **periodic refresh**

- **DRAM Read Disturbance**:
  Accessing a DRAM cell disturbs other **physically nearby cells** and exacerbates their **charge leakage**

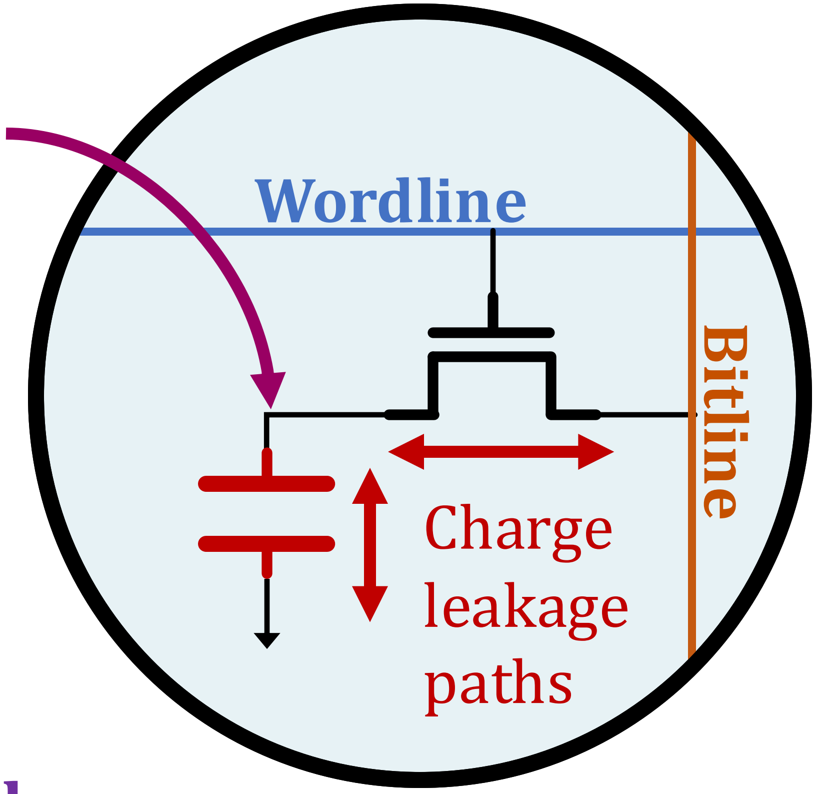

RowHammer: An Example of DRAM Read Disturbance

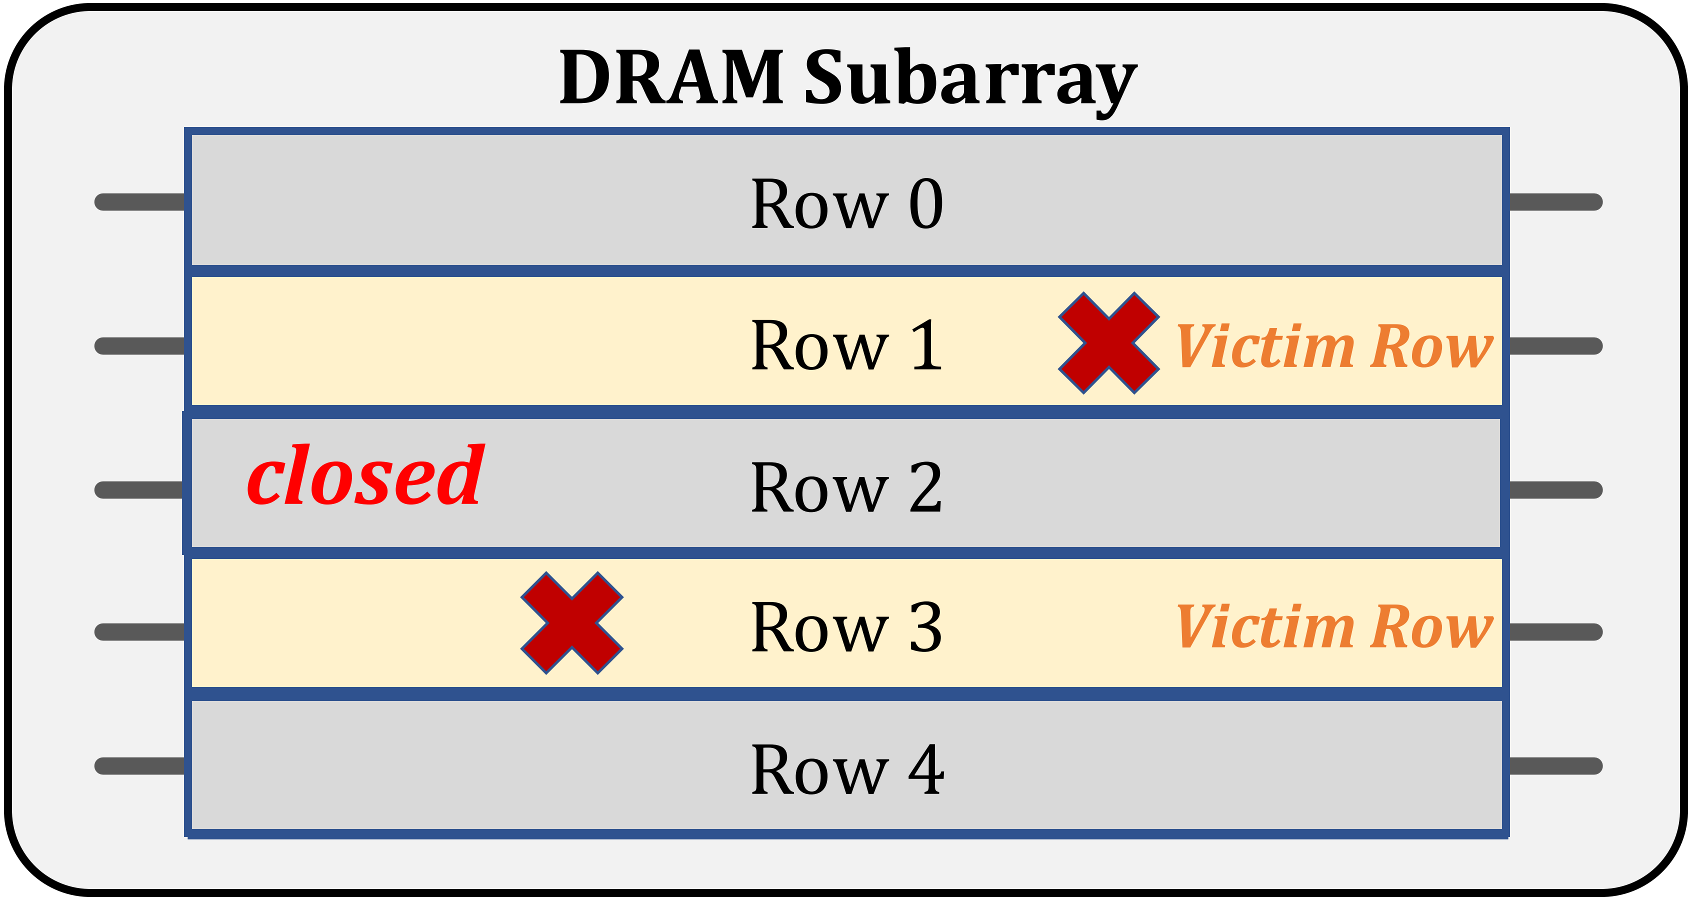

Repeatedly **opening** (activating) and **closing** (precharging) a DRAM row causes **RowHammer bitflips** in nearby cells and breaks **memory isolation**

One Can Take Over an Otherwise-Secure System

Flipping Bits in Memory Without Accessing Them: An Experimental Study of DRAM Disturbance Errors

Abstract. *Memory isolation is a key property of a reliable and secure computing system — an access to one memory address should not have unintended side effects on data stored in other addresses. However, as DRAM process technology*

[Flipping Bits in Memory Without Accessing Them: An Experimental Study of DRAM Disturbance Errors](#) (Kim et al., ISCA 2014)

Project Zero

[Exploiting the DRAM rowhammer bug to gain kernel privileges](#) (Seaborn, 2015)

News and updates from the Project Zero team at Google

Induce bit flips in page table entries (PTEs).
Gain write access to its own page table,
and hence gain read-write access to all of physical memory.

Exploiting the DRAM rowhammer bug to gain kernel privileges

More Security Implications (I)

"We can gain unrestricted access to systems of website visitors."

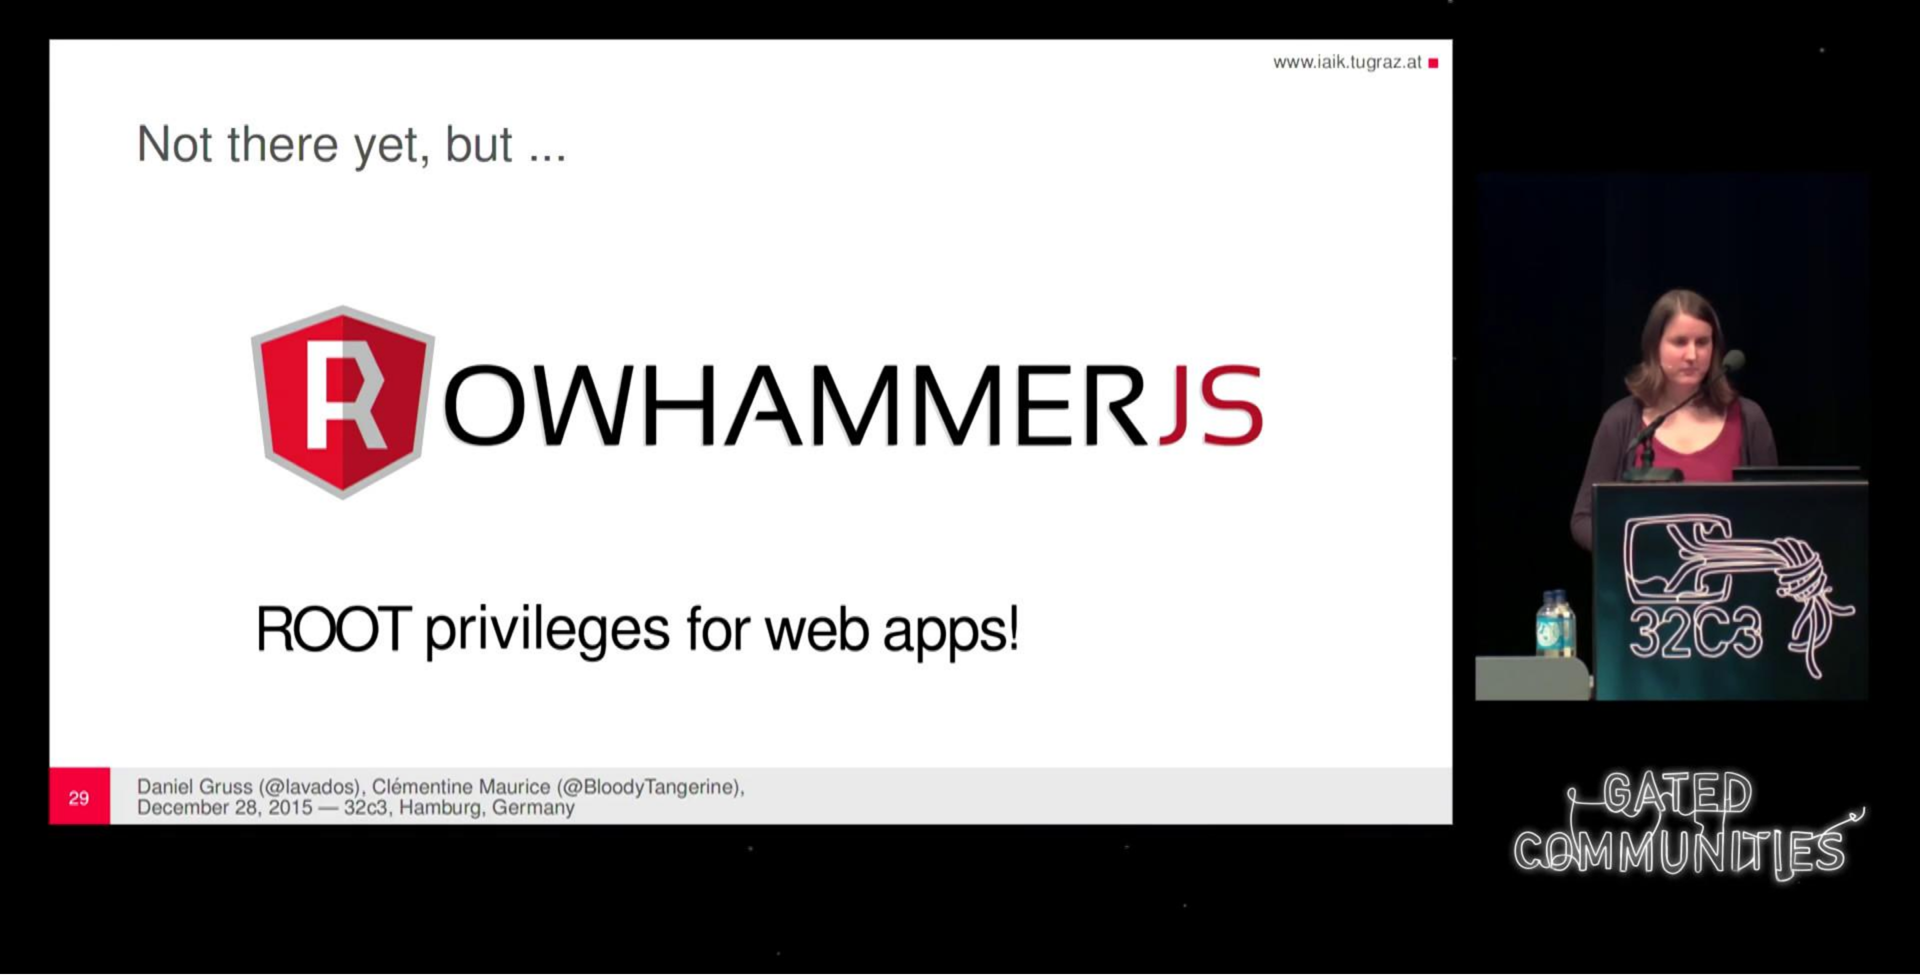

Rowhammer.js: A Remote Software-Induced Fault Attack in JavaScript (DIMVA'16)

More Security Implications (II)

"Can gain control of a smart phone deterministically"

Drammer: Deterministic Rowhammer Attacks on Mobile Platforms, CCS'16

Source: https://fossbytes.com/drammer-rowhammer-attack-android-root-devices/

More Security Implications (III)

- Using an integrated GPU in a mobile system to remotely escalate privilege via the WebGL interface. IEEE S&P 2018

ars TECHNICA BIZ & IT TECH SCIENCE POLICY CARS GAMING & CULTURE

"GRAND PWNING UNIT" —

Drive-by Rowhammer attack uses GPU to compromise an Android phone

JavaScript based GLitch pwns browsers by flipping bits inside memory chips.

Grand Pwning Unit: Accelerating Microarchitectural Attacks with the GPU

Pietro Frigo, Vrije Universiteit Amsterdam, p.frigo@vu.nl

Cristiano Giuffrida, Vrije Universiteit Amsterdam, giuffrida@cs.vu.nl

Herbert Bos, Vrije Universiteit Amsterdam, herbertb@cs.vu.nl

Kaveh Razavi, Vrije Universiteit Amsterdam, kaveh@cs.vu.nl

More Security Implications (IV)

- Rowhammer over RDMA (I)

ars TECHNICA

BIZ & IT TECH SCIENCE POLICY CARS GAMING & CULTURE

THROWHAMMER —

Packets over a LAN are all it takes to trigger serious Rowhammer bit flips

The bar for exploiting potentially serious DDR weakness keeps getting lower.

DAN GOODIN - 5/10/2018, 5:26 PM

Throwhammer: Rowhammer Attacks over the Network and Defenses

Andrei Tatar
VU Amsterdam

Radhesh Krishnan
VU Amsterdam

Elias Athanasopoulos
University of Cyprus

Cristiano Giuffrida
VU Amsterdam

Herbert Bos
VU Amsterdam

Kaveh Razavi
VU Amsterdam

More Security Implications (V)

- Rowhammer over RDMA (II)

The Hacker News™
Security in a serious way

Nethammer—Exploiting DRAM Rowhammer Bug Through Network Requests

Nethammer: Inducing Rowhammer Faults through Network Requests

Moritz Lipp
Graz University of Technology

Misiker Tadesse Aga
University of Michigan

Michael Schwarz
Graz University of Technology

Daniel Gruss
Graz University of Technology

Clémentine Maurice
Univ Rennes, CNRS, IRISA

Lukas Raab
Graz University of Technology

Lukas Lamster
Graz University of Technology

[Lipp+ EuroS&PW, 2020]

More Security Implications (VI)

JackHammer: Efficient Rowhammer on Heterogeneous FPGA-CPU Platforms

Zane Weissman[1], Thore Tiemann[2], Daniel Moghimi[1], Evan Custodio[3], Thomas Eisenbarth[2] and Berk Sunar[1]

[1] Worcester Polytechnic Institute, MA, USA
zweissman@wpi.edu, amoghimi@wpi.edu, sunar@wpi.edu
[2] University of Lübeck, Lübeck, Germany
thore.tiemann@student.uni-luebeck.de, thomas.eisenbarth@uni-luebeck.de
[3] Intel Corporation, Hudson, MA, USA
evan.custodio@intel.com

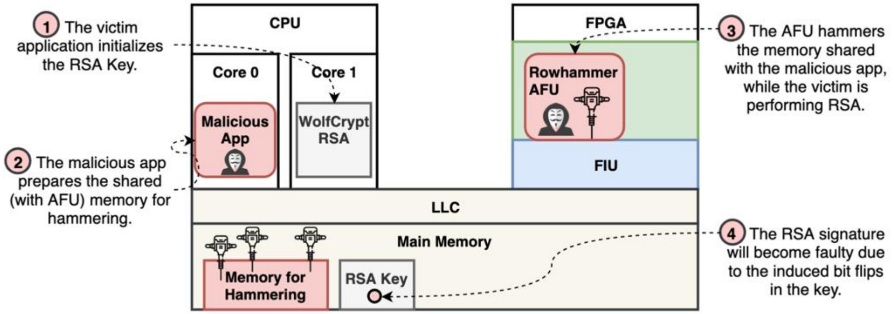

An **FPGA-based** RowHammer attack recovering **private keys** twice as fast compared to **CPU-based** attacks

More Security Implications (VII)

RAMBleed

RAMBleed: Reading Bits in Memory Without Accessing Them

Andrew Kwong
University of Michigan
ankwong@umich.edu

Daniel Genkin
University of Michigan
genkin@umich.edu

Daniel Gruss
Graz University of Technology
daniel.gruss@iaik.tugraz.at

Yuval Yarom
University of Adelaide and Data61
yval@cs.adelaide.edu.au

More Security Implications (VIII)

Terminal Brain Damage: Exposing the Graceless Degradation in Deep Neural Networks Under Hardware Fault Attacks

Sanghyun Hong, Pietro Frigo[†], Yiğitcan Kaya, Cristiano Giuffrida[†], Tudor Dumitraș

University of Maryland, College Park
[†]*Vrije Universiteit Amsterdam*

A Single Bit-flip Can Cause Terminal Brain Damage to DNNs

One specific bit-flip in a DNN's representation leads to accuracy drop over 90%

Our research found that a specific bit-flip in a DNN's bitwise representation can cause the accuracy loss up to 90%, and the DNN has 40-50% parameters, on average, that can lead to the accuracy drop over 10% when individually subjected to such single bitwise corruptions...

Read More

More Security Implications (IX)

DeepHammer: Depleting the Intelligence of Deep Neural Networks through Targeted Chain of Bit Flips

Fan Yao
University of Central Florida
fan.yao@ucf.edu

Adnan Siraj Rakin
Arizona State University
asrakin@asu.edu

Deliang Fan
Arizona State University
dfan@asu.edu

Degrade the **inference accuracy** to the level of **Random Guess**

Example: ResNet-20 for CIFAR-10, **10** output classes

Before attack, **Accuracy: 90.2%** After attack, **Accuracy: ~10% (1/10)**

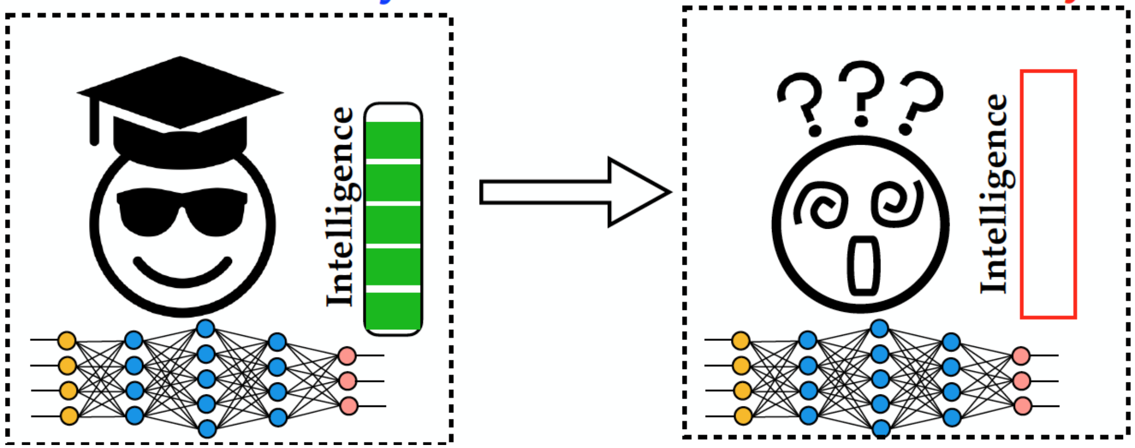

[Yao+ USENIX Security 2020]

More Security Implications (X)

HammerScope: Observing DRAM Power Consumption Using Rowhammer

Yaakov Cohen*
Ben-Gurion University of the Negev and Intel
Beer-Sheva, Israel
yaakoc@post.bgu.ac.il

Kevin Sam Tharayil*
Georgia Institute of Technology
Atlanta, Georgia, USA
kevinsam@gatech.edu

Arie Haenel*
Jerusalem College of Technology and Intel
Jerusalem, Israel
arie.haenel@jct.ac.il

Daniel Genkin
Georgia Institute of Technology
Atlanta, Georgia, USA
genkin@gatech.edu

Angelos D. Keromytis
Georgia Institute of Technology
Atlanta, Georgia, USA
angelos@gatech.edu

Yossi Oren
Ben-Gurion University of the Negev and Intel
Beer-Sheva, Israel
yos@bgu.ac.il

Yuval Yarom
University of Adelaide
Adelaide, Australia
yval@cs.adelaide.edu.au

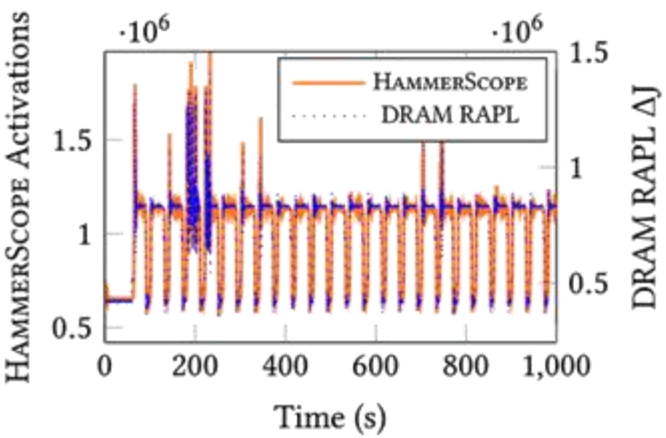

HammerScope is a **software-based power analysis** method using **RowHammer** as a side channel

More Security Implications (XI)

- **[Exploiting Correcting Codes: On the Effectiveness of ECC Memory Against Rowhammer Attacks]**. Cojocar, L. .; Razavi, K.; Giuffrida, C.; and Bos, H. In *S&P*, May 2019 *Best Practical Paper Award, Pwnie Award Nomination for Most Innovative Research*
  [Paper] [Slides]

Exploiting Correcting Codes: On the Effectiveness of ECC Memory Against Rowhammer Attacks

Lucian Cojocar, Kaveh Razavi, Cristiano Giuffrida, Herbert Bos
Vrije Universiteit Amsterdam

Thus, many believed that Rowhammer on ECC memory, even if plausible in theory, is simply impractical. This paper shows this to be false: while harder, ***Rowhammer attacks are still a realistic threat even to modern ECC-equipped systems****.*

More Security Implications (XII)

Hasan Hassan, Yahya Can Tugrul, Jeremie S. Kim, Victor van der Veen, Kaveh Razavi, and Onur Mutlu, **"Uncovering In-DRAM RowHammer Protection Mechanisms: A New Methodology, Custom RowHammer Patterns, and Implications,"** ***MICRO***, 2021. [Slides (pptx) (pdf)] [Short Talk Slides (pptx) (pdf)] [Lightning Talk Slides (pptx) (pdf)] [Full Talk (25 mins)] [Lightning Talk (1.5 mins)]

Uncovering In-DRAM RowHammer Protection Mechanisms: A New Methodology, Custom RowHammer Patterns, and Implications

Hasan Hassan[†] Yahya Can Tuğrul[†‡] Jeremie S. Kim[†] Victor van der Veen[σ] Kaveh Razavi[†] Onur Mutlu[†]

[†]*ETH Zürich* [‡]*TOBB University of Economics & Technology* [σ]*Qualcomm Technologies Inc.*

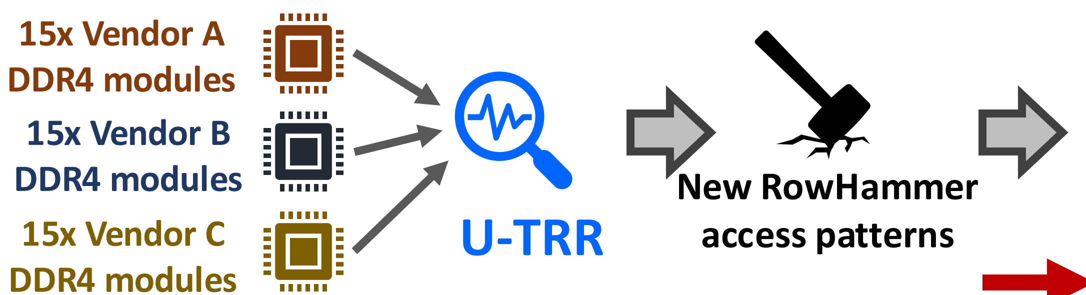

All 45 modules we test are **vulnerable**

99.9% of rows in a DRAM bank experience **at least one RowHammer bit flip**

Up to **7** RowHammer **bitflips** in an 8-byte dataword, **making ECC ineffective**

TRR does not provide security against RowHammer

U-TRR can facilitate the development of **new RowHammer attacks** and **more secure RowHammer protection** mechanisms

A RowHammer Survey:

Onur Mutlu, Ataberk Olgun, and A. Giray Yaglikci, **"Fundamentally Understanding and Solving RowHammer"** *Invited Special Session Paper at the 28th Asia and South Pacific Design Automation Conference (**ASP-DAC**)*, Tokyo, Japan, January 2023.
[arXiv version]
[Slides (pptx) (pdf)]
[Talk Video (26 minutes)]

Fundamentally Understanding and Solving RowHammer

Onur Mutlu
onur.mutlu@safari.ethz.ch
ETH Zürich
Zürich, Switzerland

Ataberk Olgun
ataberk.olgun@safari.ethz.ch
ETH Zürich
Zürich, Switzerland

A. Giray Yağlıkcı
giray.yaglikci@safari.ethz.ch
ETH Zürich
Zürich, Switzerland

https://arxiv.org/pdf/2211.07613.pdf

Breaking Samsung's Best Practice in 2024

- Salman Qazi and Daniel Moghimi, "SoothSayer: Bypassing DSAC Mitigation by Predicting Counter Replacement," DRAMSec, 2024

SoothSayer: Bypassing DSAC Mitigation by Predicting Counter Replacement

Salman Qazi
Google

Daniel Moghimi
Google

***Abstract*—In-DRAM Stochastic and Approximate Counting (DSAC) is a recently published algorithm that aims to mitigate Rowhammer at low cost. Existing in-DRAM counter-based schemes keep track of row activations and issue Targeted Row Refresh (TRR) upon detecting a concerning pattern. However, due to insufficiency of the tracking ability they are vulnerable to attacks utilizing decoy rows. DSAC claims to improve upon existing TRR mitigation by filtering out decoy-row accesses, so they cannot saturate the limited number of counters available for detecting Rowhammer, promising a reliable mitigation without the area cost of deterministic and provable schemes such as per-row activation counting (PRAC).**

In this paper, we analyze DSAC and discover some gaps that make it vulnerable to Rowhammer and Rowpress attacks.

The main focus of this work is a novel attack named SoothSayer that targets the counter replacement policy in DSAC by cloning the random number generator. We describe and simulate

literature such as Graphene [18] (implemented in the memory controller) and ProTRR [17] (implemented in DRAM) that utilize frequent item counting schemes. These can account for all Rowhammer activity if the threshold is sufficiently large and enough counters are provided. As the Rowhammer threshold decreases, the number of counters required for a correct implementation increases. According to the authors of DSAC, who are affiliated with a memory vendor, the number of counters used in these implementations are unacceptably large for a memory vendor to implement within their designs. To avoid this cost, deployed counter-based mitigations employ fewer counters than necessary and are often probabilistic. Due to this limitation, recent Rowhammer techniques [2], [8], [13] have managed to bypass TRR with decoy DRAM

RowHammer in DDR5

Patrick Jattke; Max Wipfli; Flavien Solt; Michele Marazzi; Matej Bölcskei; and Kaveh Razavi, **"ZenHammer: Rowhammer Attacks on AMD Zen-based Platforms,"** in *USENIX Security*, 2024.
[Paper] [URL]

ZENHAMMER: Rowhammer Attacks on AMD Zen-based Platforms

Patrick Jattke† Max Wipfli† Flavien Solt Michele Marazzi Matej Bölcskei Kaveh Razavi

ETH Zurich

†Equal contribution first authors

We found bit flips on only 1 of 10 tested devices (S1), suggesting that the changes in DDR5 such as improved Rowhammer mitigations, on-die error correction code (ECC), and a higher refresh rate (32 ms) make it harder to trigger bit flips.

DRAM Read Disturbance: A Critical Challenge

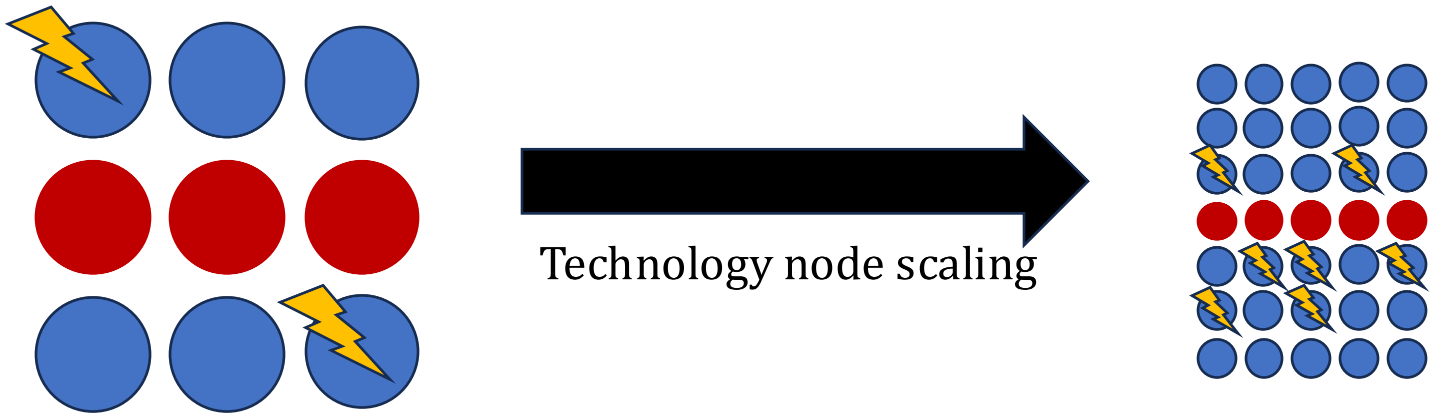

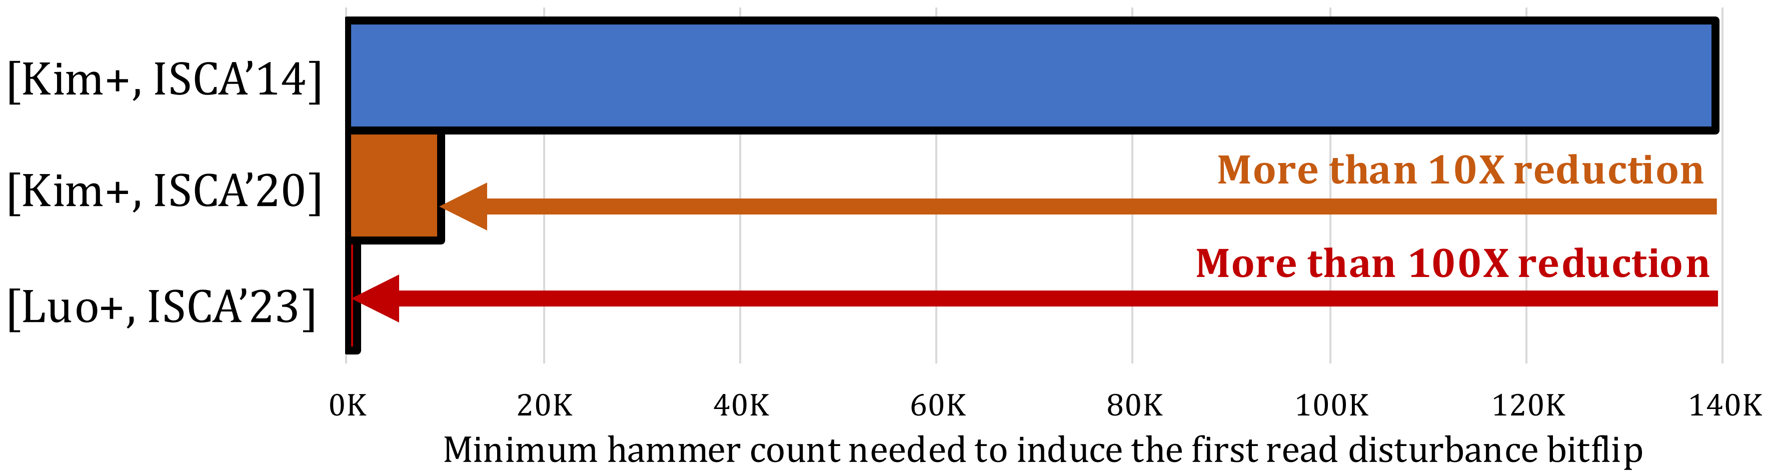

DRAM cells become **increasingly**
more vulnerable to **read disturbance**

Existing RowHammer Mitigations (I)

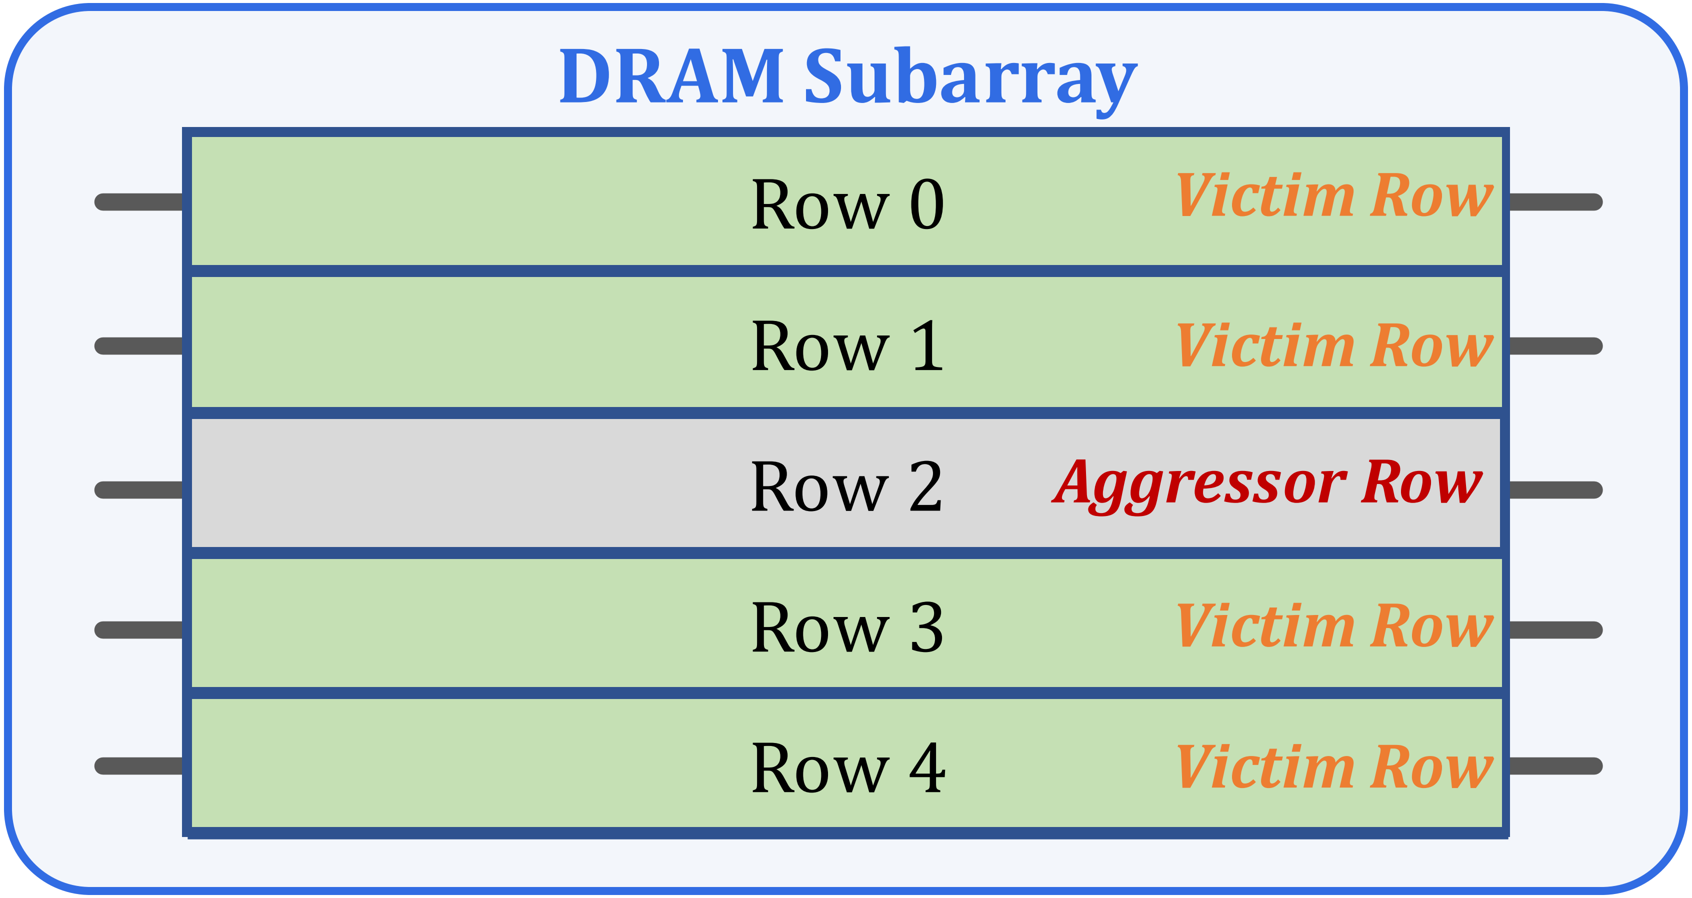

Refreshing potential victim rows
mitigates read disturbance bitflips

[Kim+ ISCA'20]

Existing RowHammer Mitigations (II)

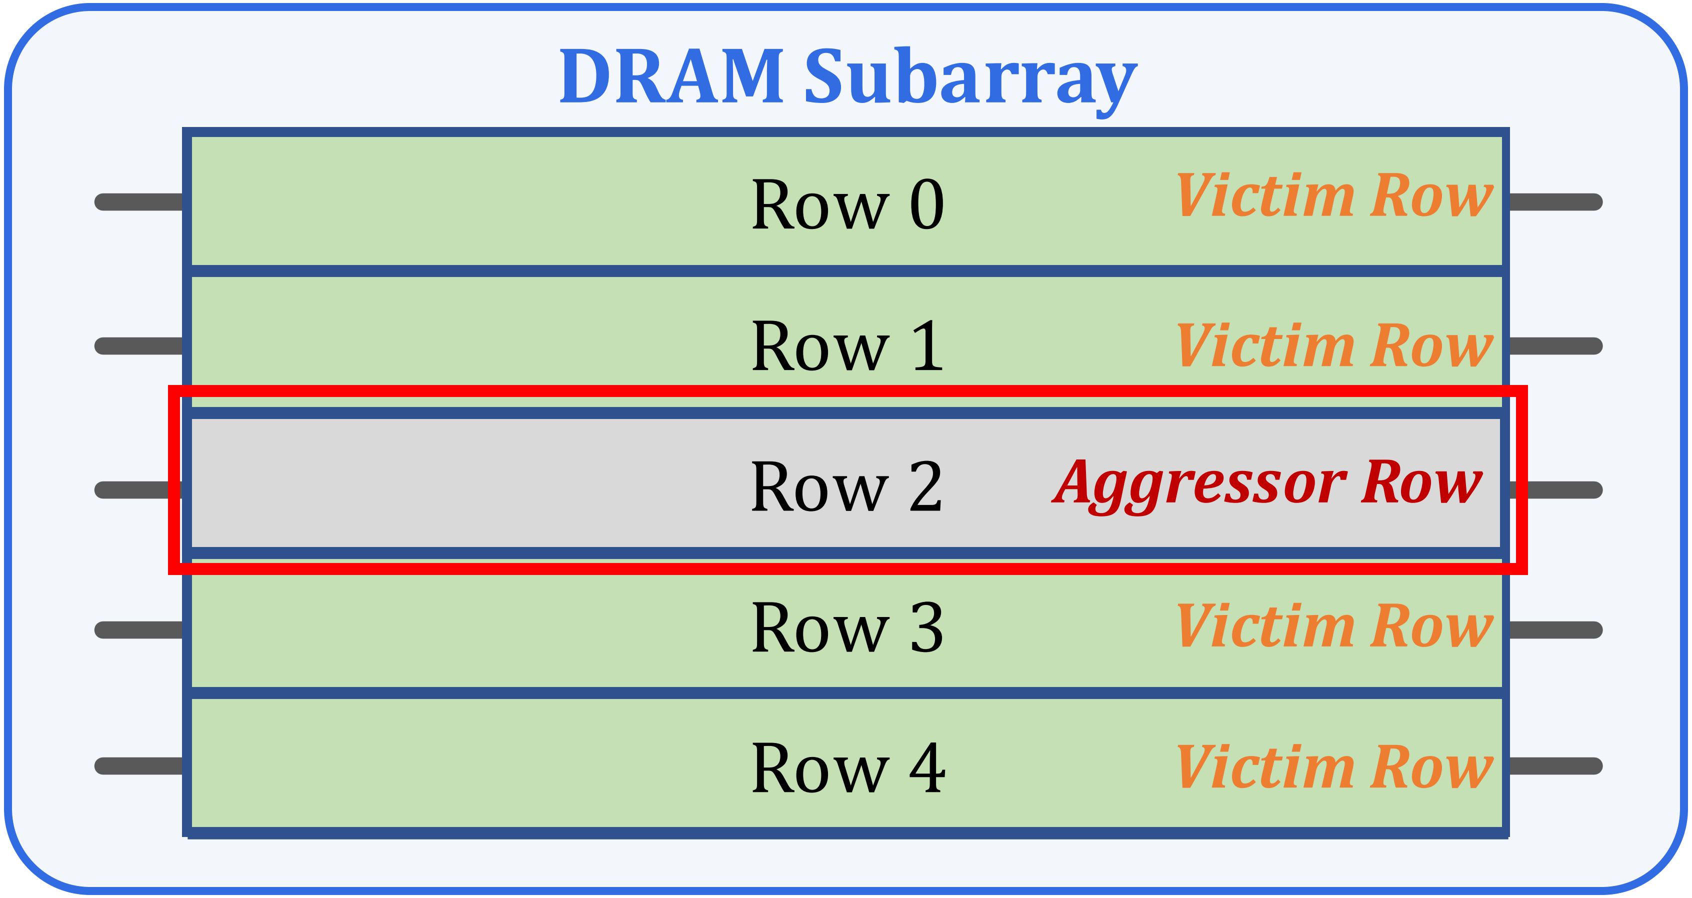

Mitigation techniques track DRAM row activation counts (of aggressor rows) to preventively refresh potential victim rows

[Kim+ ISCA'20]

Scalability with Worsening RowHammer Vulnerability

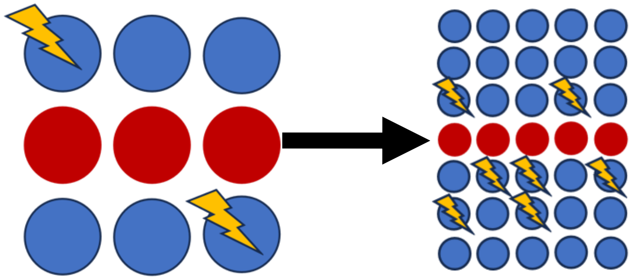

DRAM cells become **increasingly more vulnerable** to read disturbance at device level

Existing solutions become **prohibitively expensive** or **ineffective** going forward

[Kim+, ISCA'20] [Frigo+, IEEE S&P'20] [Hassan+, MICRO'21]

We need **more efficient and scalable** solutions to DRAM read disturbance

A **more detailed understanding** of DRAM read disturbance can help

Our Approach

We can **mitigate DRAM read disturbance efficiently** and **scalably** by

1 building a **detailed understanding** of DRAM read disturbance

MICRO'21 DSN'22

2 leveraging **insights into modern DRAM chips** and **memory controllers**

HPCA'24 MICRO'22

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

HPCA'21 MICRO'24

BreakHammer

Core Contributions

1 building a **detailed understanding** of DRAM read disturbance

MICRO'21

DSN'22

2 leveraging **insights into modern DRAM chips** and **memory controllers**

HPCA'24

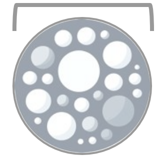

MICRO'22

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

HPCA'21

MICRO'24

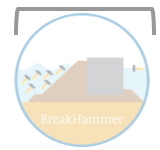

SAFARI

A Deeper Look into RowHammer

- Lois Orosa*, **Abdullah Giray Yağlıkçı***, Haocong Luo, Ataberk Olgun, Jisung Park, Hasan Hassan, Minesh Patel, Jeremie S. Kim, and Onur Mutlu,
  "A Deeper Look into RowHammer's Sensitivities: Experimental Analysis of Real DRAM Chips and Implications on Future Attacks and Defenses"
  *Proceedings of the 54th International Symposium on Microarchitecture (**MICRO**)*, Virtual, October 2021.
  [Slides (pptx) (pdf)] [Talk Video (21 minutes)]
  [Short Talk Slides (pptx) (pdf)]
  [Lightning Talk Slides (pptx) (pdf)] [Lightning Talk Video (1.5 minutes)]
  [arXiv version]

A Deeper Look into RowHammer's Sensitivities: Experimental Analysis of Real DRAM Chips and Implications on Future Attacks and Defenses

Lois Orosa* — ETH Zürich
A. Giray Yağlıkçı* — ETH Zürich
Haocong Luo — ETH Zürich
Ataberk Olgun — ETH Zürich, TOBB ETÜ
Jisung Park — ETH Zürich
Hasan Hassan — ETH Zürich
Minesh Patel — ETH Zürich
Jeremie S. Kim — ETH Zürich
Onur Mutlu — ETH Zürich

DRAM Testing Infrastructure

DRAM Bender on a Xilinx Virtex UltraScale+ XCU200

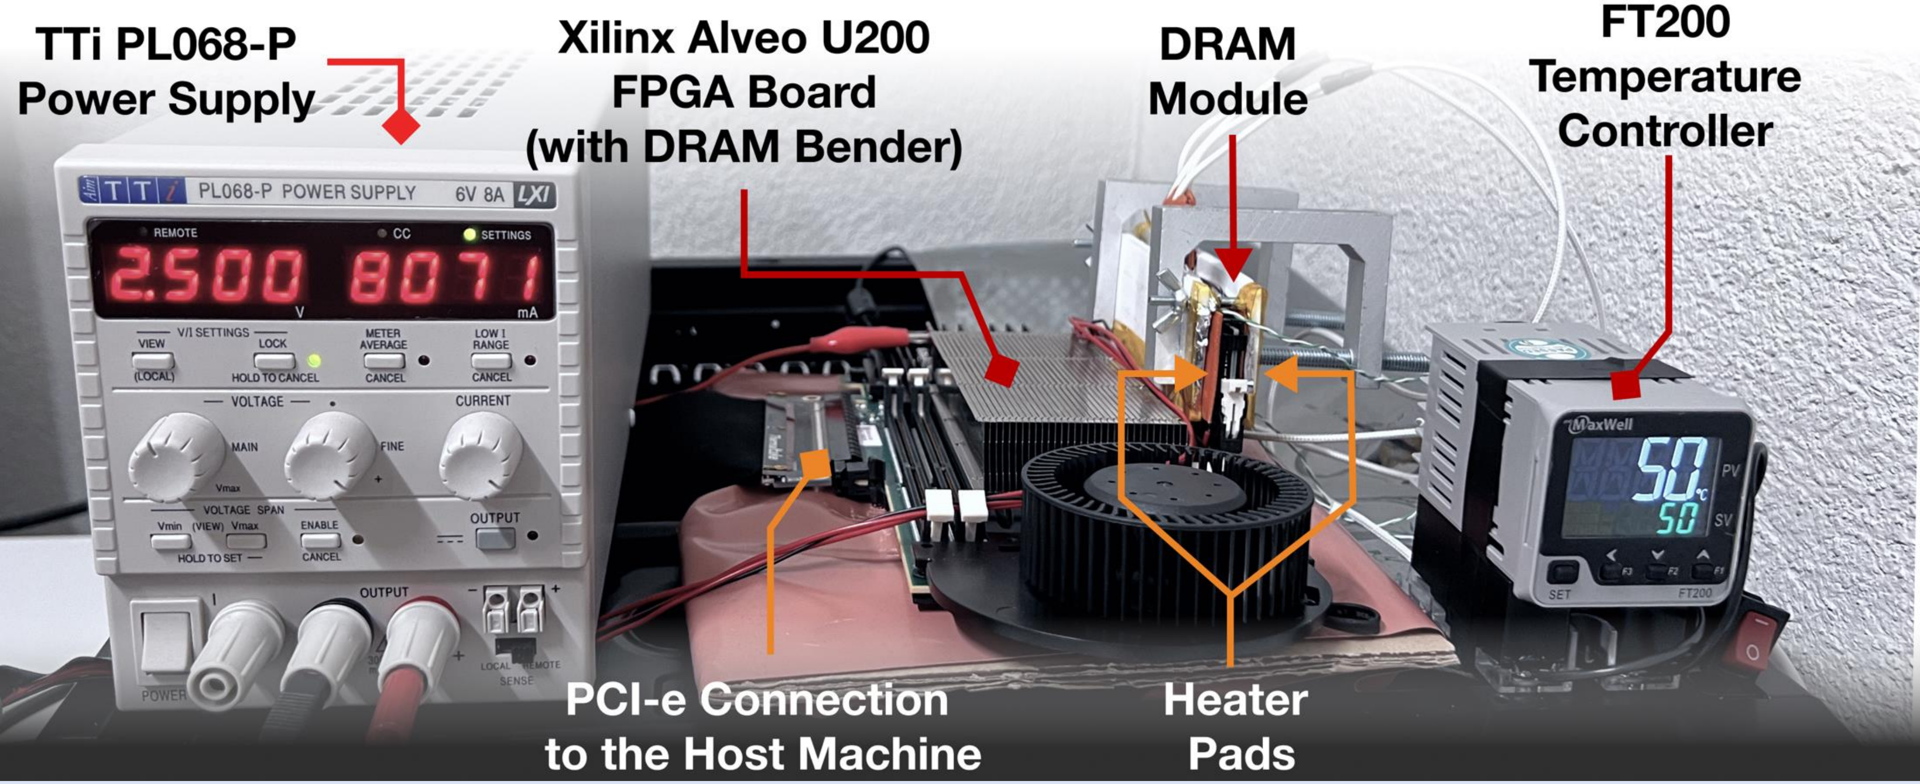

Fine-grained control over **DRAM commands**, **timing parameters (±1.5ns)**, **temperature (±0.5°C)**, and **voltage (±1mV)**

*Olgun et al., "DRAM Bender: An Extensible and Versatile FPGA-based Infrastructure to Easily Test State-of-the-art DRAM Chips," in TCAD, 2023. [GitHub: https://github.com/CMU-SAFARI/DRAM-Bender]

DRAM Testing Methodology

To characterize our DRAM chips at **worst-case** conditions:

1. Prevent sources of interference during core test loop

- **No DRAM refresh**: to avoid refreshing victim row
- **No DRAM calibration events**: to minimize variation in test timing
- **No RowHammer mitigation mechanisms**: to observe circuit-level effects
- Test for **less than a refresh window (32ms)** to avoid retention failures

2. Worst-case access sequence

- We use **worst-case** access sequence based on prior works' observations
- For each row, **repeatedly access the two physically-adjacent rows as fast as possible**

Tested DRAM Chips

- **272 DRAM Chips** (**24 DDR3** and **248 DDR4** DRAM Chips)

- **4 major manufacturers**:
  Micron, Samsung, SK Hynix, and Nanya

2 DRAM standards

Mfr.	DDR4 DIMMs	DDR3 SODIMMs	# Chips	Density	Die	Org.
A (Micron)	9	1	144 (8)	8Gb (4Gb)	B (P)	x4 (x8)
B (Samsung)	4	1	32 (8)	4Gb (4Gb)	F (Q)	x8 (x8)
C (SK Hynix)	5	1	40 (8)	4Gb (4Gb)	B (B)	x8 (x8)
D (Nanya)	4	-	32 (-)	8Gb (-)	C (-)	x8 (-)

4 major manufacturers

272 DRAM chips

Key Finding 1: Temperature Range

- DRAM read disturbance is more effective within a **bounded vulnerable temperature range**

- Vulnerable temperature range **varies across DRAM cells**

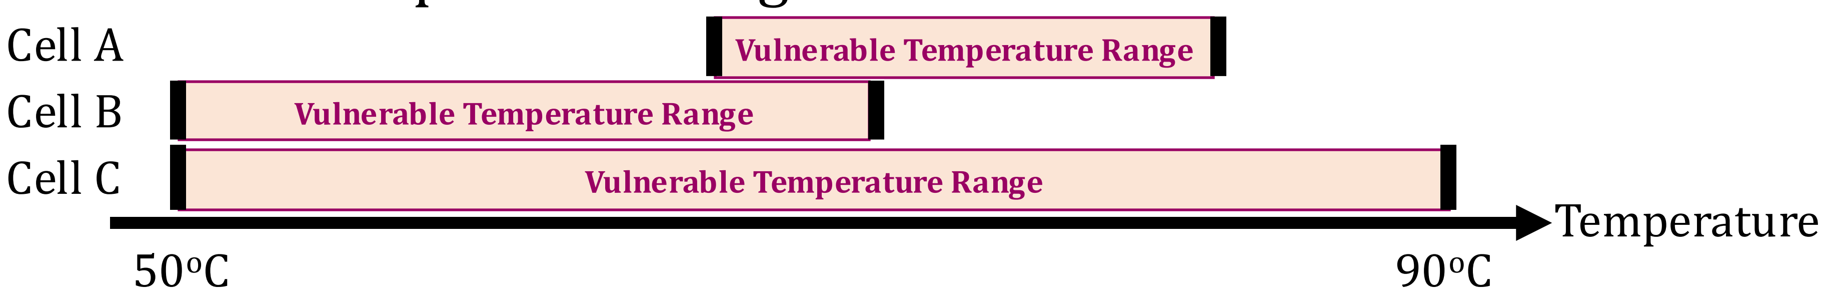

- Most cells are vulnerable in a **continuous temperature range**

Micron	Samsung	SK Hynix	Nanya
99.1%	98.9%	98.0%	99.2%

Orosa, Yağlıkçı et al., **"A Deeper Look into RowHammer's Sensitivities: Experimental Analysis of Real DRAM Chips and Implications on Future Attacks and Defenses"** in MICRO, 2021.

Circuit-Level Justification

We hypothesize that our observations are caused by the **non-monotonic behavior of charge trapping** characteristics of DRAM cells

3D TCAD model [Yang+, EDL'19]

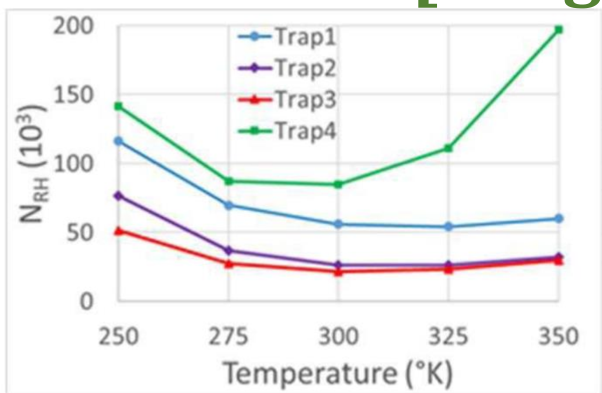

	Row	Column	Energy
Trap1	1	4	0.50eV
Trap2	2	3	0.45eV
Trap3	3	3	0.45eV
Trap4	4	3	0.45eV

Fig. 6. Hammering threshold N_{RH} vs. temperature from 250 to 350°K for different traps. Location in row and column refers to matrix in Fig. 2b.

***HC_{first}* decreases as temperature increases**, until a temperature inflection point where ***HC_{first}* starts to increase as temperature increases**

A **cell is more vulnerable** to RowHammer at **temperatures close to its temperature inflection point**

Key Finding 1: Temperature Range

Implication on testing

A DRAM cell should be tested
at **each possible** operating temperature

Implication on defenses

A DRAM cell **can be retired**
based on its **vulnerable temperature range**

Follow-up

This finding paved the way to **SpyHammer**,
a new attack that spies on temperature **[Orosa+, 2022]**

SAFARI

Orosa, Yağlıkçı et al., **"A Deeper Look into RowHammer's Sensitivities: Experimental Analysis of Real DRAM Chips and Implications on Future Attacks and Defenses"** in MICRO, 2021.

Key Finding 2: Aggressor Row On Time

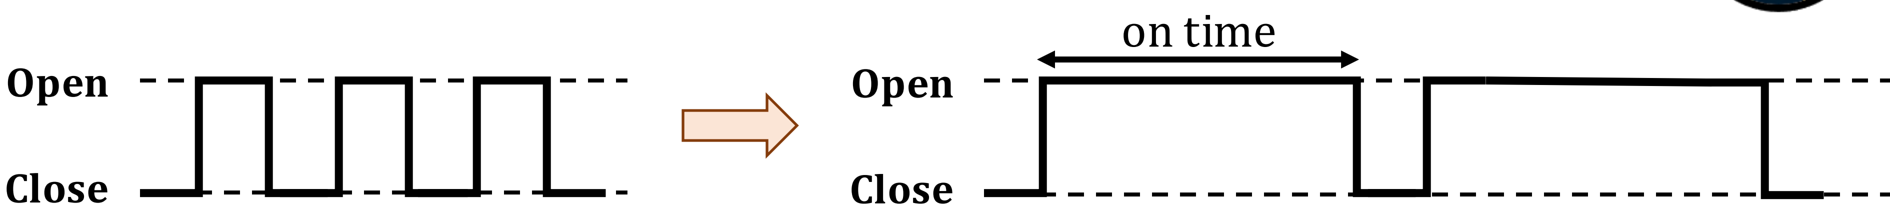

A smaller hammer count is sufficient to induce bitflips with increased **aggressor row on time**

Implication on defenses

Existing mitigations are **ineffective** without this insight

Follow-up

This finding paved the way to **the discovery of RowPress [Luo+, ISCA'23]**

Orosa, Yağlıkçı et al., **"A Deeper Look into RowHammer's Sensitivities: Experimental Analysis of Real DRAM Chips and Implications on Future Attacks and Defenses"** in MICRO, 2021.

Key Finding 3: Variation across DRAM Rows

Weakest 10% of rows

90% of rows

*Say, **bitflips occur** at **N hammers** in these rows*

*We need **at least 2N hammers** to induce bitflips in these rows*

RowHammer vulnerability **significantly varies** across DRAM rows

Implication on defenses

Configuring a solution for **the worst-case overprotects** many rows and makes it **expensive**

Follow-up

This finding paved the way to **the design of Svärd [Yaglikci+, HPCA'24]** (will be covered)

Orosa, Yağlıkçı et al., **"A Deeper Look into RowHammer's Sensitivities: Experimental Analysis of Real DRAM Chips and Implications on Future Attacks and Defenses"** in MICRO, 2021.

Core Contributions

1 building a **detailed understanding** of DRAM read disturbance

MICRO'21

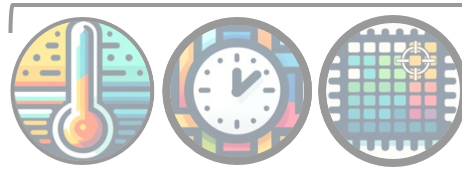

DSN'22

Voltage

2 leveraging **insights into modern DRAM chips** and **memory controllers**

HPCA'24

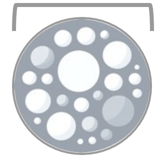

MICRO'22

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

HPCA'21

RowHammer Under Reduced Voltage

- A. Giray Yağlıkçı, Haocong Luo, Geraldo F. de Oliviera, Ataberk Olgun, Minesh Patel, Jisung Park, Hasan Hassan, Jeremie S. Kim, Lois Orosa, and Onur Mutlu, **"Understanding RowHammer Under Reduced Wordline Voltage: An Experimental Study Using Real DRAM Devices"**
  *Proceedings of the 52nd Annual IEEE/IFIP International Conference on Dependable Systems and Networks (**DSN**)*, Baltimore, MD, USA, June 2022.
  [Slides (pptx) (pdf)]
  [Lightning Talk Slides (pptx) (pdf)]
  [arXiv version]
  [Talk Video (34 minutes, including Q&A)]
  [Lightning Talk Video (2 minutes)]

Understanding RowHammer Under Reduced Wordline Voltage: An Experimental Study Using Real DRAM Devices

A. Giray Yağlıkçı[1] Haocong Luo[1] Geraldo F. de Oliviera[1] Ataberk Olgun[1] Minesh Patel[1]
Jisung Park[1] Hasan Hassan[1] Jeremie S. Kim[1] Lois Orosa[1,2] Onur Mutlu[1]
[1] *ETH Zürich* [2] *Galicia Supercomputing Center (CESGA)*

A Closer Look into RowHammer

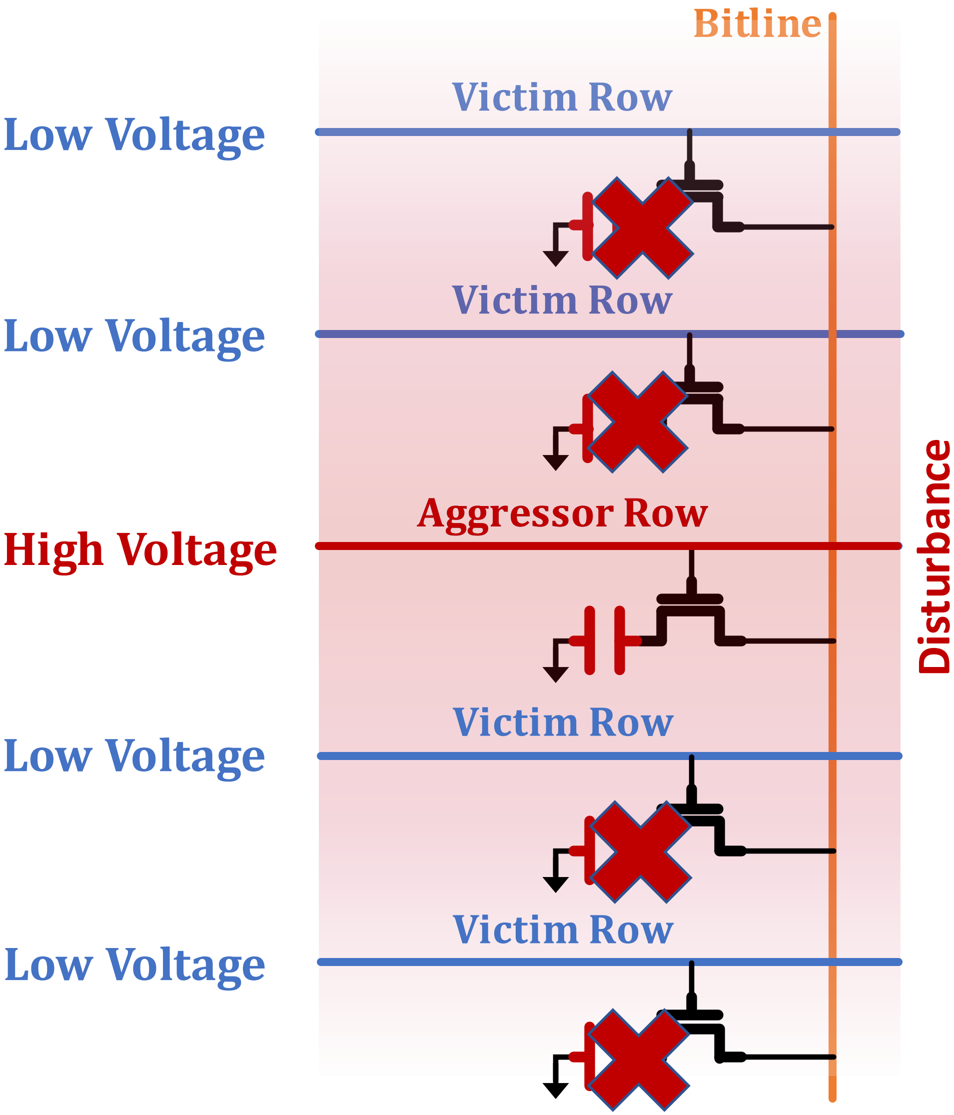

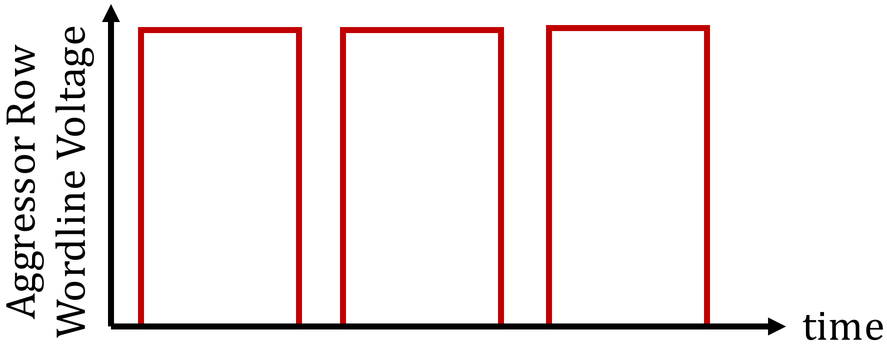

Repeatedly toggling wordline voltage leads to **RowHammer bitflips**

Yağlıkçı et al., **"Understanding RowHammer Under Reduced Wordline Voltage: An Experimental Study Using Real DRAM Devices,"** in DSN, 2022.

Effect of Reducing Wordline Voltage

Reducing **wordline voltage**
can **reduce RowHammer vulnerability**
at the cost of **weaker cells**

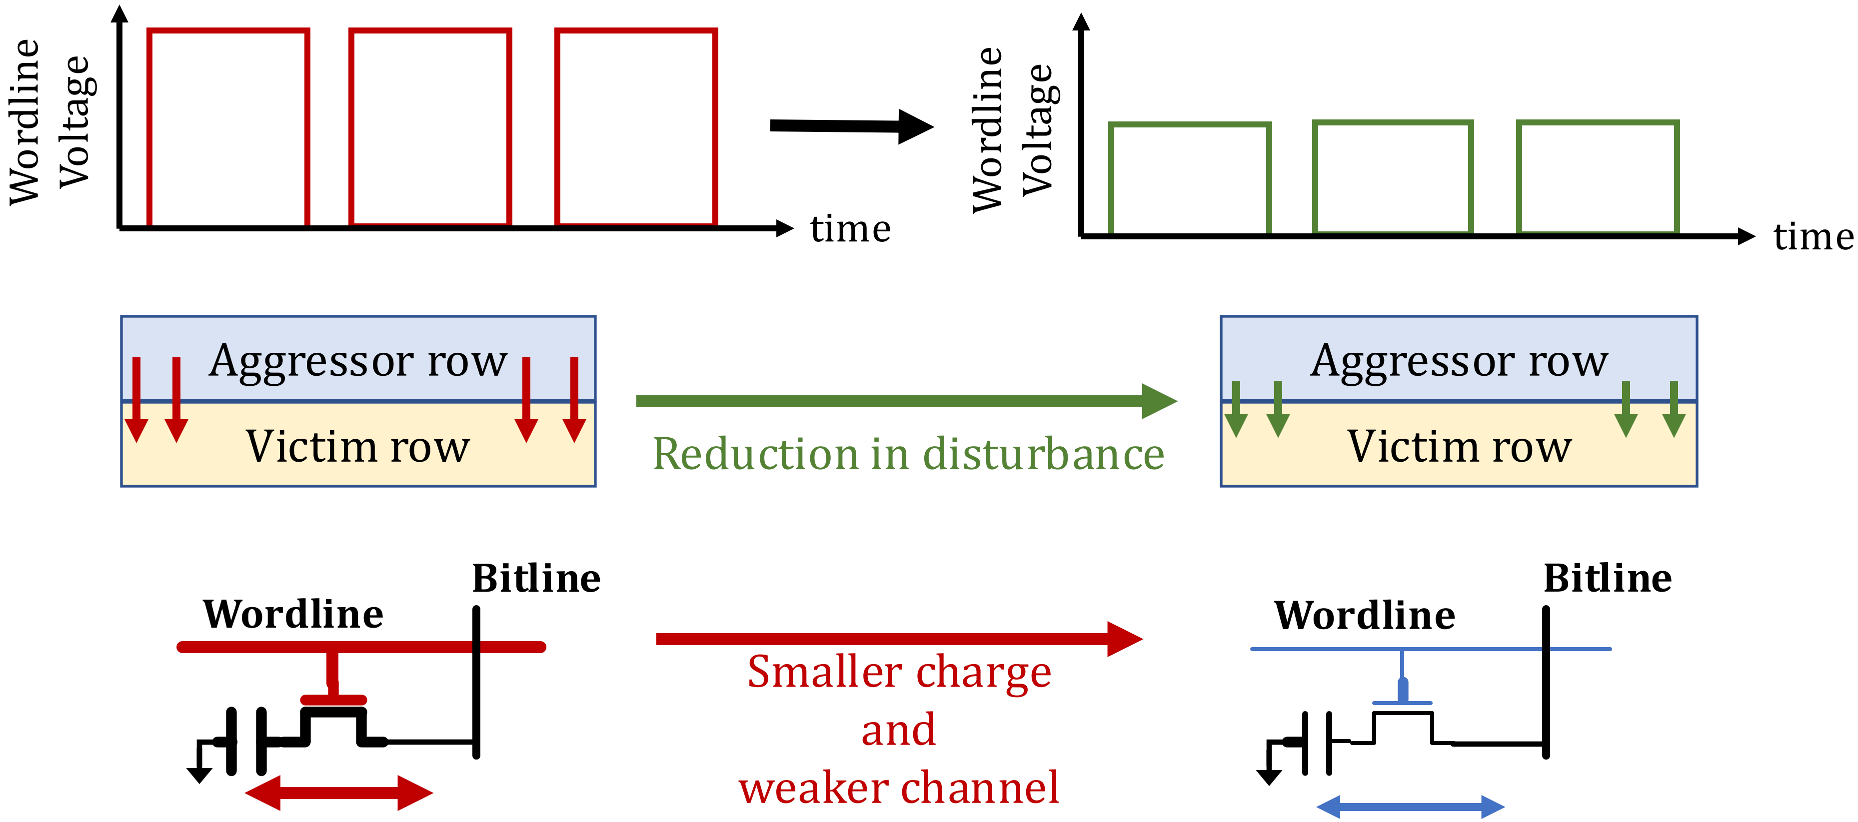

Key Results of Voltage Scaling Study

- Reducing wordline voltage can reduce RowHammer vulnerability
  - **15.2% (66.9% max)** fewer bitflips occur
  - **7.4% (85.8% max)** more activations needed to induce a bitflip

- Row activation latency **increases** with reduced **wordline voltage**
  - **208 / 272** DRAM chips have **no bitflips** at **nominal latency**
  - Changing timing constraint from 13.5ns to 24ns **avoids bitflips**

- **More DRAM cells** tend to experience **data retention bitflips** when **wordline voltage is reduced**
  - **216 / 272** DRAM chips have **no bitflips** at **nominal refresh rate**
  - **SECDED ECC at nominal refresh rate** avoids bitflips
  - **16% increase in refresh rate** avoids bitflips

Yağlıkçı et al., **"Understanding RowHammer Under Reduced Wordline Voltage: An Experimental Study Using Real DRAM Devices,"** in DSN, 2022.

Core Contributions

1 building a **detailed understanding** of DRAM read disturbance

MICRO'21

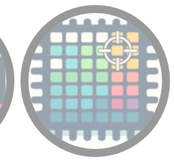

DSN'22

2 leveraging **insights into modern DRAM chips** and **memory controllers**

HPCA'24

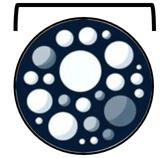

MICRO'22

Svärd: Leveraging heterogeneity

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

HPCA'21

MICRO'24

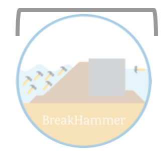

Spatial Variation-Aware Read Disturbance Defenses

- A. Giray Yağlıkçı, Geraldo F. de Oliviera, Yahya Can Tuğrul, İsmail Emir Yüksel, Ataberk Olgun, Haocong Luo, and Onur Mutlu,
"Spatial Variation-Aware Read Disturbance Defenses: Experimental Analysis of Real DRAM Chips and Implications on Future Solutions,"
*Proceedings of the 30th Edition of The **International Symposium on High-Performance Computer Architecture** (HPCA), Edinburgh, Scotland, UK, March 2024.* [Slides (pptx) (pdf)] [arXiv version]

Spatial Variation-Aware Read Disturbance Defenses: Experimental Analysis of Real DRAM Chips and Implications on Future Solutions

Abdullah Giray Yağlıkçı Geraldo F. Oliveira Yahya Can Tuğrul
İsmail Emir Yüksel Ataberk Olgun Haocong Luo Onur Mutlu
ETH Zürich

Tested DRAM Chips

144 DDR4 DRAM chips from
SK Hynix, Micron, and Samsung

All rows in **four different banks**

Mfr.	DIMM ID	# of Chips	Density Die Rev.	Chip Org.	Date (ww-yy)
Mfr. H (SK Hynix)	H0	8	16Gb – A	x8	51-20
	H1, H2, H3	3×8	16Gb – C	x8	48-20
	H4	8	8Gb – D	x8	48-20
Mfr. M (Micron)	M0	4	16Gb – E	x16	46-20
	M1, M3	2×16	8Gb – B	x4	N/A
	M2	16	16Gb – E	x4	14-20
	M4	4	16Gb – B	x16	26-21
Mfr. S (Samsung)	S0, S1	2×8	8Gb – B	x8	52-20
	S2	8	8Gb – B	x8	10-21
	S3	8	4Gb – F	x8	N/A
	S4	16	8Gb – C	x4	35-21

Spatial Variation-Aware Mitigation

- **Significant** and **irregular variation** in read disturbance vulnerability **across DRAM rows**

- Read disturbance solutions:
  - Configured **for the worst row**
  - **Overprotect** many rows
  - Incur **large performance overheads**

- **Key Idea: Dynamically tune** the aggressiveness of existing solutions to **the victim row's read disturbance vulnerability**

- **Svärd:** Spatial Variation-Aware Read Disturbance Defenses **Fewer preventive actions (e.g., refresh)** for **stronger rows**

*Svärd is the Swedish word for sword, a weapon with a long blade for cutting or thrusting [Merriam-Webster]

Integrating Svärd with Existing Solutions

PARA: Probabilistic Row Activation [Kim+, ISCA'14]

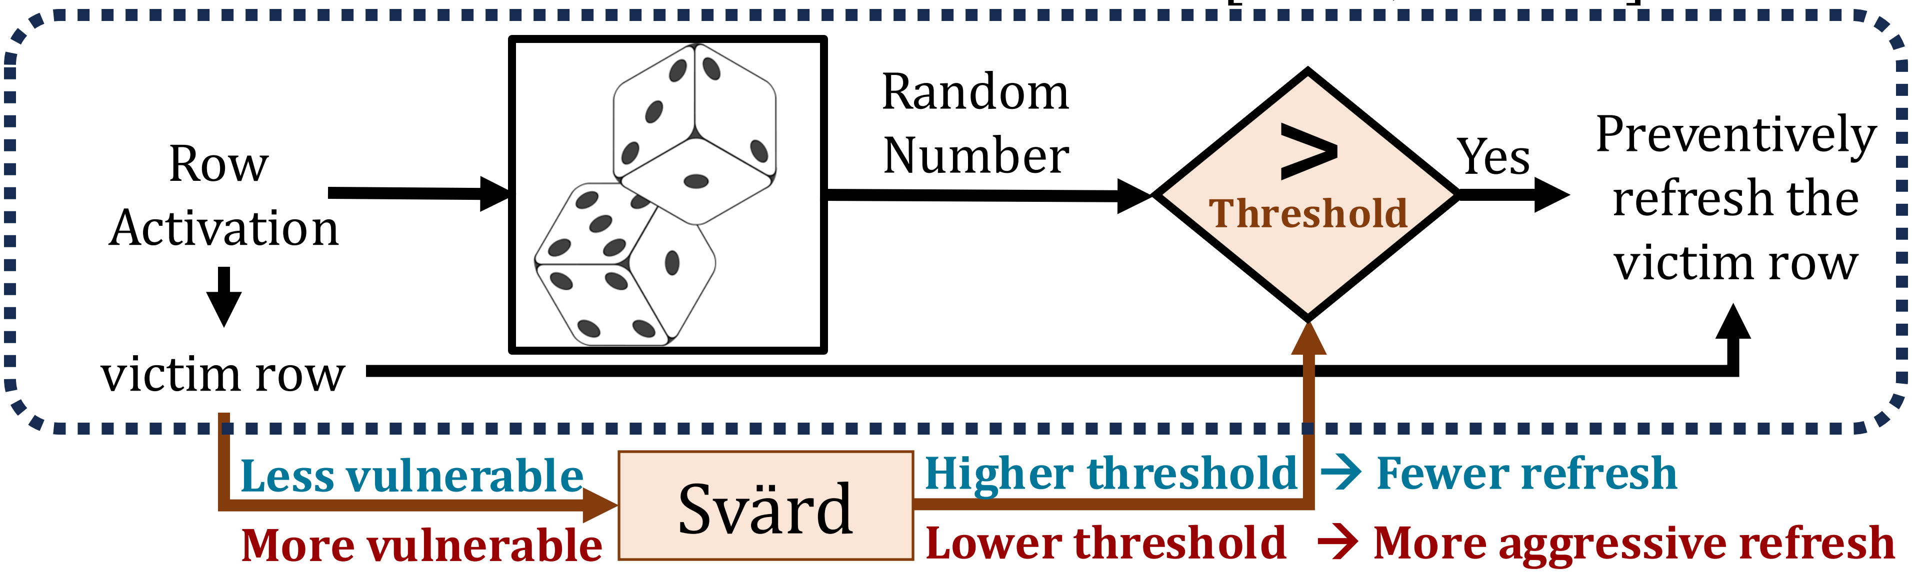

Svärd works with **many other** read disturbance solutions, including:

AQUA	BlockHammer	Hydra	RRS
[Saxena+, MICRO'22]	[Yaglikci+, HPCA'21]	[Qureshi+, ISCA'22]	[Saileshwar+, ASPLOS'22]

Performance Evaluation

- Cycle-level simulations using **Ramulator 2.0** [Luo+, CAL 2023]

- **System Configuration**:

Processor	3.2 GHz, 8 core, 4-wide issue, 128-entry instr. window
Last-Level Cache	64-byte cache line, 8-way set-associative, 8 MB
Memory Scheduler	FR-FCFS
Address Mapping	Minimalistic Open Pages
Main Memory	DDR4, 4 bank group, 4 banks per bank group (16 banks per rank)

- **Workloads**: **120** different **8-core** multiprogrammed workloads from **SPEC CPU2006, SPEC CPU2017, TPC, MediaBench,** and **YCSB** benchmark suites

- Integrated with **AQUA, BlockHammer, PARA, Hydra,** and **RRS**

- **HC_{first}**: {4K, 2K, 1K, 512, 256, 128, 64} hammers
  The **minimum hammer count** needed to induce **the first bitflip**

Yağlıkçı et al., **"Spatial Variation-Aware Read Disturbance Defenses: Experimental Analysis of Real DRAM Chips and Implications on Future Solutions,"** in HPCA, 2024.

Key Results

- Svärd **reduces the performance overhead** of existing solutions and **significantly improves system throughput**

AQUA
[Saxena+, MICRO'22]
1.6x

BlockHammer
[Yaglikci+, HPCA'21]
4.9x

Hydra
[Qureshi+, ISCA'22]
1.1x

PARA
[Kim+, ISCA'14]
2.0x

RRS
[Saileshwar+, ASPLOS'22]
4.8x

- Svärd's hardware complexity:

No Additional Latency

In-DRAM
Implementation
0.006%

In-Processor
Implementation
0.027%

Yağlıkçı et al., **"Spatial Variation-Aware Read Disturbance Defenses: Experimental Analysis of Real DRAM Chips and Implications on Future Solutions,"** in HPCA, 2024.

Core Contributions

1 building a **detailed understanding** of DRAM read disturbance

MICRO'21

DSN'22

2 leveraging **insights into modern DRAM chips** and **memory controllers**

HPCA'24

MICRO'22

HiRA: Parallelizing Preventive Actions

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

HPCA'21

MICRO'24

HiRA: Hidden Row Activation

- Abdullah Giray Yağlıkçı, Ataberk Olgun, Minesh Patel, Haocong Luo, Hasan Hassan, Lois Orosa, Oguz Ergin, and Onur Mutlu, **"HiRA: Hidden Row Activation for Reducing Refresh Latency of Off-the-Shelf DRAM Chips,"** in ***MICRO*** 2022.
  [Slides (pptx) (pdf)]
  [Longer Lecture Slides (pptx) (pdf)]
  [Lecture Video (36 minutes)]
  [arXiv version]

HiRA: Hidden Row Activation for Reducing Refresh Latency of Off-the-Shelf DRAM Chips

A. Giray Yağlıkçı[1] Ataberk Olgun[1] Minesh Patel[1] Haocong Luo[1] Hasan Hassan[1]
Lois Orosa[1,3] Oğuz Ergin[2] Onur Mutlu[1]
[1]*ETH Zürich* [2]TOBB University of Economics and Technology [3]*Galicia Supercomputing Center (CESGA)*

Subarray-Level Parallelism

Each subarray is mostly independent…

- except occasionally sharing ***global structures***

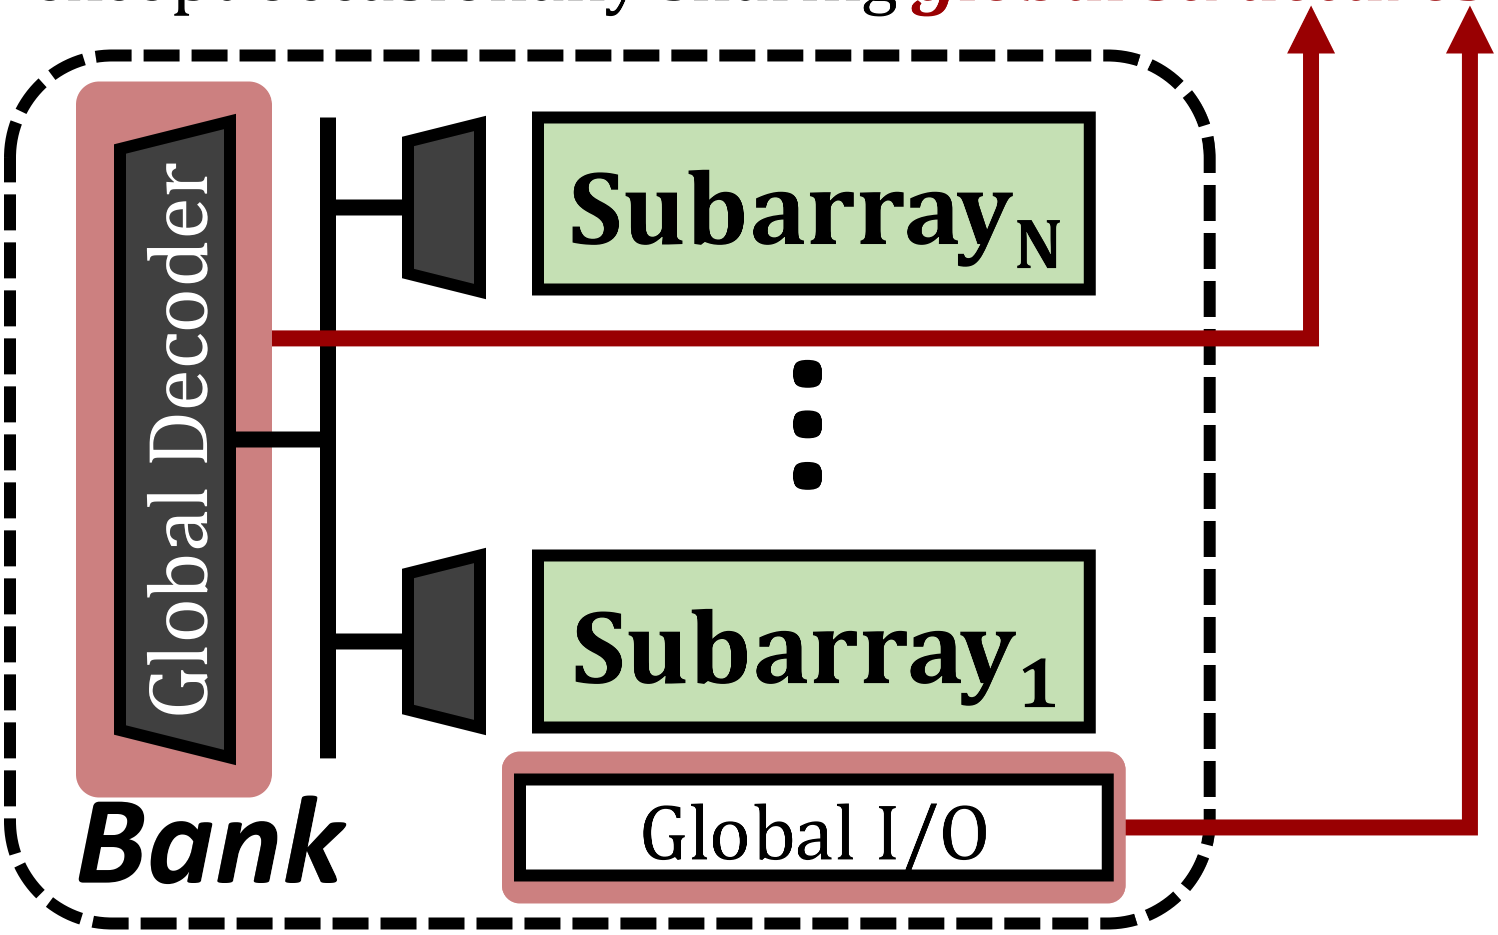

Yoongu Kim, et al., **"A Case for Exploiting Subarray-Level Parallelism (SALP) in DRAM,"** **in *ISCA***, 2012

HiRA: Hidden Row Activation – Key Insight

Concurrently activating two rows
in **different subarrays** of the **same bank**
can **refresh one row** while **opening the other row**

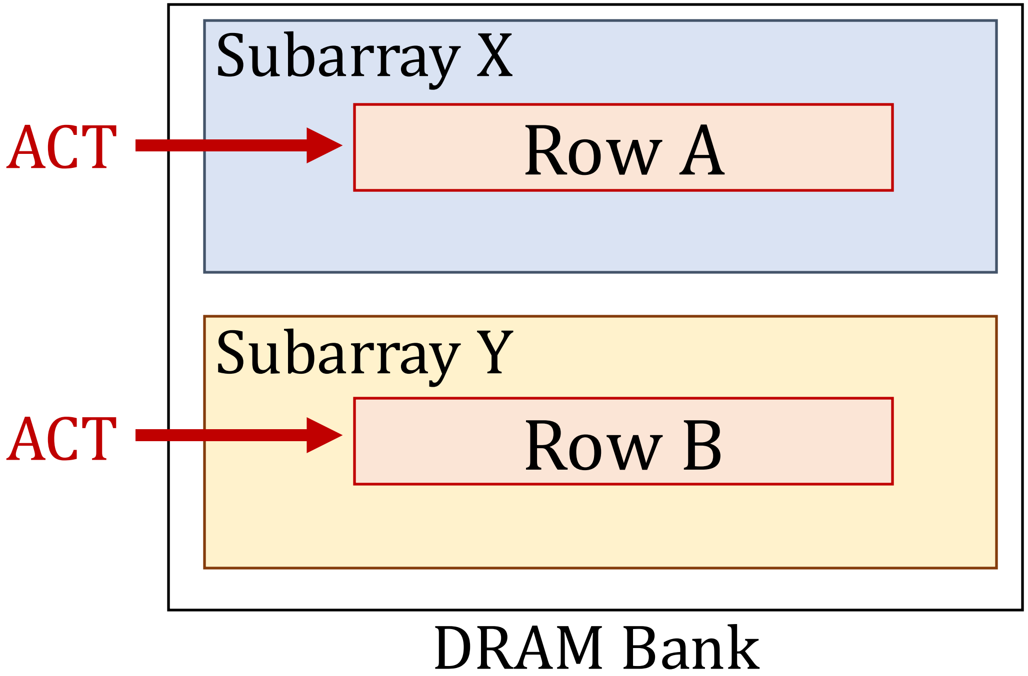

Refreshes RowA

concurrently with

opening RowB

Yağlıkçı et al., **"HiRA: Hidden Row Activation for Reducing Refresh Latency of Off-the-Shelf DRAM Chips,"** in MICRO, 2022.

HiRA: Hidden Row Activation

Refresh RowA concurrently with Activating RowB

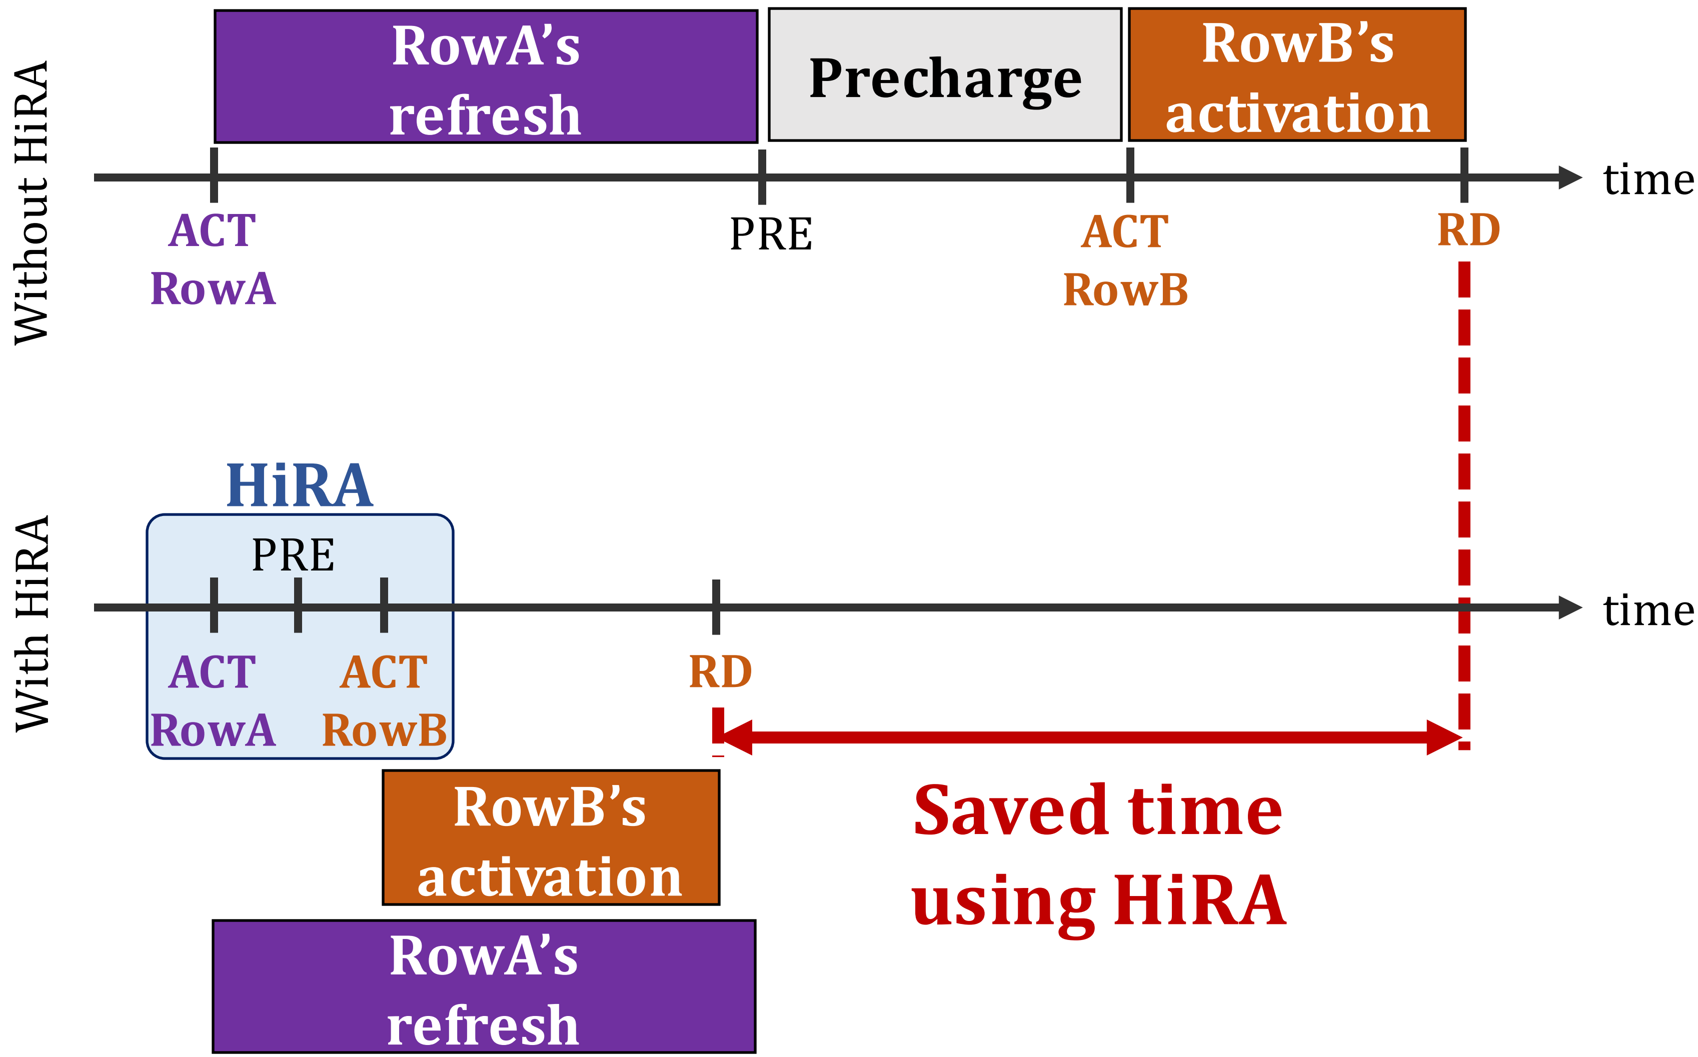

HiRA in Off-the-Shelf DRAM Chips

Table 4: Characteristics of the tested DDR4 DRAM modules.

Module Label	Module Vendor	Module Identifier Chip Identifier	Freq (MT/s)	Date Code	Chip Cap.	Die Rev.	Chip Org.	HiRA Coverage Min.	HiRA Coverage Avg.	HiRA Coverage Max.	Norm. N_{RH} Min.	Norm. N_{RH} Avg.	Norm. N_{RH} Max.
A0	G.SKILL	DWCW (Partial Marking)*	2400	42-20	4Gb	B	x8	24.8%	25.0%	25.5%	1.75	1.90	2.52
A1		F4-2400C17S-8GNT [39]						24.9%	26.6%	28.3%	1.72	1.94	2.55
B0	Kingston	H5AN8G8NDJR-XNC	2400	48-20	4Gb	D	x8	25.1%	32.6%	36.8%	1.71	1.89	2.34
B1		KSM32RD8/16HDR [87]						25.0%	31.6%	34.9%	1.74	1.91	2.51
C0	SK Hynix	H5ANAG8NAJR-XN	2400	51-20	4Gb	F	x8	25.3%	35.3%	39.5%	1.47	1.89	2.23
C1		HMAA4GU6AJR8N-XN [109]						29.2%	38.4%	49.9%	1.09	1.88	2.27
C2								26.5%	36.1%	42.3%	1.49	1.96	2.58

* The chip identifier is partially removed on these modules. We infer the chip manufacturer and die revision based on the remaining part of the chip identifier.

- HiRA works in **56 off-the-shelf DRAM chips** from **SK Hynix**
- **51.4% reduction** in the time spent for refresh operations

HiRA **effectively reduces the time spent** for **refresh** operations in **off-the-shelf** DRAM chips

HiRA-MC: HiRA Memory Controller

1 Generates each **periodic refresh** and **RowHammer-preventive refresh** **with a deadline**

2 **Buffers** each **refresh request** and **performs** the refresh request **until** the **deadline**

3 Finds if it can **refresh a DRAM row** concurrently with **a DRAM access** or **another refresh**

Performance Evaluation

- Cycle-level simulations using **Ramulator** [Kim+, CAL 2015]

- **System Configuration**:

Processor	3.2 GHz, 8 core, 4-wide issue, 128-entry instr. window
Last-Level Cache	64-byte cache line, 8-way set-associative, 8 MB
Memory Scheduler	FR-FCFS
Address Mapping	Minimalistic Open Pages
Main Memory	DDR4, 4 bank group, 4 banks per bank group (16 banks per rank)
Timing Parameters	t_1=t_2=3ns, t_{RC}= 46.25ns, t_{FAW}=16ns

- **Workloads**: **125** different **8-core** multiprogrammed workloads from the SPEC2006 benchmark suite

- **DRAM Chip Capacity**: {2, 4, 8, 16, 32, 64, 128} Gb

- **RowHammer Threshold**: {1024, 512, 256, 128, 64} activations
  The **minimum hammer count** needed to induce **the first RowHammer bitflip**

Key Results of HiRA

- **Performance Evaluation**
  - **3.7x speedup** by reducing time spent on **RowHammer-preventive refreshes**
  - **12.6% speedup** by reducing time spent on **periodic refreshes**

- **Hardware Complexity Analysis**
  - Chip **area cost of 0.0023%** of a processor die
  - **No additional latency** overhead

Yağlıkçı et al., **"HiRA: Hidden Row Activation for Reducing Refresh Latency of Off-the-Shelf DRAM Chips,"** in MICRO, 2022.

Core Contributions

1 building a **detailed understanding** of DRAM read disturbance

MICRO'21 DSN'22

2 leveraging **insights into modern DRAM chips** and **memory controllers**

HPCA'24 MICRO'22

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

HPCA'21 MICRO'24

BlockHammer: Selectively throttling unsafe accesses

BlockHammer

- A. Giray Yaglikci, Minesh Patel, Jeremie S. Kim, Roknoddin Azizi, Ataberk Olgun, Lois Orosa, Hasan Hassan, Jisung Park, Konstantinos Kanellopoulos, Taha Shahroodi, Saugata Ghose, and Onur Mutlu, **"BlockHammer: Preventing RowHammer at Low Cost by Blacklisting Rapidly-Accessed DRAM Rows"** *Proceedings of the 27th International Symposium on High-Performance Computer Architecture* (***HPCA***), Virtual, February-March 2021.
[Slides (pptx) (pdf)]
[Short Talk Slides (pptx) (pdf)]
[Intel Hardware Security Academic Awards Short Talk Slides (pptx) (pdf)]
[Talk Video (22 minutes)]
[Short Talk Video (7 minutes)]
[Intel Hardware Security Academic Awards Short Talk Video (2 minutes)]
[BlockHammer Source Code]
Intel Hardware Security Academic Award Finalist (one of 4 finalists out of 34 nominations)

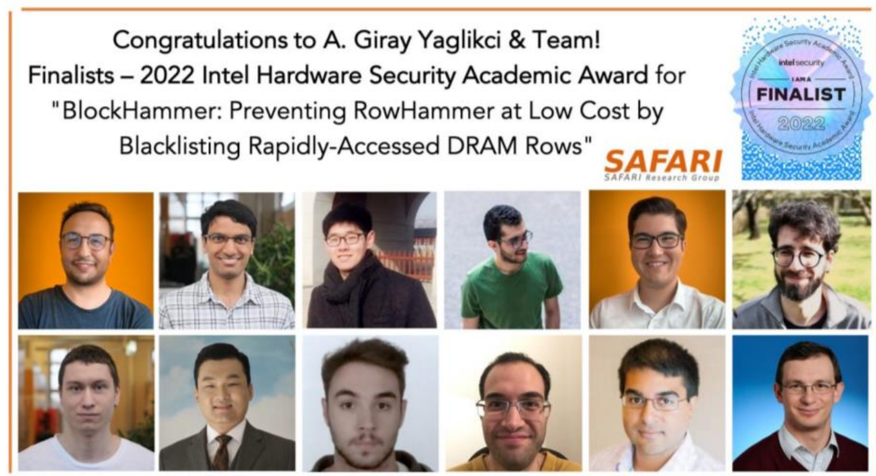

BlockHammer: Preventing RowHammer at Low Cost by Blacklisting Rapidly-Accessed DRAM Rows

A. Giray Yağlıkçı[1] Minesh Patel[1] Jeremie S. Kim[1] Roknoddin Azizi[1] Ataberk Olgun[1] Lois Orosa[1]
Hasan Hassan[1] Jisung Park[1] Konstantinos Kanellopoulos[1] Taha Shahroodi[1] Saugata Ghose[2] Onur Mutlu[1]
[1]*ETH Zürich* [2]*University of Illinois at Urbana–Champaign*

BlockHammer: Key Idea

- A RowHammer attack hammers Row A
- BlockHammer detects hammered rows using area-efficient Bloom filters
- **Selectively throttles** accesses targeting hammered rows

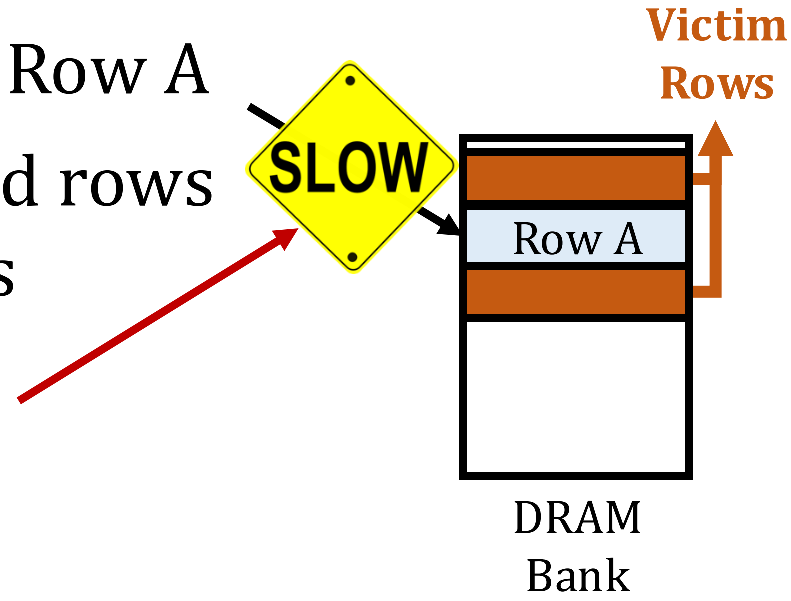

BlockHammer **prevents** read disturbance bitflips by **selectively throttling** the memory requests that would **otherwise cause** read disturbance **bitflips** and memory **contention**

BlockHammer: Evaluation (1/2)

- **Qualitative** comparison against **14 mechanisms**
  - Comprehensive protection
  - Compatibility with commodity DRAM chips
  - Scalability with worsening RowHammer vulnerability
  - Deterministic protection

BlockHammer is the **only mechanism** that satisfies all **four properties**

- **Quantitative** comparison against **6 state-of-the-art mechanisms**
  - PARA [Kim+, ISCA'14]
  - ProHIT [Son+, DAC'17]
  - MRLoc [You+, DAC'19]
  - CBT [Seyedzadeh+, ISCA'18]
  - TWiCe [Lee+, ISCA'19]
  - Graphene [Park+, MICRO'20]

BlockHammer is **low cost** and **competitive**

BlockHammer: Evaluation (2/2)

- Cycle-level simulations using **Ramulator** and **DRAMPower**
- System Configuration:

Processor	3.2 GHz, {1,8} core, 4-wide issue, 128-entry instr. window
LLC	64-byte cache line, 8-way set-associative, {2,16} MB
Memory scheduler	FR-FCFS
Address mapping	Minimalistic Open Pages
DRAM	DDR4 1 channel, 1 rank, 4 bank group, 4 banks per bank group
RowHammer Threshold	32K → 1K

- Single-Core Benign Workloads:
  - 22 SPEC CPU 2006
  - 4 YCSB Disk I/O
  - 2 Network Accelerator Traces
  - 2 Bulk Data Copy with Non-Temporal Hint (movnti)
- Randomly Chosen Multiprogrammed Workloads:
  - 125 workloads containing **8 benign applications**
  - 125 workloads containing **7 benign applications** and **1 RowHammer attack**

BlockHammer: Evaluation (2/2)

- Cycle-level simulations using **Ramulator** and **DRAMPower**

DRAM	DDR4 1 channel, 1 rank, 4 bank group, 4 banks per bank group
RowHammer Threshold	32K

No RowHammer attack:
Negligible (<0.6%) performance and energy overheads

- Single-Core Benign Workloads:

RowHammer attack present:
Significant improvement on system performance (71%) and energy consumption (32%)

- Randomly Chosen Multiprogrammed Workloads:
  - 125 workloads containing **8 benign applications**
  - 125 workloads containing **7 benign applications** and **1 RowHammer attack**

Core Contributions

1 building a **detailed understanding** of DRAM read disturbance

MICRO'21



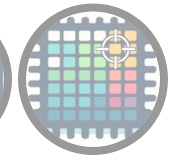

DSN'22

2 leveraging **insights into modern DRAM chips** and **memory controllers**

HPCA'24

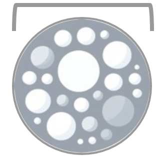

MICRO'22

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

HPCA'21

MICRO'24

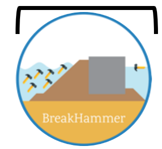

BreakHammer: Throttling malicious threads

Leveraging Adversarial Detection to Enable Scalable and Low Overhead RowHammer Mitigations

Oğuzhan Canpolat[§†] A. Giray Yağlıkçı[§] Ataberk Olgun[§] İsmail Emir Yüksel[§] Yahya Can Tuğrul[§†]
Konstantinos Kanellopoulos[§] Oğuz Ergin[†] Onur Mutlu[§]

[§]*ETH Zürich* [†]*TOBB University of Economics and Technology* [*]*SAFARI Research Group*

available on arXiv and will appear in MICRO 2024

https://arxiv.org/abs/2404.13477

Oğuzhan Canpolat, et al., "Leveraging Adversarial Detection to Enable Scalable and Low Overhead RowHammer Mitigations," [cs.CR] 2404.13477, 2024.

Motivation: Memory Performance Attacks

Attacker triggers **many** preventive actions
to **block** access to main memory

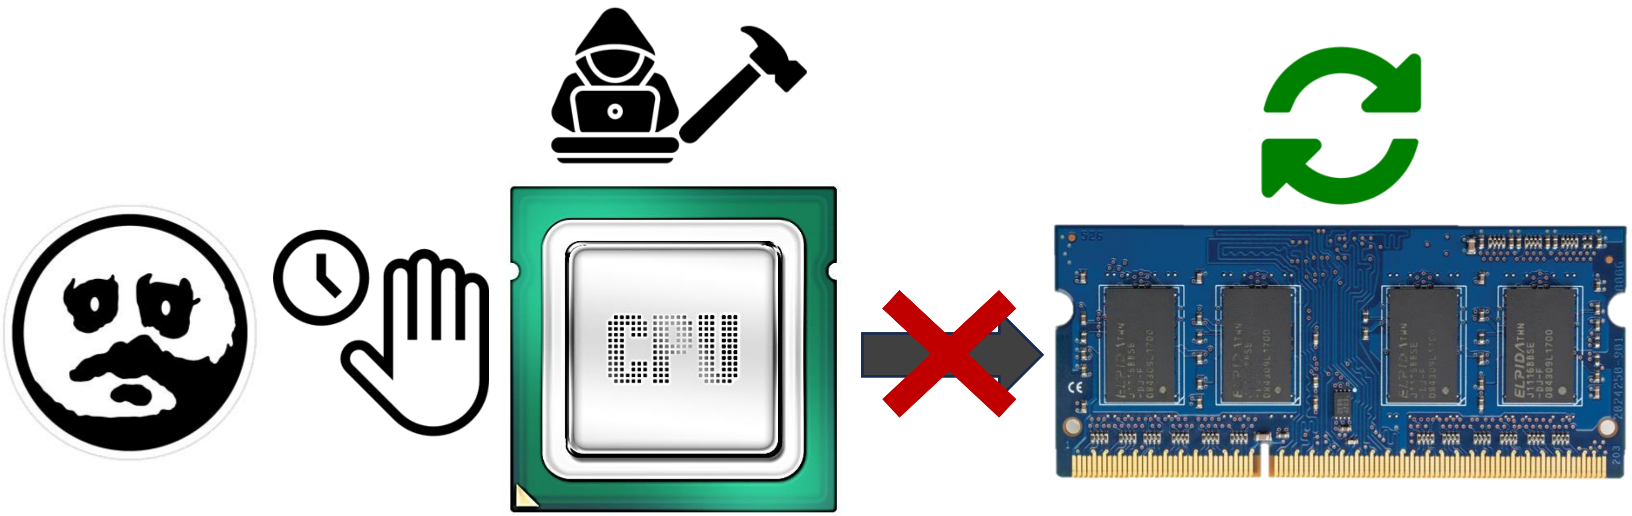

A new attack vector: Triggering preventive actions
to **reduce availability** of the main memory

BreakHammer: Key Idea

Detect and **slow down** attackers
that trigger **many** preventive actions

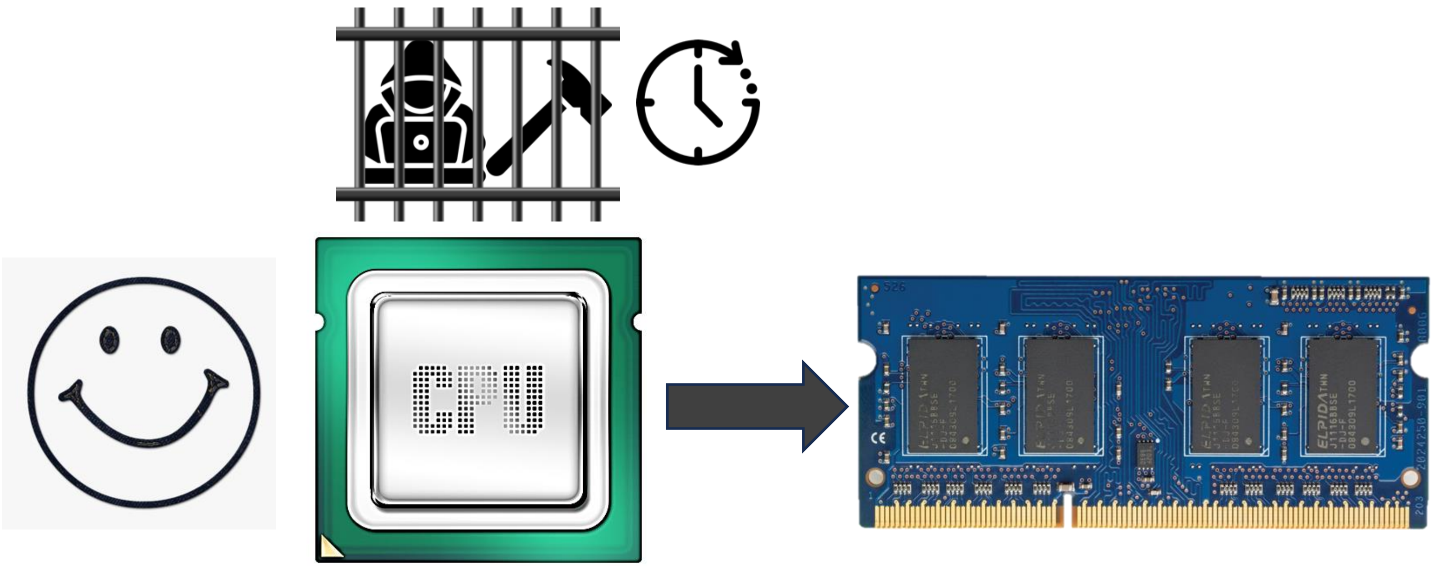

Evaluation

- **Performance and energy consumption evaluation:**
  Ramulator 2.0 [Luo+, CAL 2023] and DRAMPower [Chandrasekar+, DATE 2013]

- **System Configuration:**

Processor	4 cores, 4.2GHz clock frequency, 4-wide issue, 128-entry instruction window
DRAM	DDR5, 1 channel, 2 rank/channel, 8 bank groups, 4 banks/bank group, 64K rows/bank
Memory Ctrl.	64-entry read and write requests queues, Scheduling policy: FR-FCFS with a column cap of 4 LLC: 8 MiB (4-core)

- **Comparison Points:** Integrated with 7 state-of-the-art RowHammer mitigation mechanisms
  - PARA [Kim+ 2014]
  - Graphene [Park+ 2020]
  - AQUA [Saxena+ 2022]
  - Hydra [Quershi+ 2022]
  - TWiCe [Lee+ 2019]
  - REGA [Marazzi+ 2023]
  - RFM [JEDEC 2020]

- **Workloads:** 4-core workload mixes from SPEC CPU2006/2017, TPC, MediaBench, and YCSB
  - 60 mixes all benign
  - 60 mixes with 3 benign and 1 attacker

Evaluation

No attack

BreakHammer does **not** degrade the average system performance

Processor	4 cores, 4.2GHz clock frequency, 4-wide issue, 128-entry instruction window
DRAM	DDR5, 1 channel, 2 rank/channel, 8 bank groups, 4 banks/bank group, 64K rows/bank
Memory Ctrl.	64-entry read and write requests queues, Scheduling policy: FR-FCFS with a column cap of 4 LLC: 8 MiB (4-core)

Under attack

BreakHammer reduces preventive refresh count (by 67%)

- AQUA [Saxena+ 2022]

ershi+ 2022]

Under attack

BreakHammer improves system performance (by 40%)

s: 4-core workload mixes from SPEC CPU2006/2017, TPC, MediaBench, and YCSB

Under attack

BreakHammer reduces DRAM energy consumption (by 54%)

Conclusion

We can **mitigate DRAM read disturbance efficiently** and **scalably** by

1 building a **detailed understanding** of DRAM read disturbance

2 leveraging **insights into modern DRAM chips** and **memory controllers**

3 **throttling** memory requests that lead to **time-consuming preventive refreshes**

More Details and Discussion on YouTube

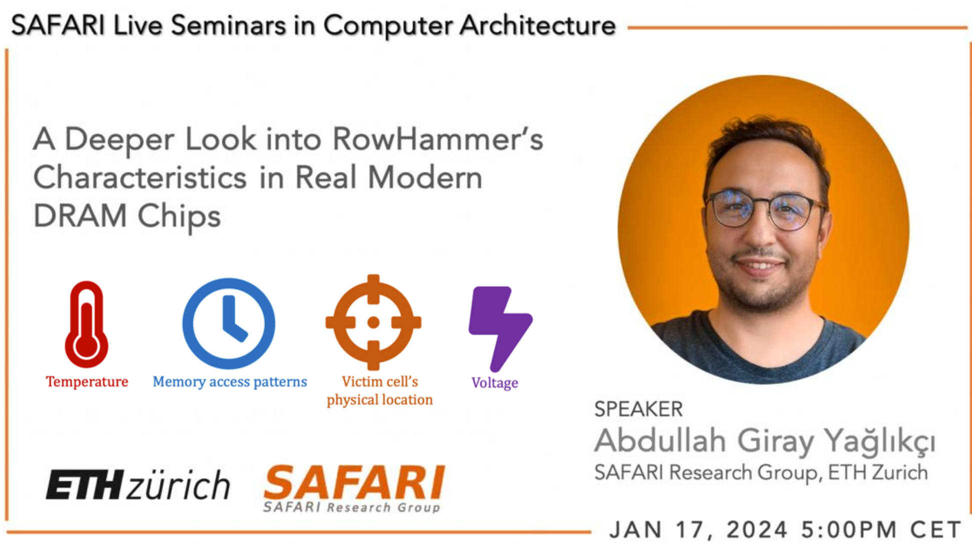

https://www.youtube.com/live/CRtm1es4n3o?si=8N5zB6e_RUc5Ejl8

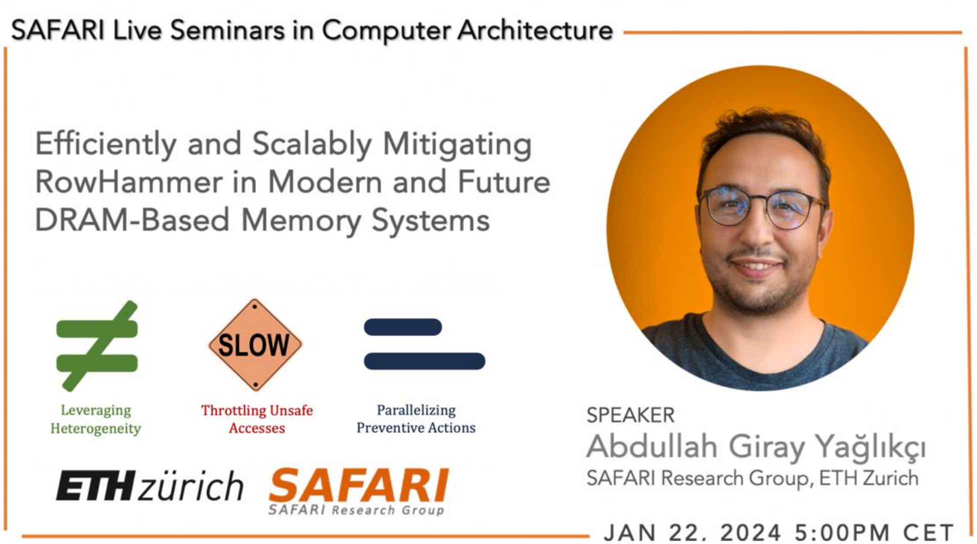

https://www.youtube.com/live/YQwRYWpCsk0?si=jXPueMHb5wgs69-q

Industry Solutions to RowHammer

Per-Row Activation Counters in DRAM

- Tanj Bennett, Stefan Saroiu, Alec Wolman, and Lucian Cojocar, "[Panopticon: A Complete In-DRAM Rowhammer Mitigation](),"
DRAMSec Workshop (co-located with ISCA), 2021.

Panopticon: A Complete In-DRAM Rowhammer Mitigation

Tanj Bennett[§], Stefan Saroiu, Alec Wolman, and Lucian Cojocar
Microsoft, [§]Avant-Gray LLC

Accurate RowHammer detection using **an activation counter per DRAM row** stored within the **DRAM array** leveraging **high density → low cost**

RowHammer in 2023: SK Hynix

ISSCC 2023 / SESSION 28 / HIGH-DENSITY MEMORIES /

28.8 A 1.1V 16Gb DDR5 DRAM with Probabilistic-Aggressor Tracking, Refresh-Management Functionality, Per-Row Hammer Tracking, a Multi-Step Precharge, and Core-Bias Modulation for Security and Reliability Enhancement

Woongrae Kim, Chulmoon Jung, Seongnyuh Yoo, Duckhwa Hong, Jeongjin Hwang, Jungmin Yoon, Ohyong Jung, Joonwoo Choi, Sanga Hyun, Mankeun Kang, Sangho Lee, Dohong Kim, Sanghyun Ku, Donhyun Choi, Nogeun Joo, Sangwoo Yoon, Junseok Noh, Byeongyong Go, Cheolhoe Kim, Sunil Hwang, Mihyun Hwang, Seol-Min Yi, Hyungmin Kim, Sanghyuk Heo, Yeonsu Jang, Kyoungchul Jang, Shinho Chu, Yoonna Oh, Kwidong Kim, Junghyun Kim, Soohwan Kim, Jeongtae Hwang, Sangil Park, Junphyo Lee, Inchul Jeong, Joohwan Cho, Jonghwan Kim

SK hynix Semiconductor, Icheon, Korea

RowHammer in 2023: Samsung

DSAC: Low-Cost Rowhammer Mitigation Using In-DRAM Stochastic and Approximate Counting Algorithm

Seungki Hong Dongha Kim Jaehyung Lee Reum Oh
Changsik Yoo Sangjoon Hwang Jooyoung Lee

DRAM Design Team, Memory Division, Samsung Electronics

https://arxiv.org/pdf/2302.03591v1.pdf

DDR5 Update in 2024

- DRAM implements per-row activation counters
  similar to Panopticon and PRHT

- DRAM asserts a back-off signal when refresh is needed
  similar to SMD, Mithril+, and Panopticon

- Memory controller issues refresh upon back-off signal

JEDEC STANDARD

DDR5 SDRAM

JESD79-5C_v1.30

(Revision of JESD79-5B_v1.20, September 2022)

April 2024

JEDEC SOLID STATE TECHNOLOGY ASSOCIATION

JEDEC®

JEDEC, "**JESD79-5C_v1.30: DDR5 SDRAM Specification**," April, 2024.

Understanding PRAC

Understanding the Security Benefits and Overheads of Emerging Industry Solutions to DRAM Read Disturbance

Oğuzhan Canpolat[§†] A. Giray Yağlıkçı[§] Geraldo F. Oliveira[§] Ataberk Olgun[§]
Oğuz Ergin[†] Onur Mutlu[§]
[§]*ETH Zürich* [†]*TOBB University of Economics and Technology*

presented in DRAMSec 2024

- PRAC is the latest update to DDR5 in April 2024
- We show that
  - PRAC is secure and has non-negligible performance, energy, and area overheads for benign workloads
  - PRAC can be exploited by memory performance attacks (up to **79% reduction** in DRAM chip's throughput)

Industry Solutions to Read Disturbance: Per Row Activation Counting DRAM Timings

Counters	DRAM Rows
0	10101010101010101010101010
0	10101010101010101010101010
0	10101010101010101010101010
0	10101010101010101010101010
0	10101010101010101010101010

Row counter updates are **not** completely parallelized with DRAM access

PRAC increases row-close time (t_{RP}) by **~140%**

Industry Solutions to Read Disturbance: Per Row Activation Counting DRAM Timings

Timing parameter changes for DDR5-3200AN speed bin
[JEDEC JESD79-5C, April 2024]

t_{RP} : +21ns (+140%)

t_{RAS} : -16ns (-50%)

t_{RTP} : -2.5ns (-33%)

t_{WR} : -20ns (-66%)

t_{RC} : +5ns (+10%)

Memory Performance Attacks

Access pattern to trigger **most** back-offs
with **fewest** activations possible by targeting a single row

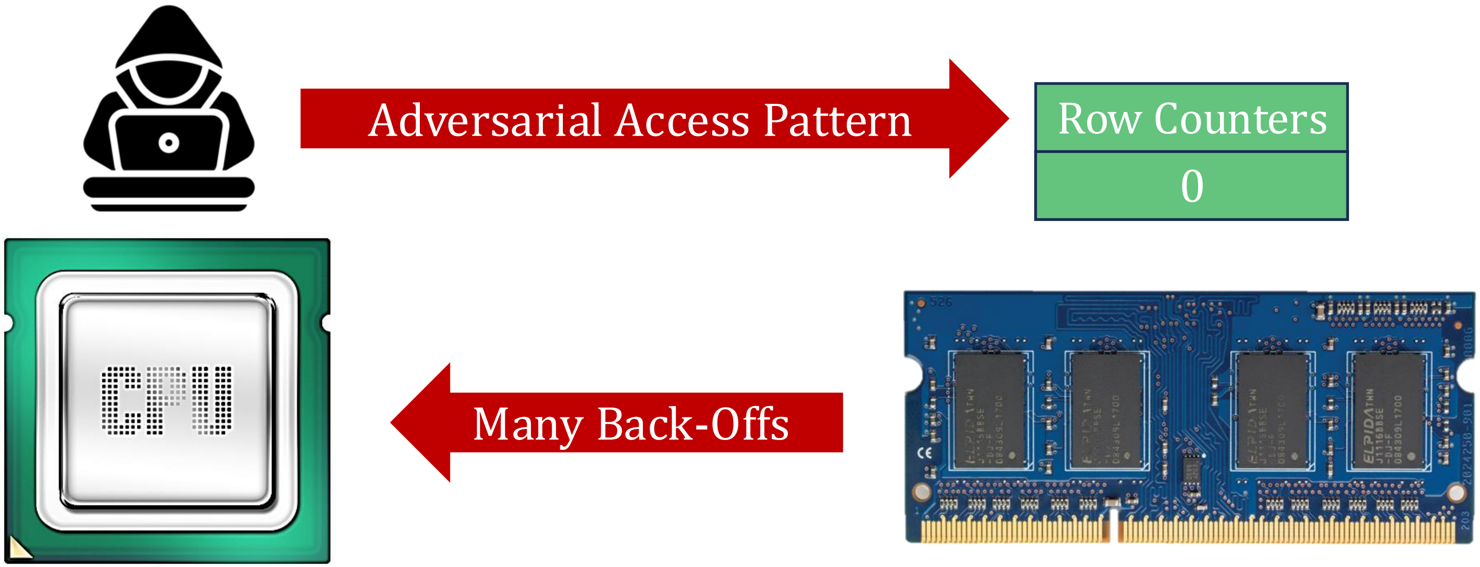

Mathematically hogs up to **79% of DRAM throughput**
of future DRAM chips

Degrades system performance by up to **65%** (53% on average)

More in the Paper

Understanding the Security Benefits and Overheads of Emerging Industry Solutions to DRAM Read Disturbance

Oğuzhan Canpolat[§†] A. Giray Yağlıkçı[§] Geraldo F. Oliveira[§] Ataberk Olgun[§]
Oğuz Ergin[†] Onur Mutlu[§]
[§]*ETH Zürich* [†]*TOBB University of Economics and Technology*

We present the first rigorous security, performance, energy, and cost analyses of the state-of-the-art on-DRAM-die read disturbance mitigation method, widely known as Per Row Activation Counting (PRAC), with respect to its description in the updated (as of April 2024) JEDEC DDR5 specification. Unlike prior state-of-the-art that advises the memory controller to periodically issue a DRAM command called refresh management (RFM), which provides the DRAM chip with time to perform its countermeasures, PRAC introduces a new back-off signal. PRAC's back-off signal propagates from the DRAM chip to the

data integrity of other physically close but *unaccessed* DRAM rows. RowHammer [1] is a prime example of DRAM read disturbance, where a DRAM row (i.e., victim row) can experience bitflips when at least one nearby DRAM row (i.e., aggressor row) is repeatedly activated (i.e., hammered) [1, 3–69] more times than a threshold, called the *minimum hammer count to induce the first bitflip* (N_{RH}). *RowPress* [70] is another prime example of DRAM read disturbance that amplifies the effect of RowHammer and consequently reduces N_{RH}.

Open Sourced

https://github.com/CMU-SAFARI/ramulator2

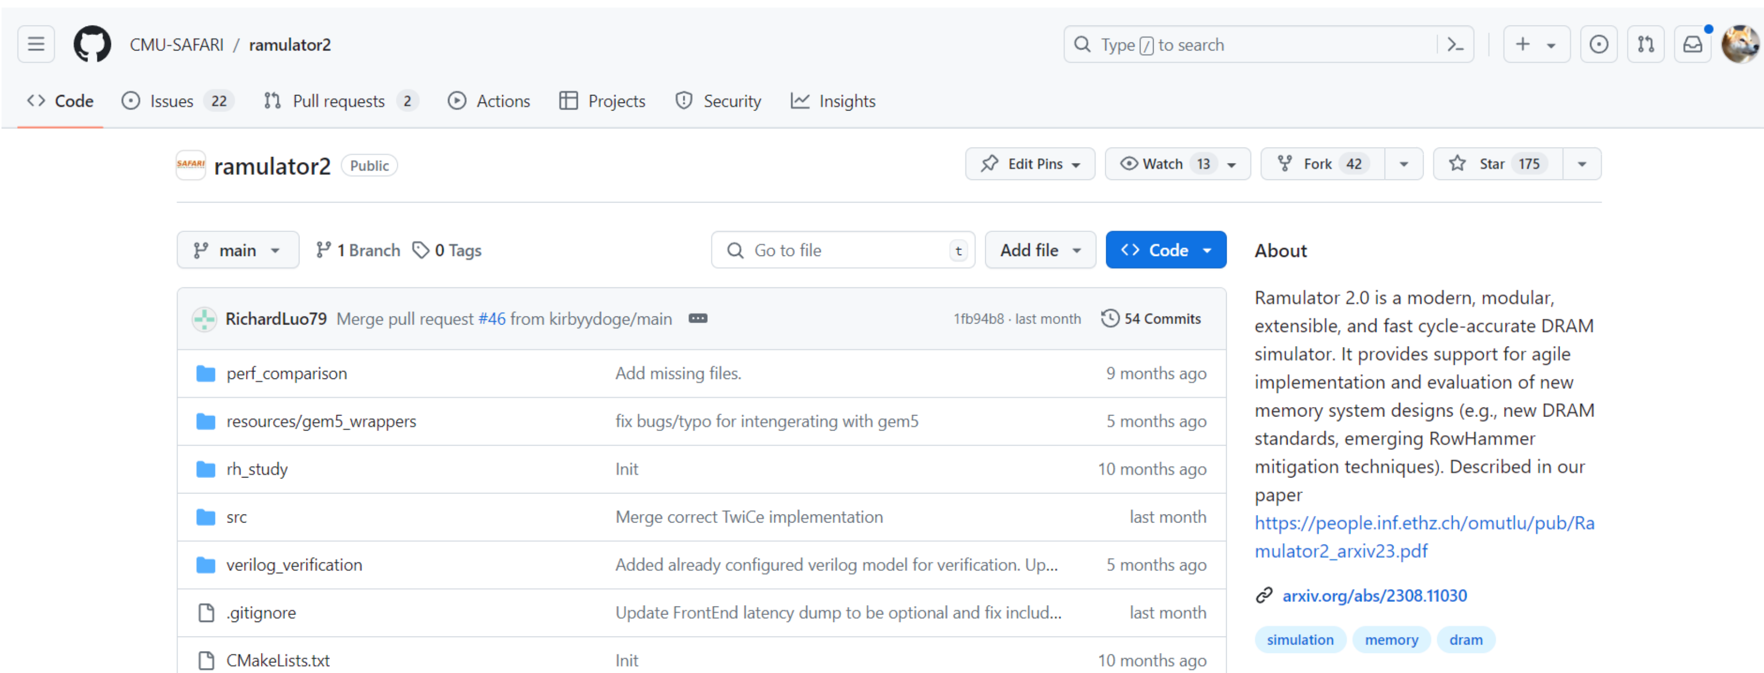

An Upcoming Paper: Self-Managing DRAM

Self-Managing DRAM: A Low-Cost Framework for Enabling Autonomous and Efficient in-DRAM Operations

Hasan Hassan[†] Ataberk Olgun[†] A. Giray Yağlıkçı Haocong Luo Onur Mutlu
ETH Zürich

available on arXiv and will appear in MICRO 2024

Enables new **in-DRAM** maintenance mechanisms without modifications other than the **ACT-NACK** signal from the DRAM chip to the memory controller

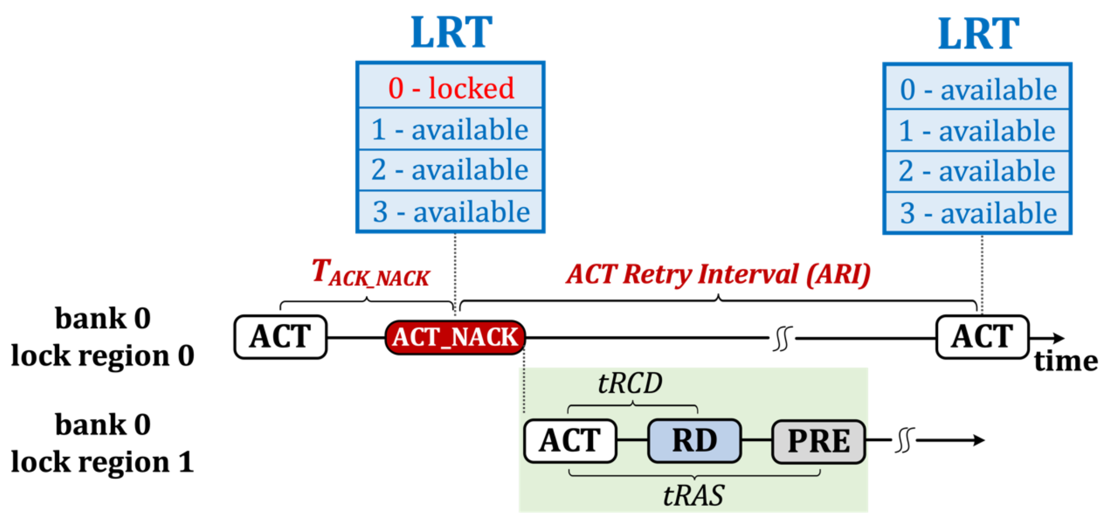

Three example DRAM maintenance operations:

- ❖ Periodic refresh
- ❖ RowHammer-preventive refresh
- ❖ Memory scrubbing

Hasan Hassan, Ataberk Olgun, A. Giray Yaglikci, Haocong Luo, and Onur Mutlu, ``Self-Managing DRAM: A Low-Cost Framework for Enabling Autonomous and Efficient in-DRAM Operations,'' cs.AR:2207.13358, 2022 (last update in 2024)

Enabling Efficient and Scalable DRAM Read Disturbance Mitigation via New Experimental Insights into Modern DRAM Chips

Abdullah Giray Yağlıkçı

Aug 14, 2024

Rambus

SAFARI

ETH zürich

Enabling Efficient and Scalable DRAM Read Disturbance Mitigation via New Experimental Insights into Modern DRAM Chips

Backup Slides

Abdullah Giray Yağlıkçı

Aug 14, 2024

Rambus

SAFARI

ETH zürich

Understanding PRAC

Understanding the Security Benefits and Overheads of Emerging Industry Solutions to DRAM Read Disturbance

Oğuzhan Canpolat[§†] A. Giray Yağlıkçı[§] Geraldo F. Oliveira[§] Ataberk Olgun[§]
Oğuz Ergin[†] Onur Mutlu[§]
[§]*ETH Zürich* [†]*TOBB University of Economics and Technology*

presented in DRAMSec 2024

- PRAC is the latest update to DDR5 in April 2024
- We show that
  - PRAC is secure and has non-negligible performance, energy, and area overheads for benign workloads
  - PRAC can be exploited by memory performance attacks (up to **79% reduction** in DRAM chip's throughput)

Industry Solutions to Read Disturbance: When To Refresh? (I)

Preventive refresh is a **blocking** operation

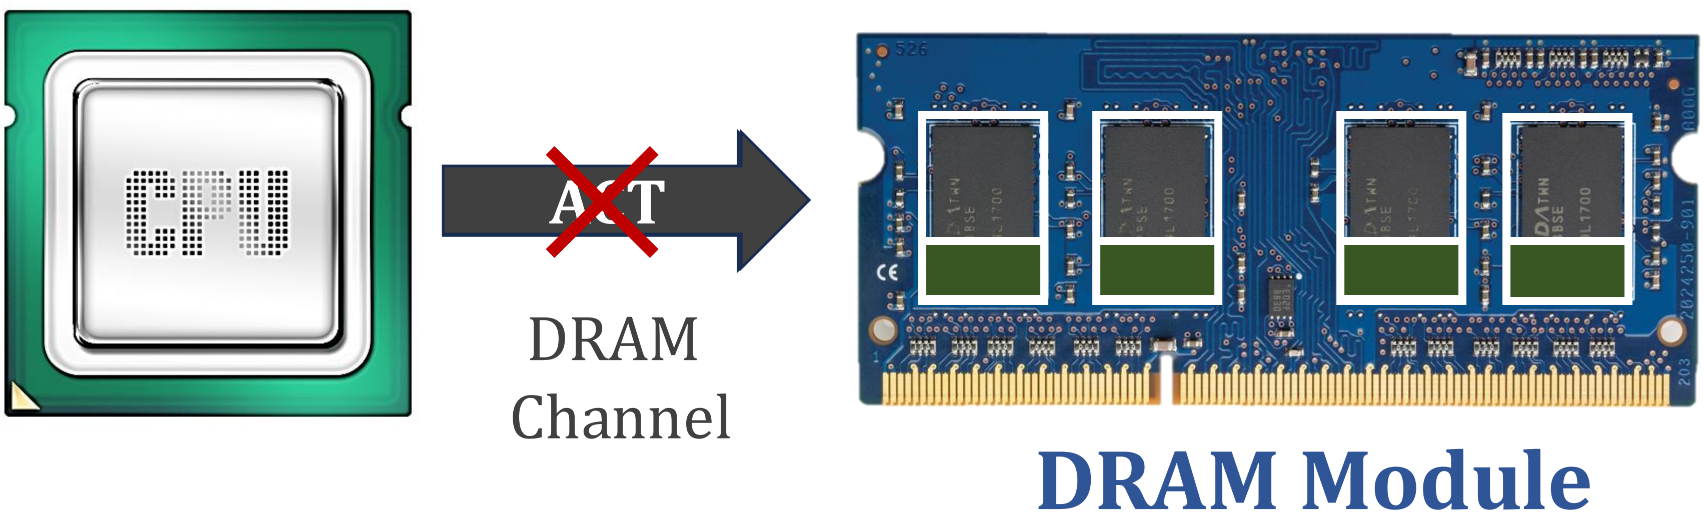

Memory controller **cannot activate** a row while **a preventive refresh blocks the bank**

Industry Solutions to Read Disturbance: When To Refresh? (II)

Earlier JEDEC DDR5 specifications introduce **Refresh Management (RFM)** commands

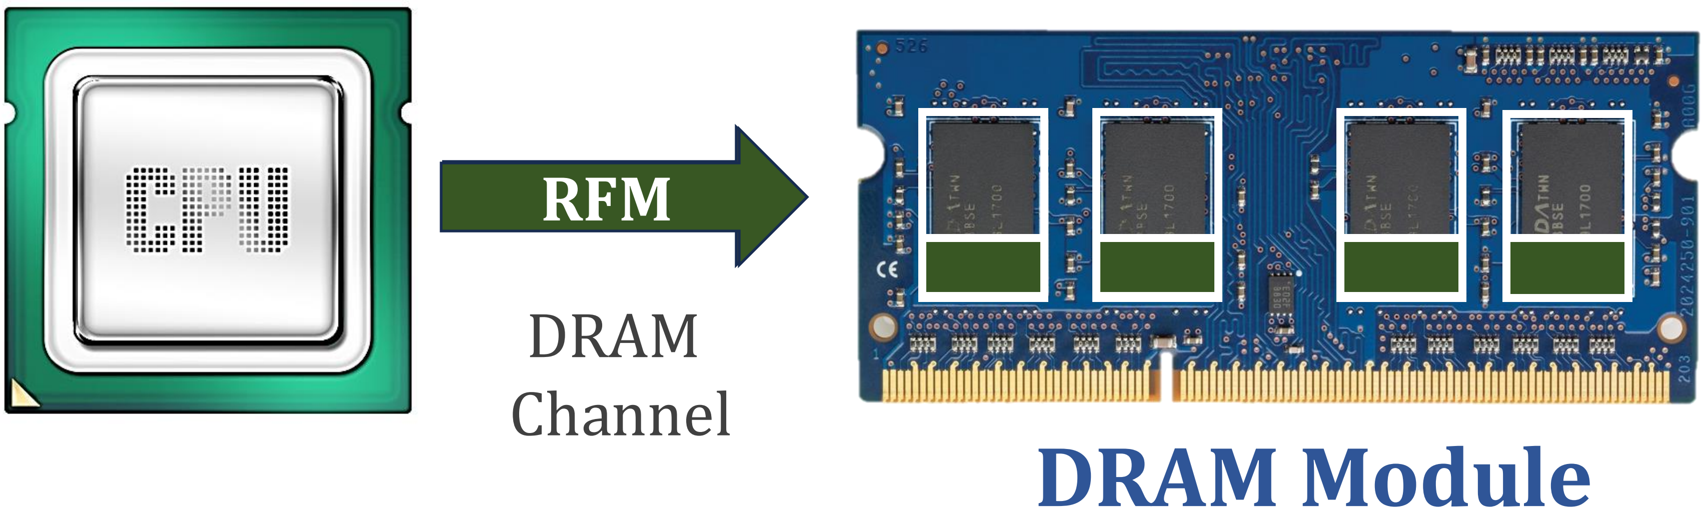

Memory controller sends an **RFM command** to allow time for preventive refreshes

Industry Solutions to Read Disturbance

Periodic Refresh Management (PRFM)

Memory controller **periodically** issues RFM commands

Per Row Activation Counting and Back-Off (PRAC)

DRAM chip **tracks** row activations and **requests** RFMs by sending **back-off** signals

Industry Solutions to Read Disturbance: Periodic Refresh Management (PRFM)

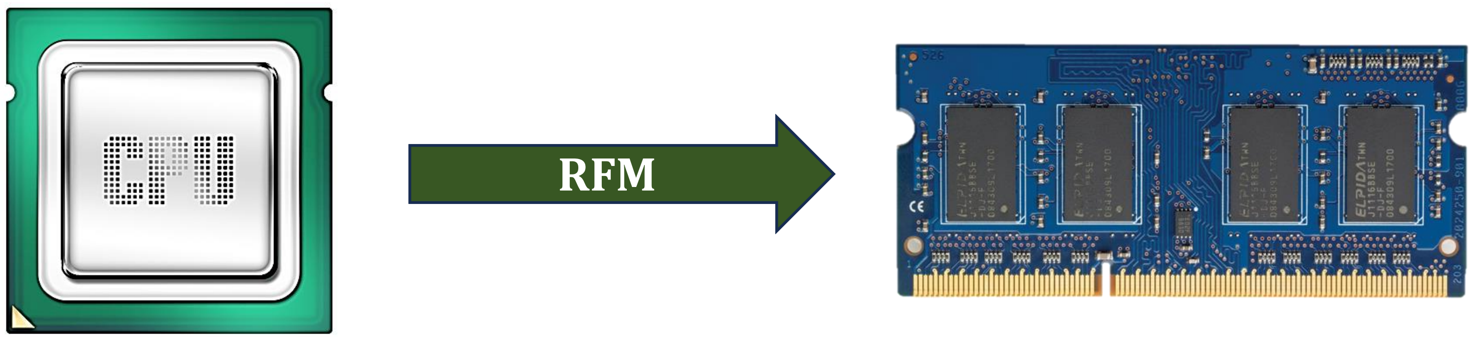

Bank Activation Counters			
0	0	0	0

RFM Threshold

PRFM tracks activations with **low accuracy**,
causing **high number** of preventive refreshes,
leading to **large** performance and energy overheads

ACT → **ACT** → **ACT** → **RFM** →

DRAM Commands

Industry Solutions to Read Disturbance

Periodic Refresh Management (PRFM)

Memory controller **periodically** issues RFM commands

Per Row Activation Counting and Back-Off (PRAC)

DRAM chip **tracks** row activations and **requests** RFMs by sending **back-off**

Industry Solutions to Read Disturbance: Per Row Activation Counting

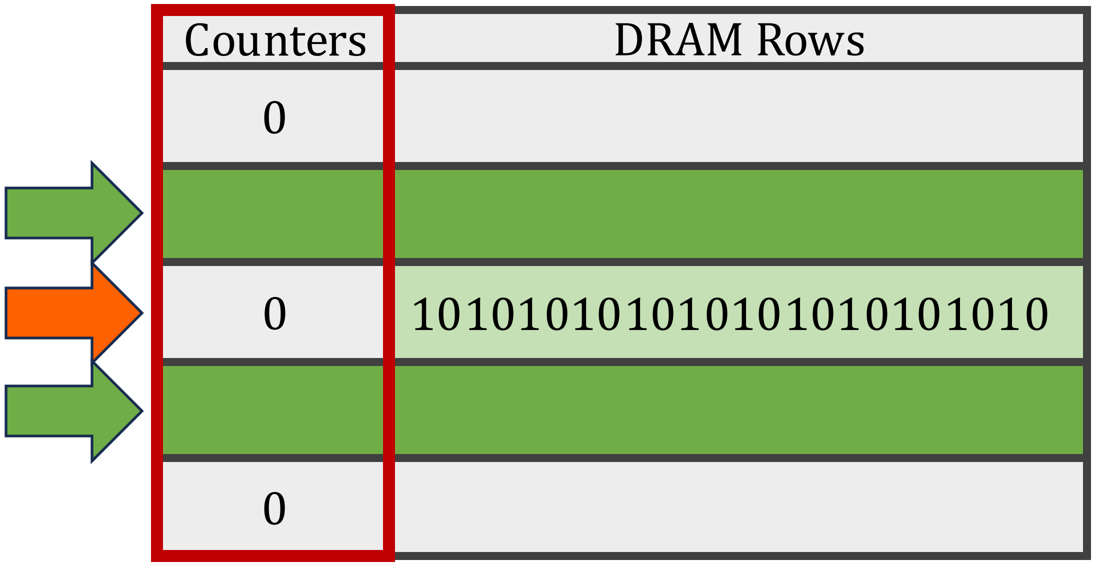

PRAC allows accurate tracking of aggressor row activations

Industry Solutions to Read Disturbance: Per Row Activation Counting (PRAC)

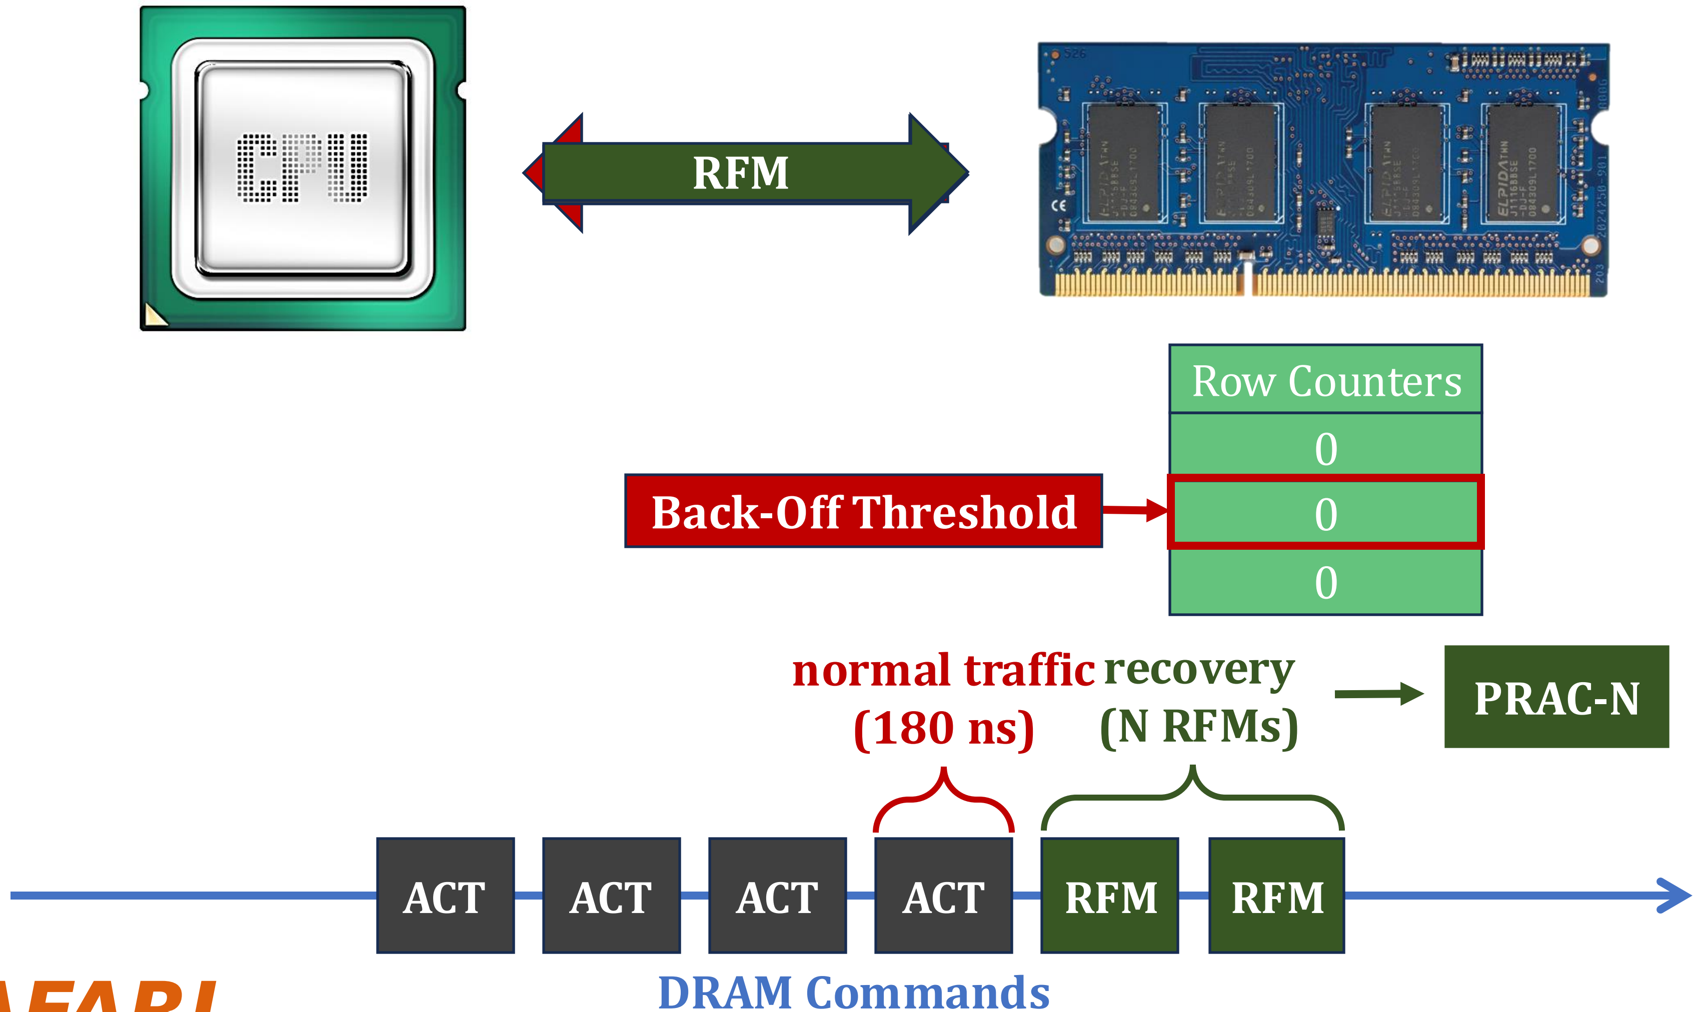

Performance Comparison: Industry Solutions

1 **PRFM**

Memory controller **periodically** issues RFM

2 **PRAC-N**

Memory controller issues **N** RFMs each with **back-off**

3 **PRAC+PRFM**

Memory controller issues RFM **periodically** and with **back-offs**

4 **PRAC-Optimistic**

PRAC-4 with **no** change in DRAM timing parameters

Experimental Results: Performance Overhead and Its Scaling

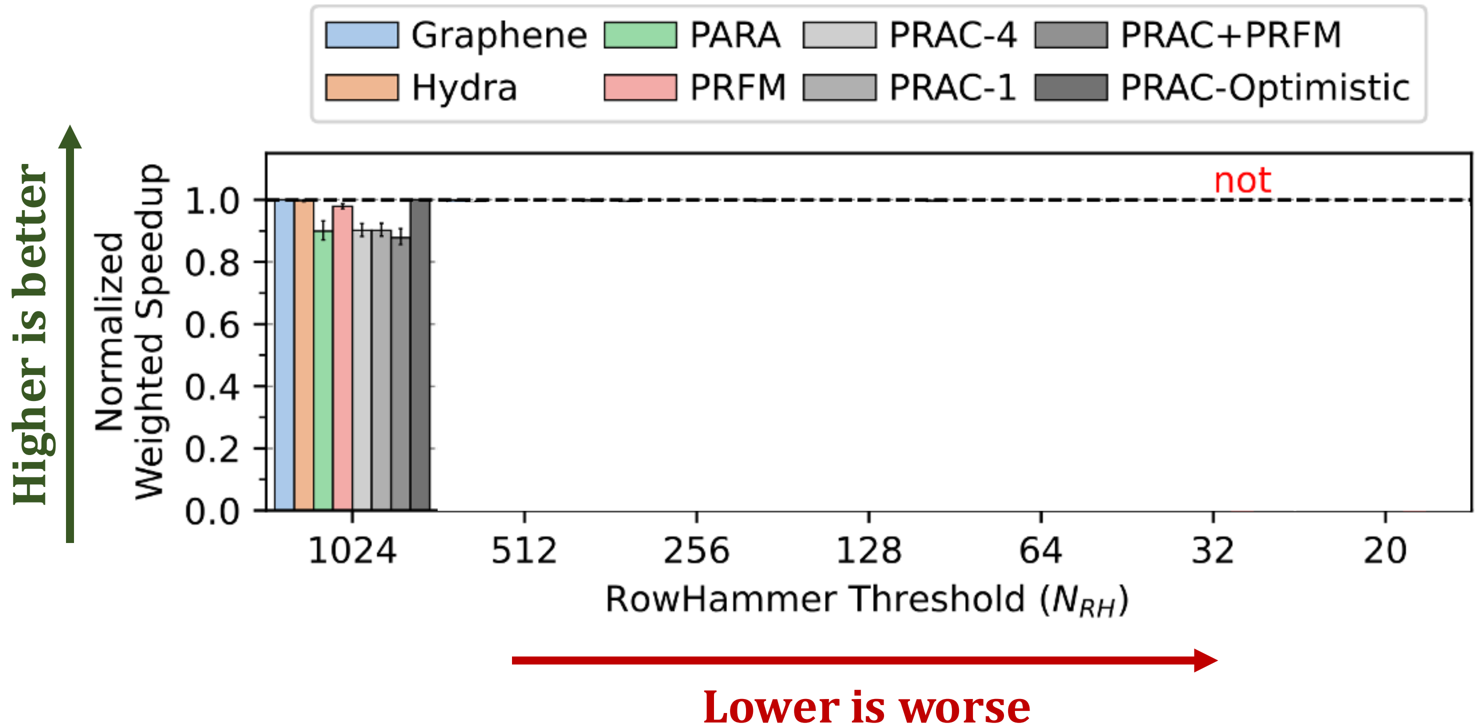

Experimental Results: Performance Overhead and Its Scaling

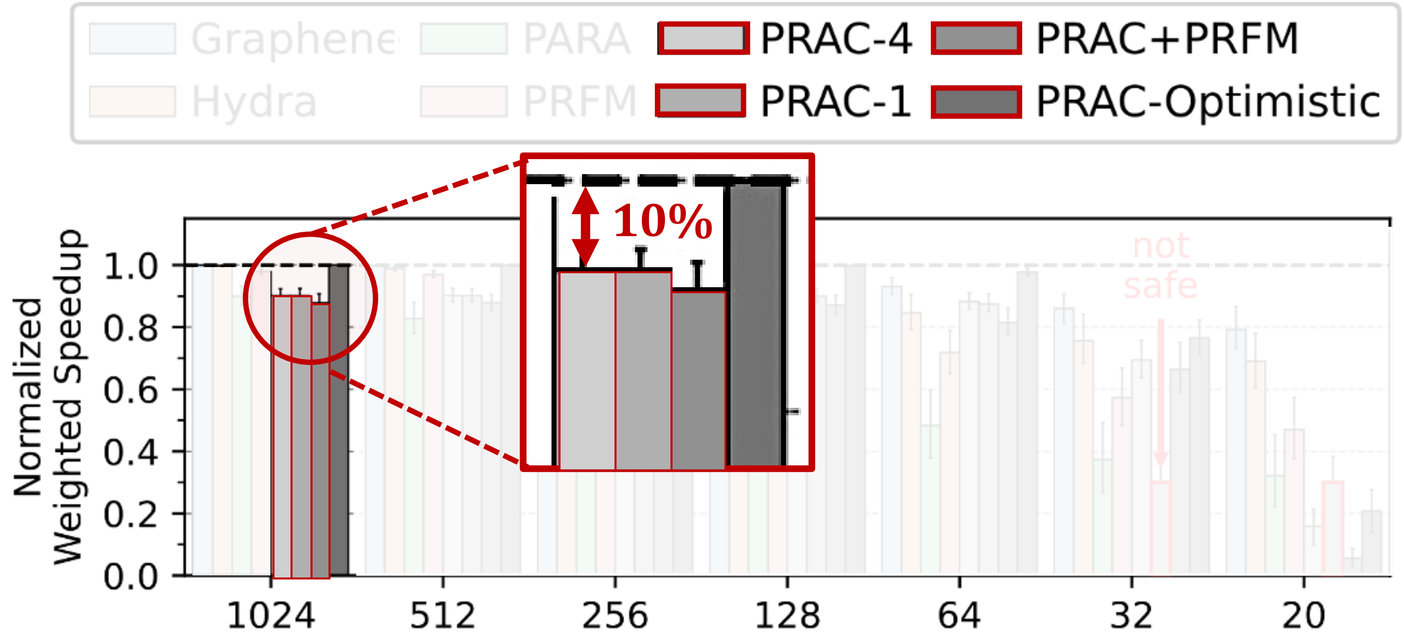

PRAC has **non-negligible** performance overhead (**10%**) due to **increased** access latency

Experimental Results: Performance Overhead and Its Scaling

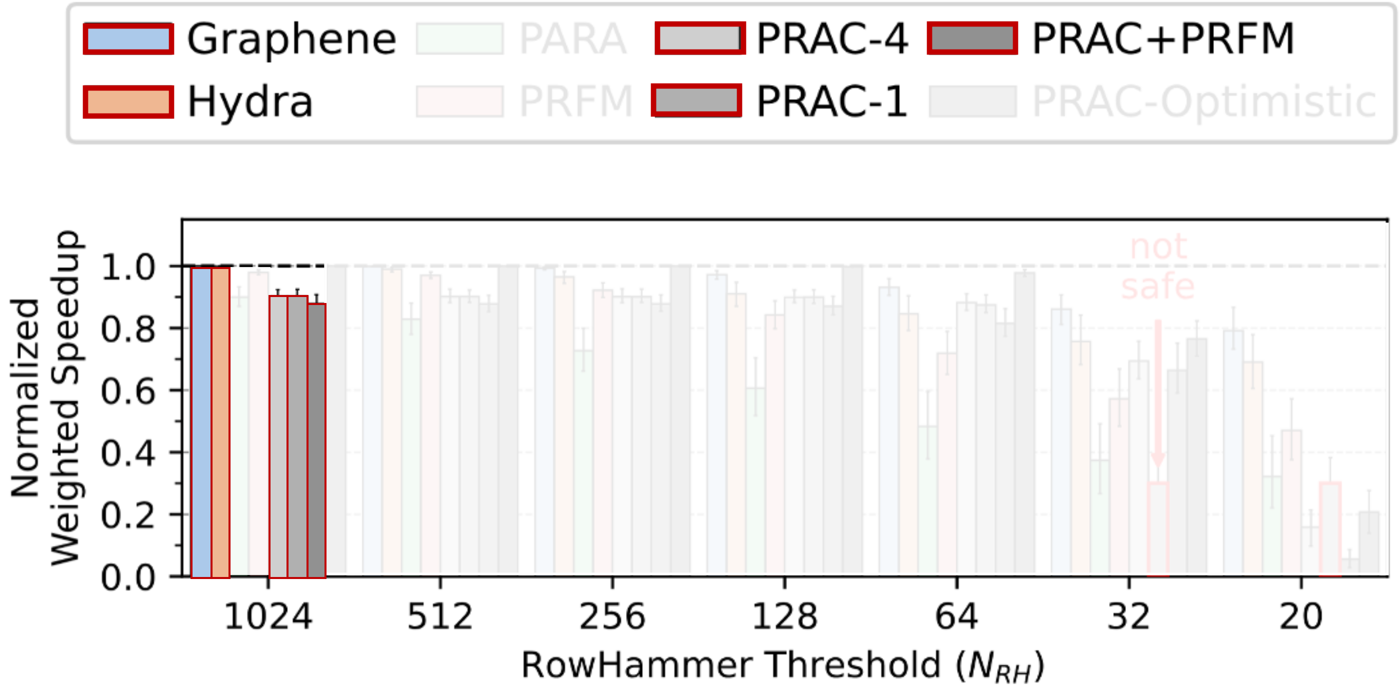

Graphene and **Hydra** outperform **PRAC**
at relatively high N_{RH} values

Experimental Results: Performance Overhead and Its Scaling

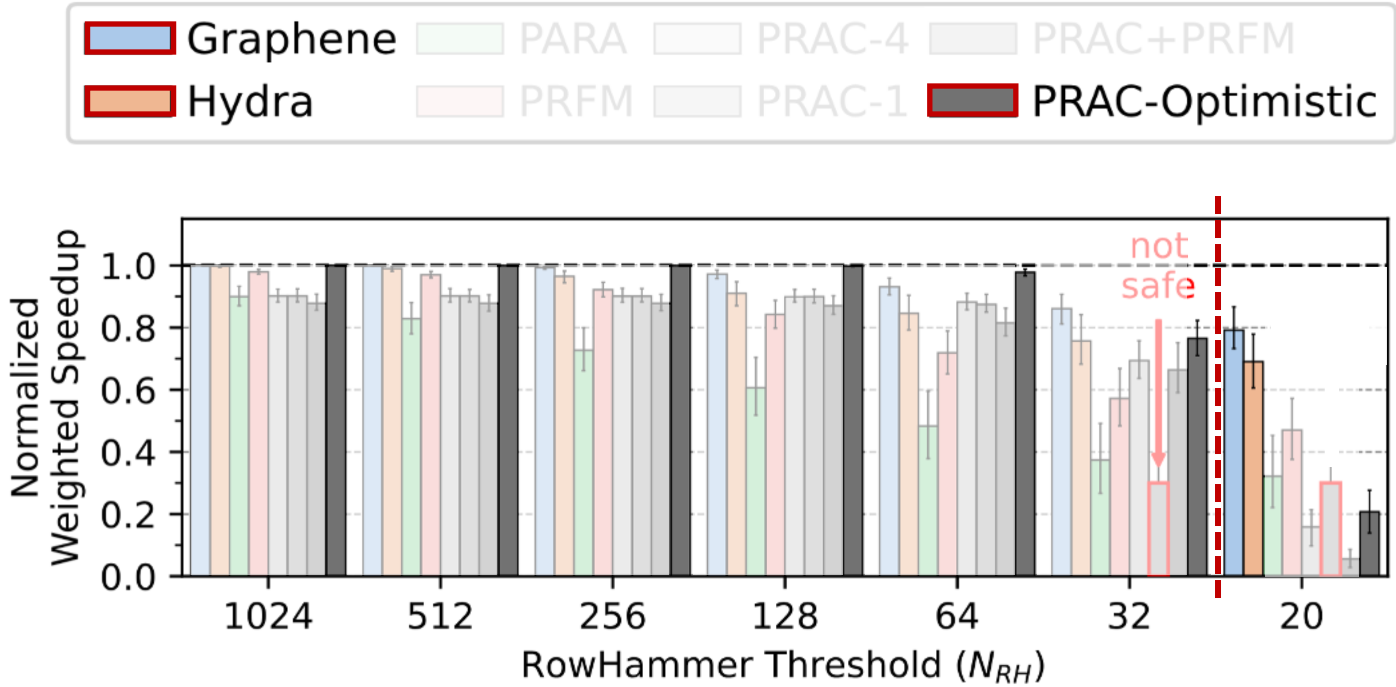

PRAC-Optimistic **outperforms** all evaluated mitigation mechanisms (above N_{RH} of 32)

Experimental Results: Performance Overhead and Its Scaling

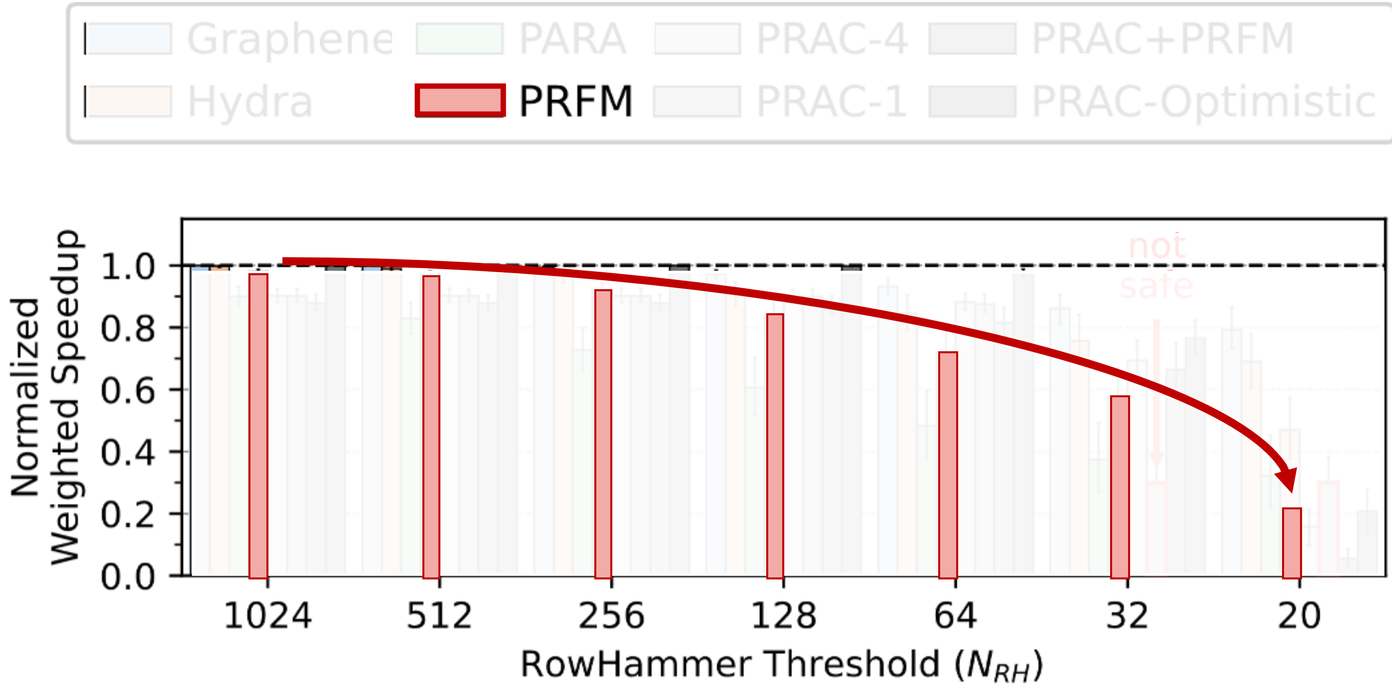

PRFM's system performance overheads **significantly increase (by 37x)** as **N_{RH} decreases**

Experimental Results: DRAM Energy Overhead and Its Scaling

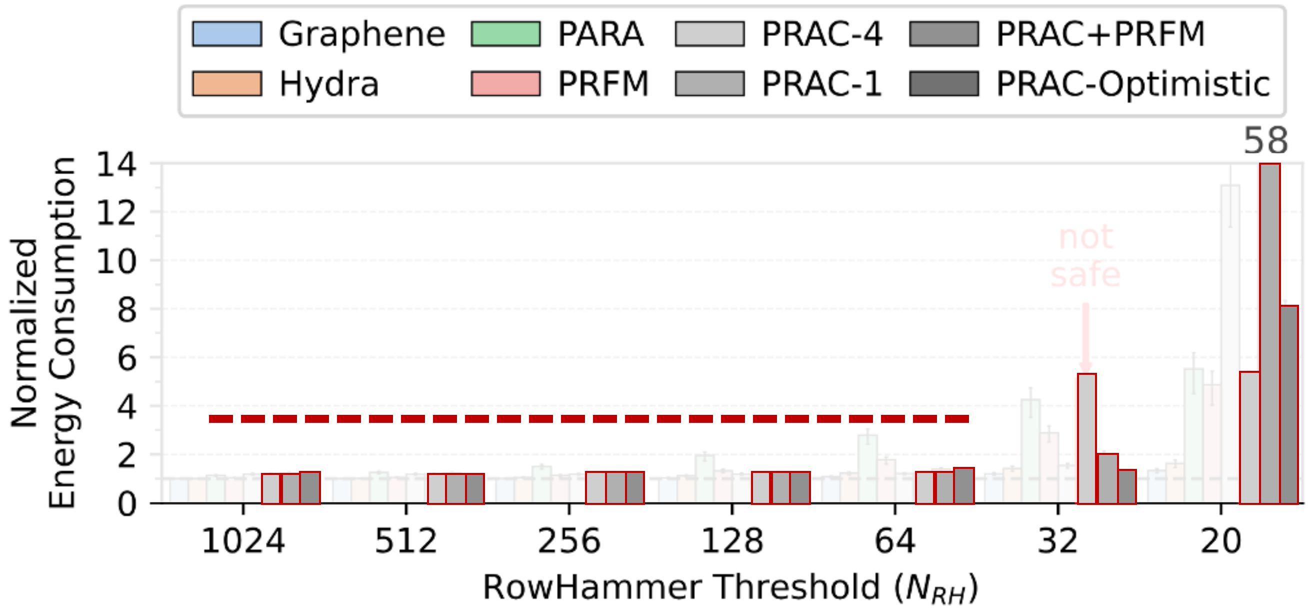

Above N_{RH} of 32, **PRAC** overhead only **slightly** increases due to **timely** preventive refreshes

Below N_{RH} of 32, **PRAC** overhead **significantly** increases due to conservative thresholds against a **wave attack**

How Large is 1000 Activations?

- Bitflips occur **at ~1000 activations**
- Mitigation mechanisms trigger **preventive actions** (e.g., preventive refresh) **at ~500 activations**
- Is 500 **a distinctive activation count**?
- Benign workloads activate **hundreds of rows** more than **500 times** in a refresh window

Memory intensive workloads
from SPEC'06, SPEC'17, TPC, YCSB, and MediaBench

Workload	MPKI	ACT-64+	ACT-128+	ACT-512+
429.mcf	68.27	2564	2564	2564
470.lbm	28.09	7089	6596	664
462.libquantum	25.95	1	0	0
549.fotonik3d	25.28	10065	88	0
459.GemsFDTD	24.93	10572	218	0
519.lbm	24.37	5824	5455	2482
434.zeusmp	22.24	11085	4825	292
510.parest	17.79	803	185	94
433.milc	17.22	321	92	0
437.leslie3d	15.82	4678	631	7
483.xalancbmk	13.67	4354	776	113
482.sphinx3	12.59	1385	762	304
505.mcf	11.35	1582	1384	732
471.omnetpp	10.72	1015	419	122
tpch2	9.09	875	307	88
520.omnetpp	9.00	1185	84	32
tpch17	7.43	1196	158	26
473.astar	5.18	5957	22	0
436.cactusADM	4.94	6151	2354	1134
jp2_encode	4.18	0	0	0

Benign workloads **might not be so benign**

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**
  Preliminary data on aging via 68-day of continuous hammering

Aging can lead to read disturbance bitflips at **smaller** hammer counts

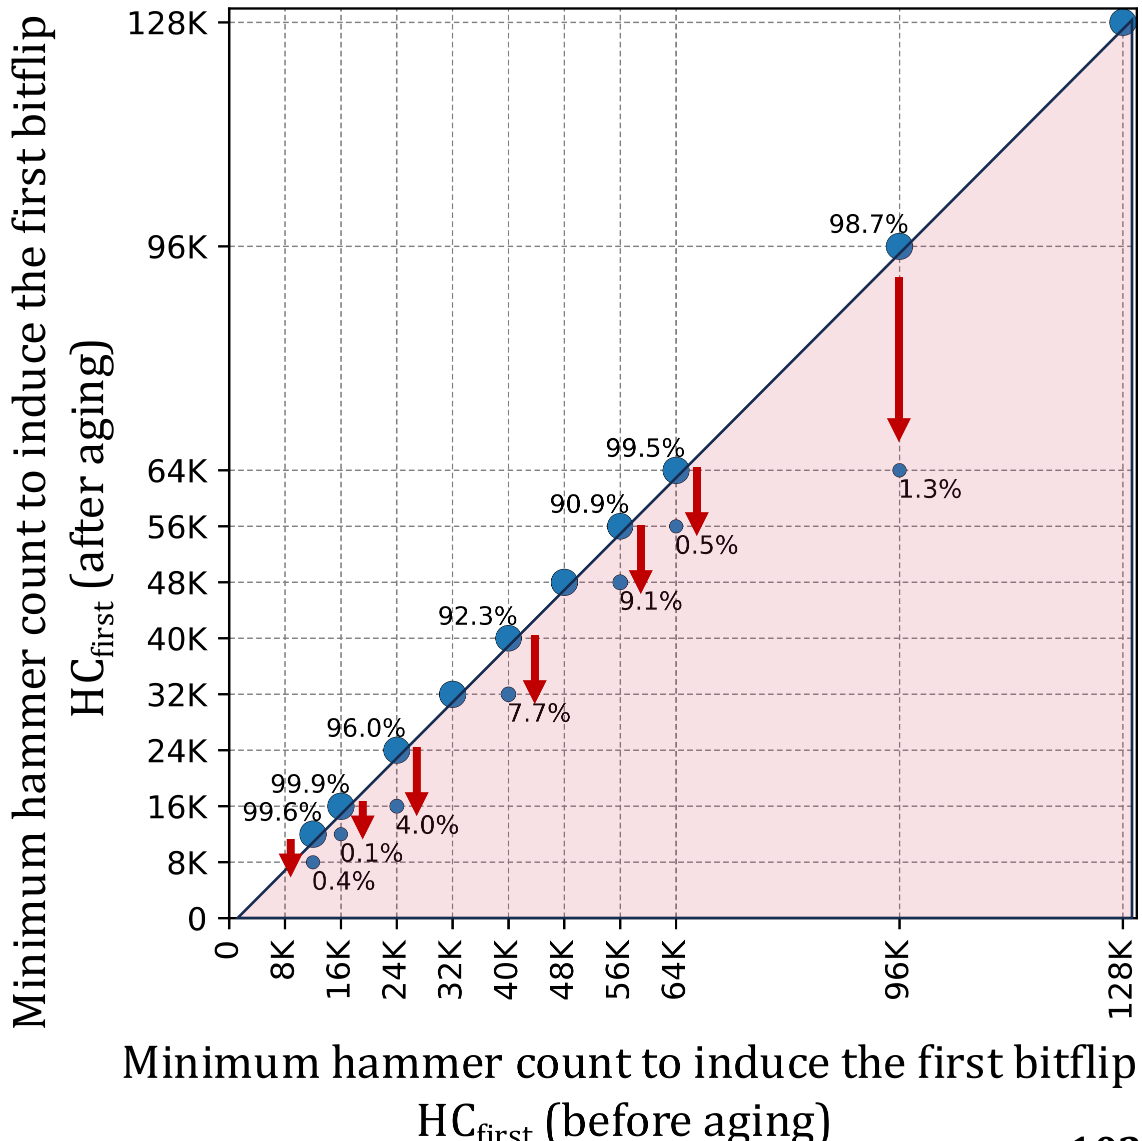

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**
  Preliminary data on aging via 68-day of continuous hammering

Aging can lead to read disturbance bitflips at **smaller** hammer counts

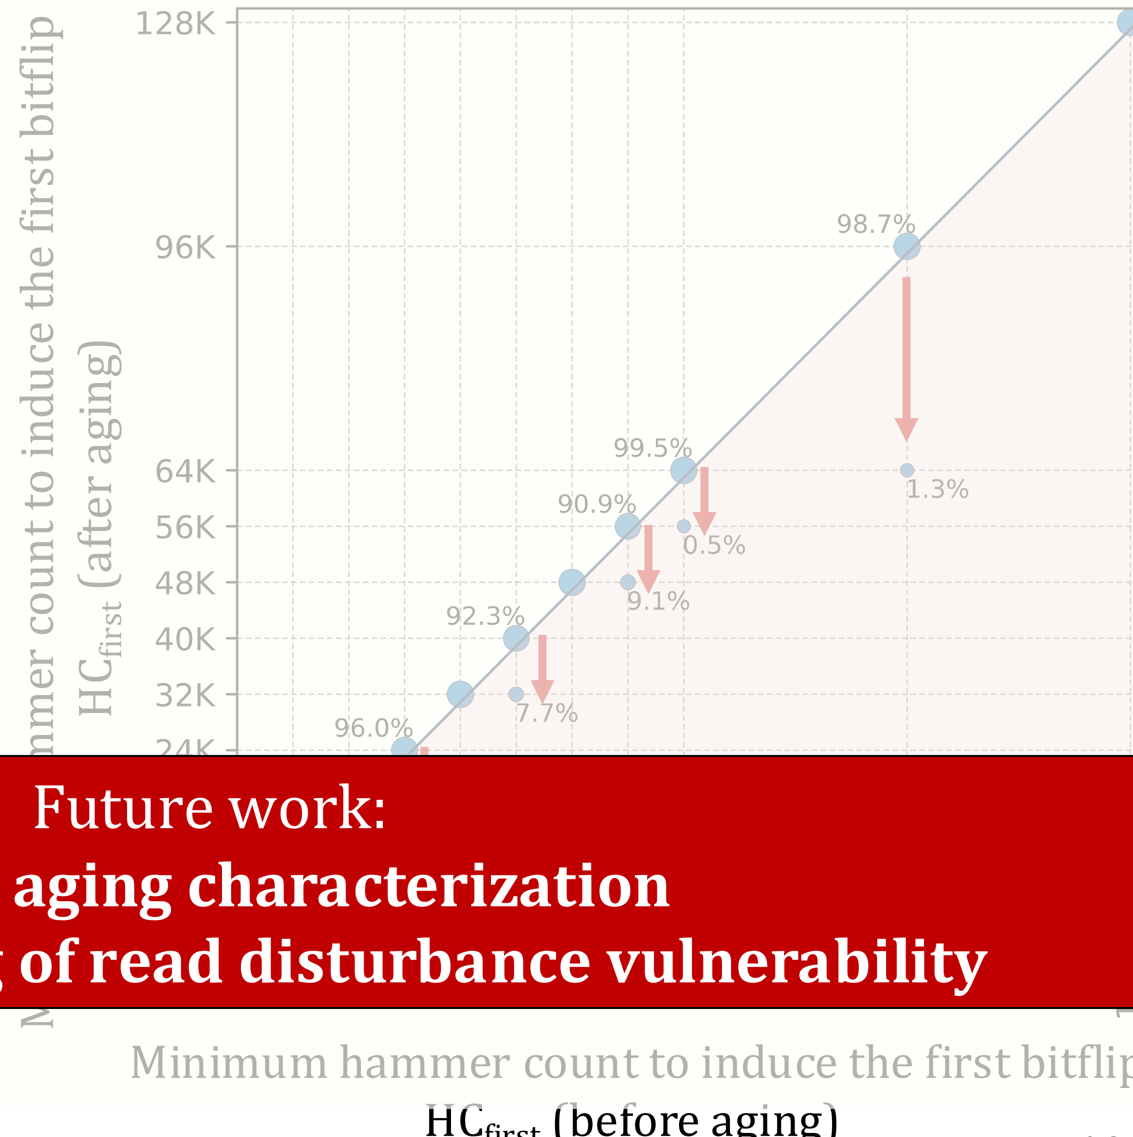

Future work:
rigorous aging characterization
and **online profiling of read disturbance vulnerability**

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**

- **Interactions** across different error mechanisms
  - RowHammer
  - RowPress
  - Data retention time errors
  - Variable retention time

  …

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**

- **Interactions** across different error mechanisms

- What is **the worst-case**?
  - Temperature
  - Data pattern
  - Memory access pattern
  - Spatial variation
  - Voltage

What is **the worst-case** considering all **these sensitivities**?

What is **the minimum hammer count** to induce a read disturbance bitflip?

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**

- **Interactions** across different error mechanisms

- What is **the worst-case**?

How reliable are our DRAM chips?

How reliable will our DRAM chips be tomorrow?

We **do not** know! This is an **open research problem**

Future Research for Better Memory Systems

- Deeper Understanding of Physics and Vulnerabilities
- Flexible and Intelligent Memory Chips, Interfaces, Controllers
- Cross-Layer Communication

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary

DRAM read disturbance solution is here

Configuration

DRAM chip is here

DRAM module can be replaced

Reconfiguration is needed

+ add more hardware counters
+ increase metadata cache capacity

Aging might change the chip's characteristics

Deployed solutions should be patchable in field

Future Research for Better Memory Systems

- Deeper Understanding of Physics and Vulnerabilities

- Flexible and Intelligent Memory Chips, Interfaces, Controllers

- Cross-Layer Communication

Cross-Layer Communication

		Detection	Mitigation
Software 	• Memory allocations • Memory access patterns • Control flow patterns • Time / power measurements	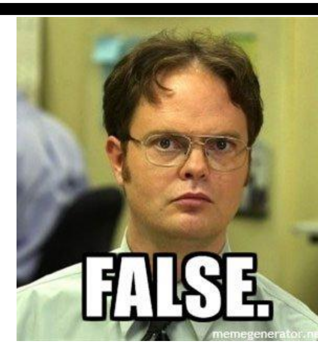	++
uArch 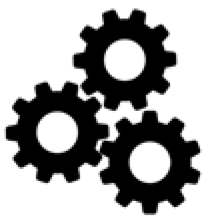	• Memory request scheduling • Speculative execution • Prefetching, branch prediction • Power management	+	+
Device 	• Bitflips occur • Memory isolation is broken	++	~

Cross-Layer Communication

		Detection	Mitigation
Software	• Memory allocations • Memory access patterns • Control flow patterns • Time / power measurements	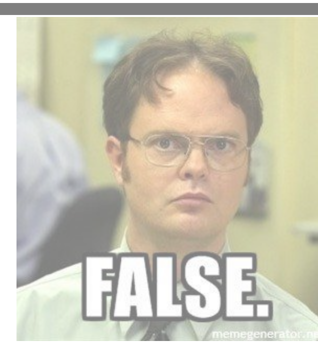	
	• Power management		
Device	• Bitflips occur • Memory isolation is broken		~

Cross-layer communication is crucial going forward

Attack Performance Overhead and Its Scaling

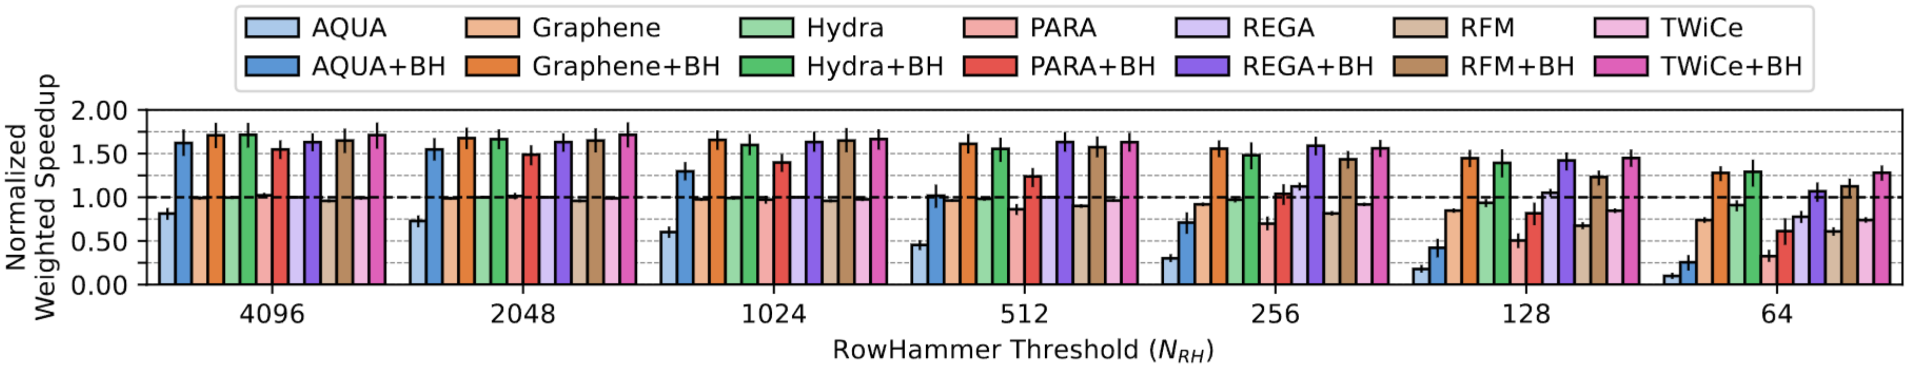

Mitigation mechanisms incur significantly increasing overheads as RowHammer threshold decreases

BreakHammer significantly (40% on average) increases system performance (even exceeding baseline with no mitigation)

Attack Energy Overhead and Its Scaling

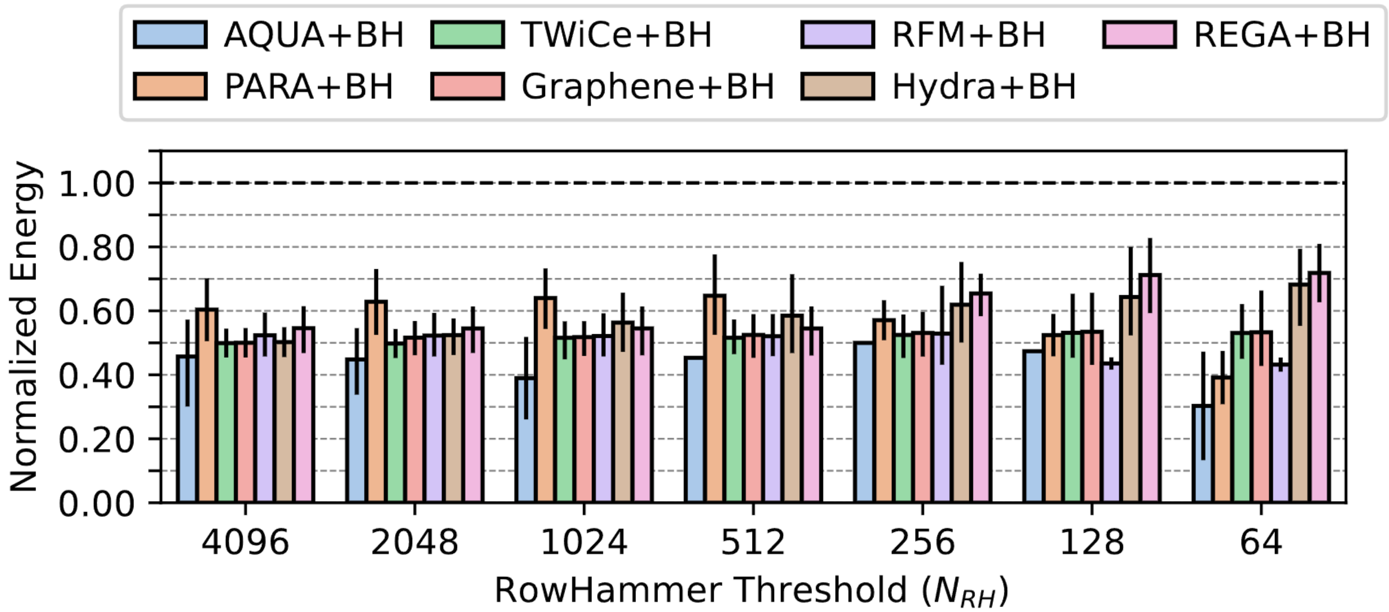

BreakHammer drastically improves (54% on average) the DRAM energy consumption

Number of Preventive Actions and Its Scaling

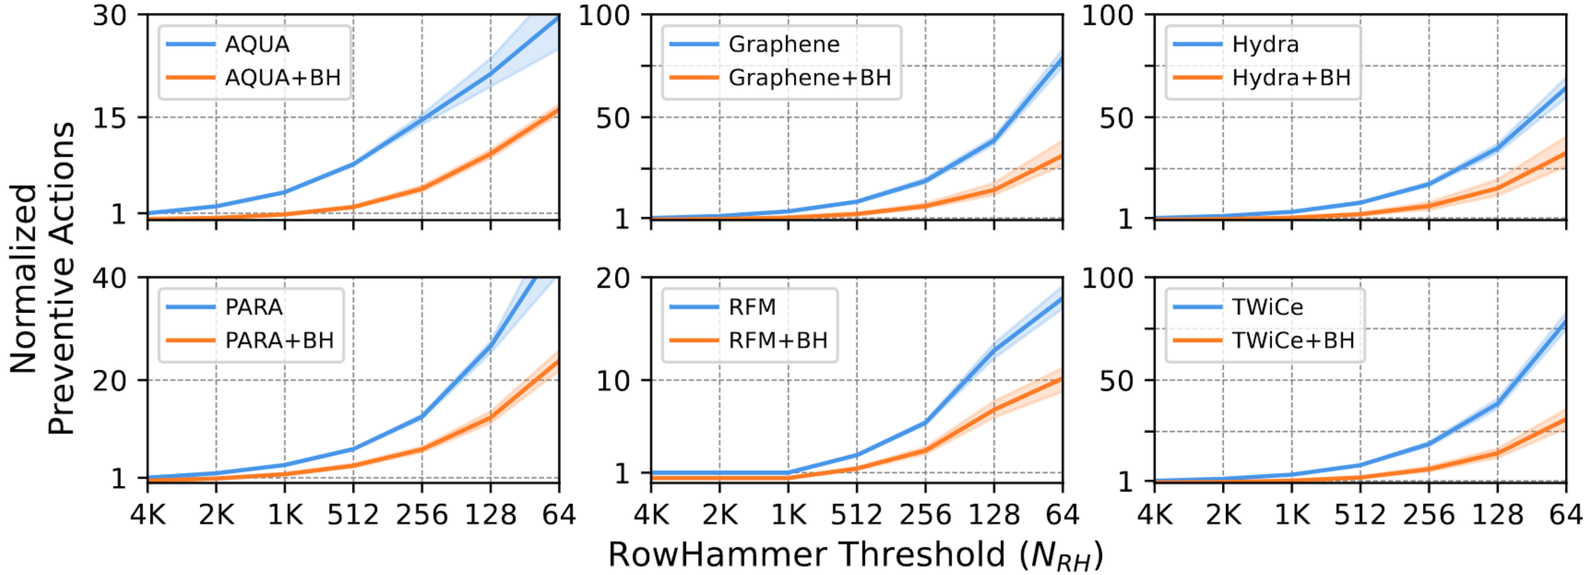

Mitigation mechanisms issue significantly **increasing** number of preventive actions as RowHammer threshold **decreases**

BreakHammer significantly (67% on average) reduces the number of preventive actions performed across all mitigations

Benign Performance and Its Scaling

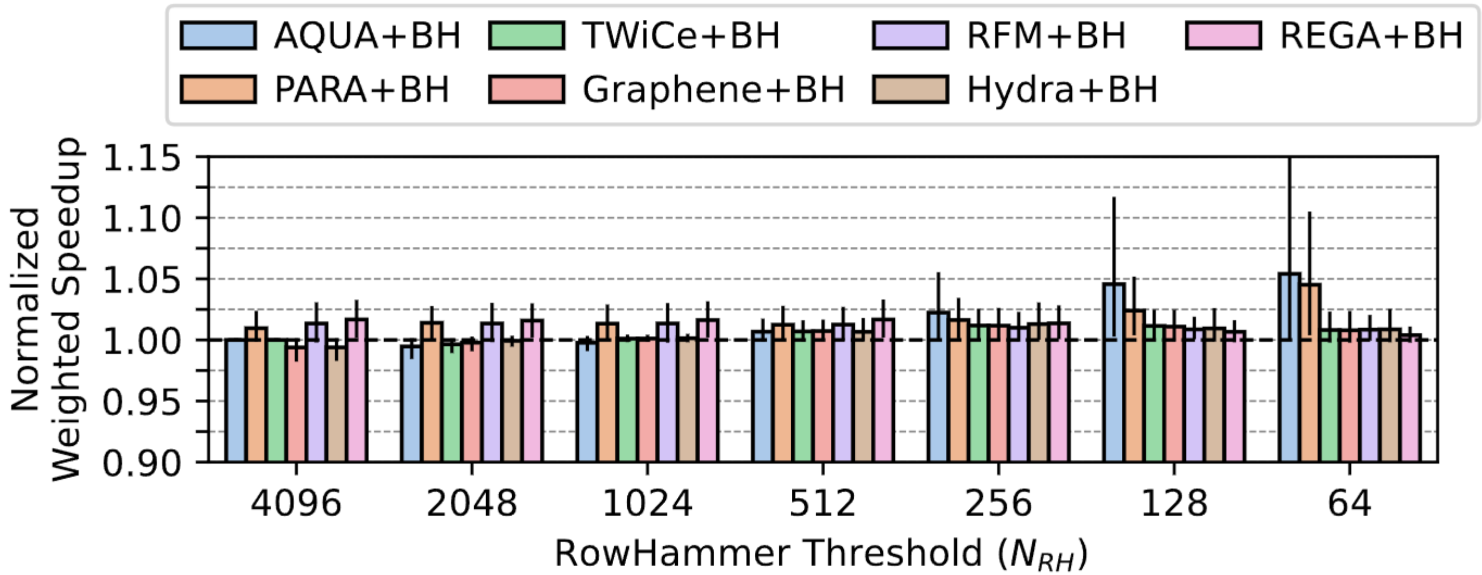

BreakHammer does **not** degrade
the average system performance of **benign** only workloads

Benign Memory Latency at $N_{RH} = 64$

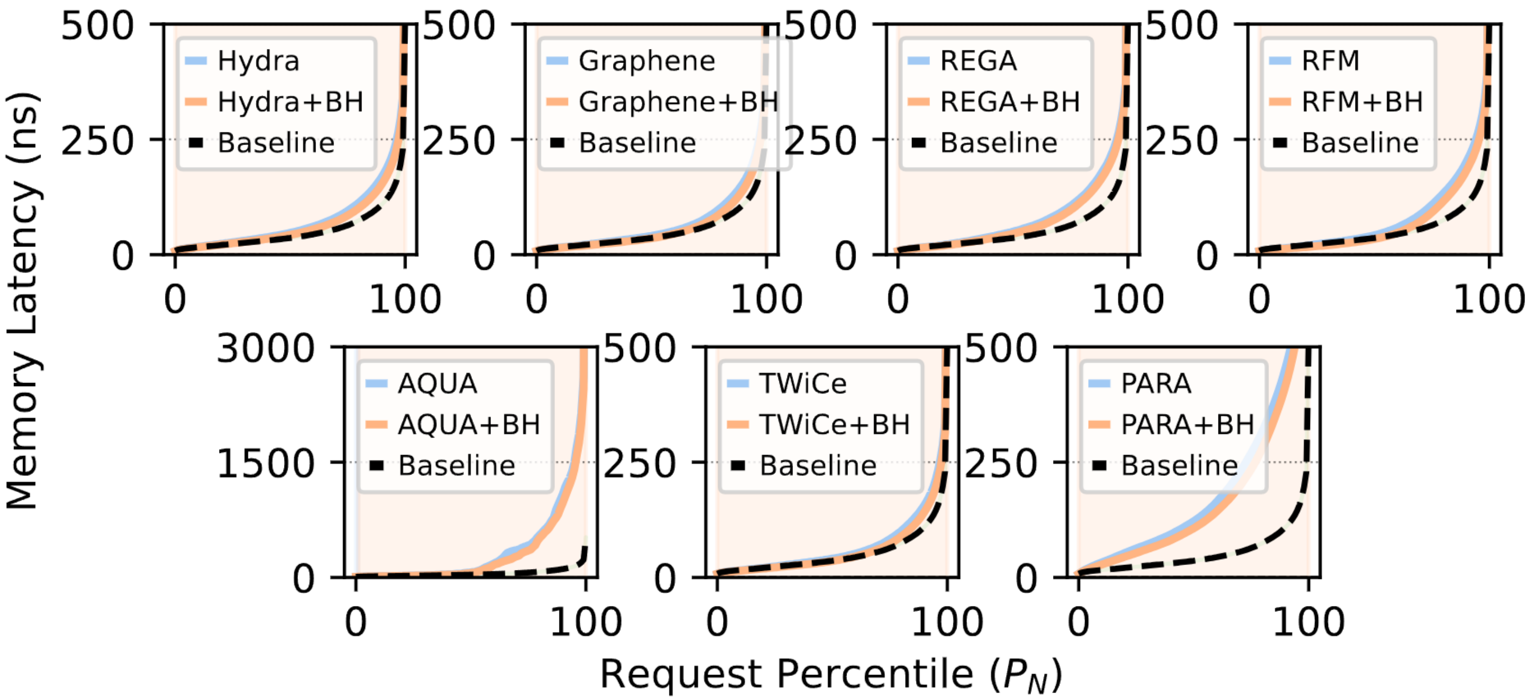

BreakHammer does **not** degrade
the memory latency of **benign** only workloads

Comparison to BlockHammer

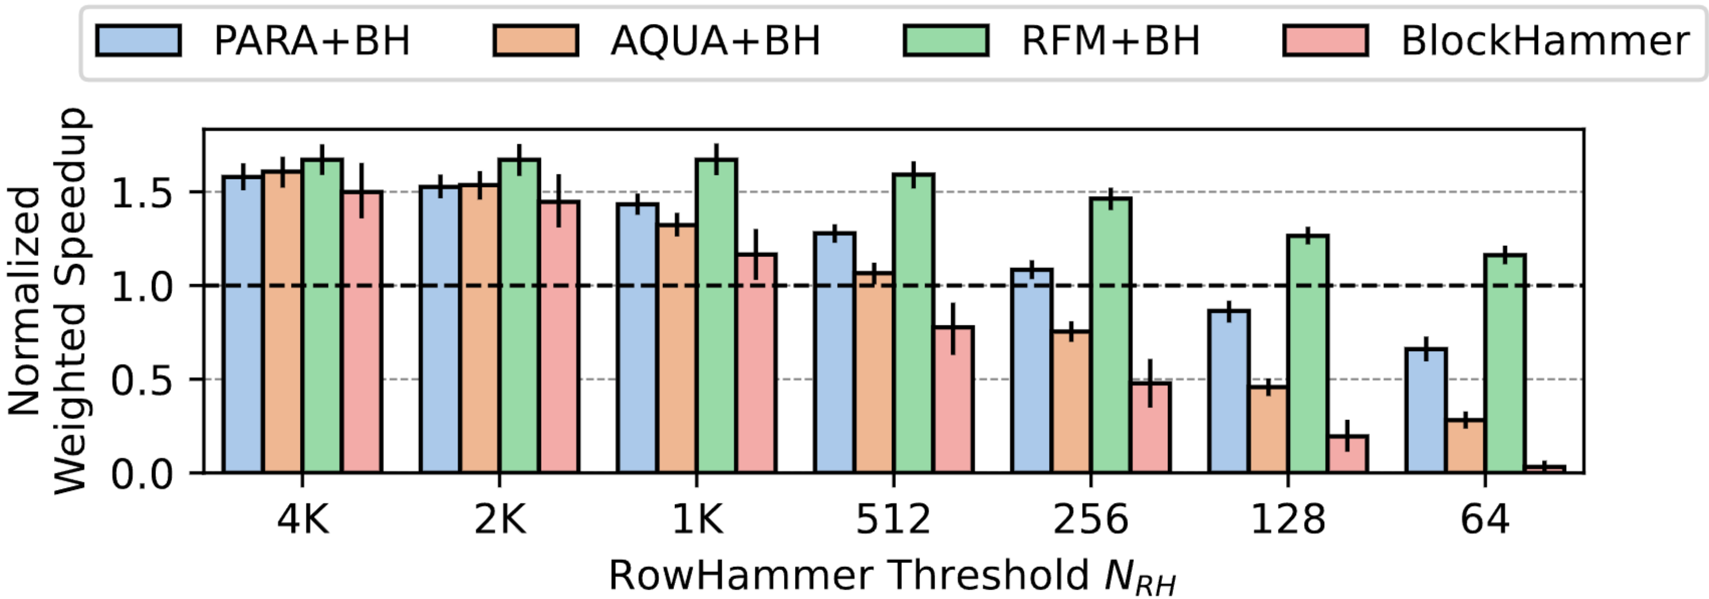

BreakHammer outperforms BlockHammer, the state-of-the-art throttling-based RowHammer mitigation mechanism

Observing Actions and Detecting Anomalies

BreakHammer **tracks the number of preventive actions** caused by each thread and **detects anomalies** in thread activity

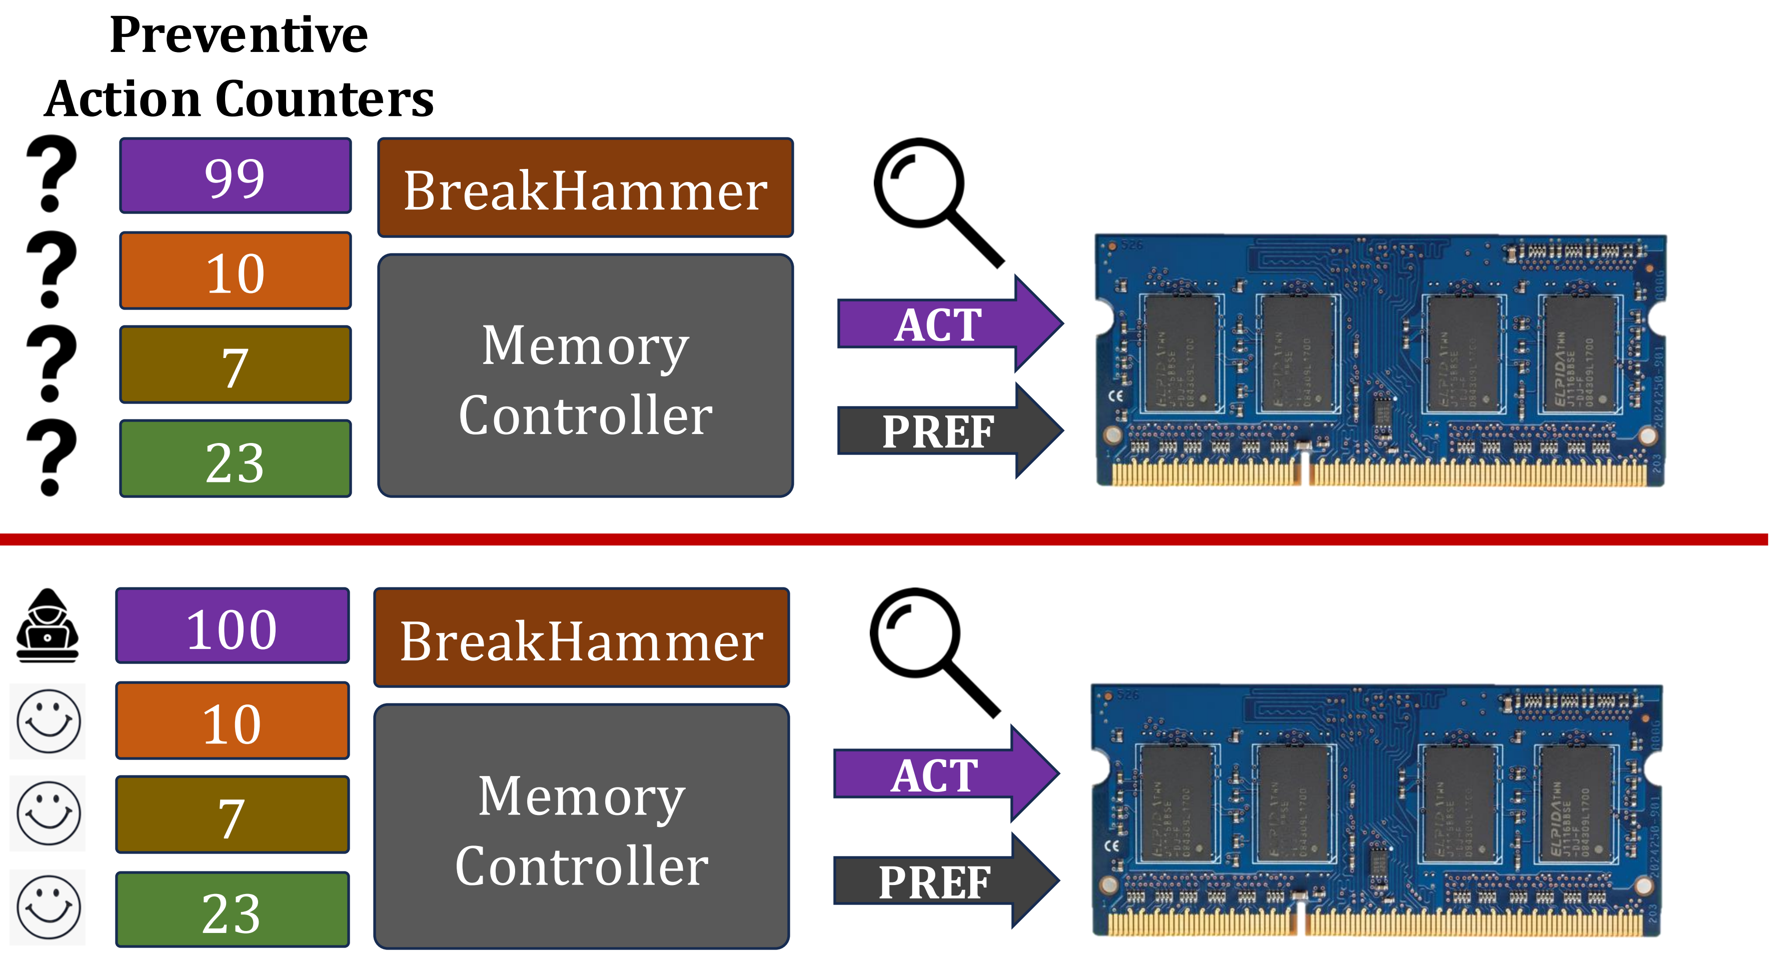

Filtering and Throttling Memory Usage

BreakHammer **slows down suspicious threads**
by reducing the number of miss status holding registers (MSHR)
they can allocate in the last level of the cache hierarchy

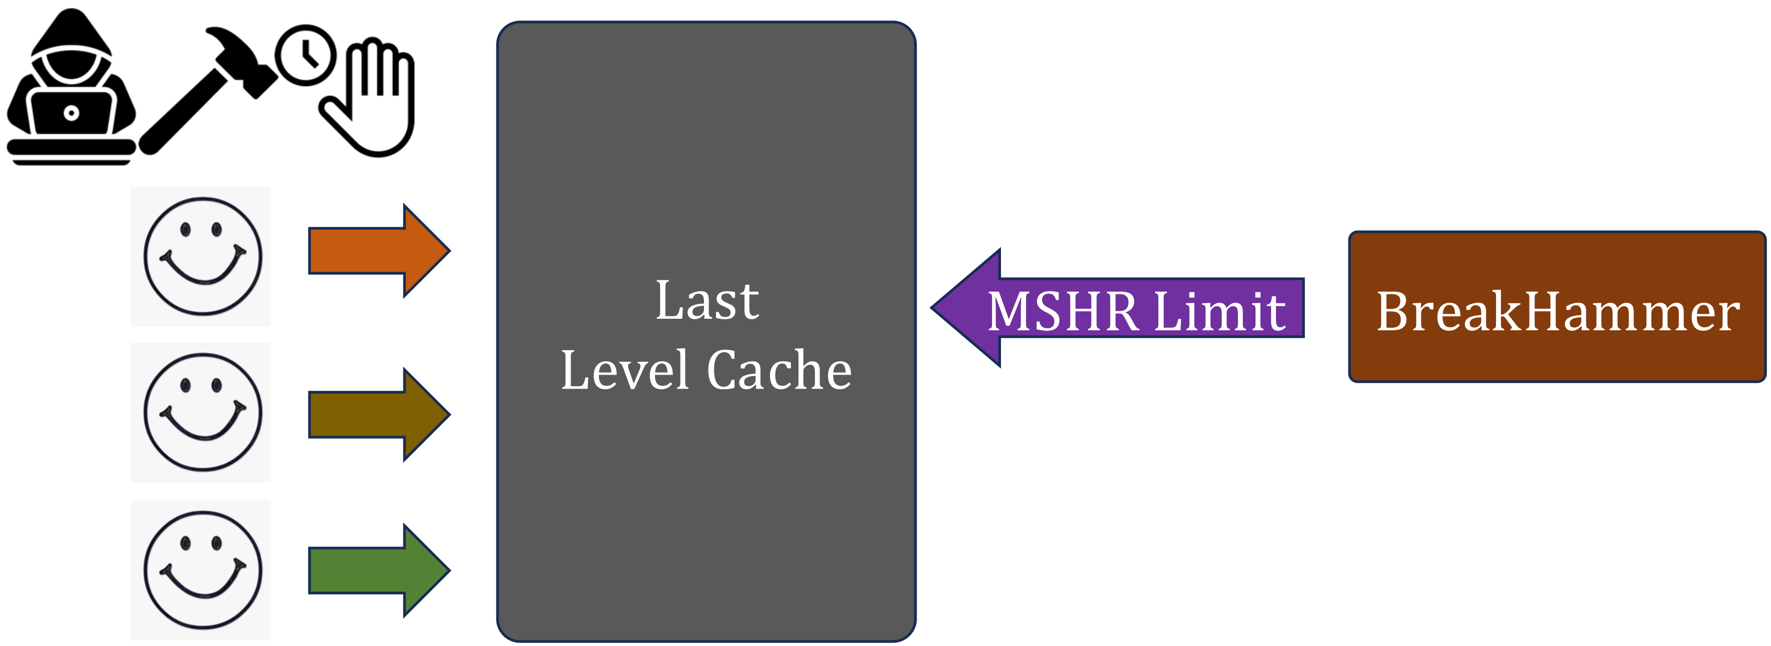

Future Research

Ongoing and Future Research on DRAM Read Disturbance

Deeper Understanding of Physics and Vulnerabilities

- Aging
- Online Profiling

Flexible and Intelligent Memory Chips, Interfaces, and Controllers

- In-field patchability
- DRAM-initiated pause
- Subarray-level parallelism
- Identify threads that cause the problem

Cross-Layer Communication

- Hardware-level detection
- System-level mitigation

Ongoing works

Broader Research Interests

Broader Research Interests

Thesis Publications

	2017	2018	2019	2020	2021	2022	2023	2024
Understanding Read Disturbance				Kim+ ISCA'20	**Orosa+ MICRO'21**	**Yaglikci+ DSN'22**	Olgun+ Disrupt'23 Luo+ ISCA'23	Olgun+, DSN'24 **Yaglikci+ HPCA'24**
Mitigating Read Disturbance					**Yaglikci+ HPCA'21**	**Yaglikci+ MICRO'22**		Olgun+, USENIX Sec'24 Bostanci+, HPCA'24
Read Disturbance Survey							Mutlu+ ASP-DAC'23	
Covert Channels and Security Primitives					Haj-Yahya+, ISCA'21 Olgun+, ISCA'21	Bostanci+ HPCA'22		Bostanci+ arXiv'24
Reducing DRAM Energy	Chang+ SIGMETRICS'17	Ghose+ SIGMETRICS'18	Koppula+ MICRO'19					
Reducing DRAM Latency			Hassan+ ISCA'19	Luo+ ISCA'20				
System Energy Saving Methods				Haj-Yahya+ ISCA'20		Haj-Yahya+ HPCA'22		
Processing Using Memory								Yuksel+, HPCA'24 Yuksel+ , DSN'24 Oliveira+, HPCA'24
Experimental Infrastructure							Olgun+, TACO'23 Luo+, CAL'23	

The Problem

Computing

is Bottlenecked by Data

Data is Key for AI, ML, Genomics, …

- Important workloads are all data intensive

- They require rapid and efficient processing of large amounts of data

- Data is increasing
  - We can generate more than we can process
  - We need to perform more sophisticated analyses on more data

Huge Demand for Performance & Efficiency

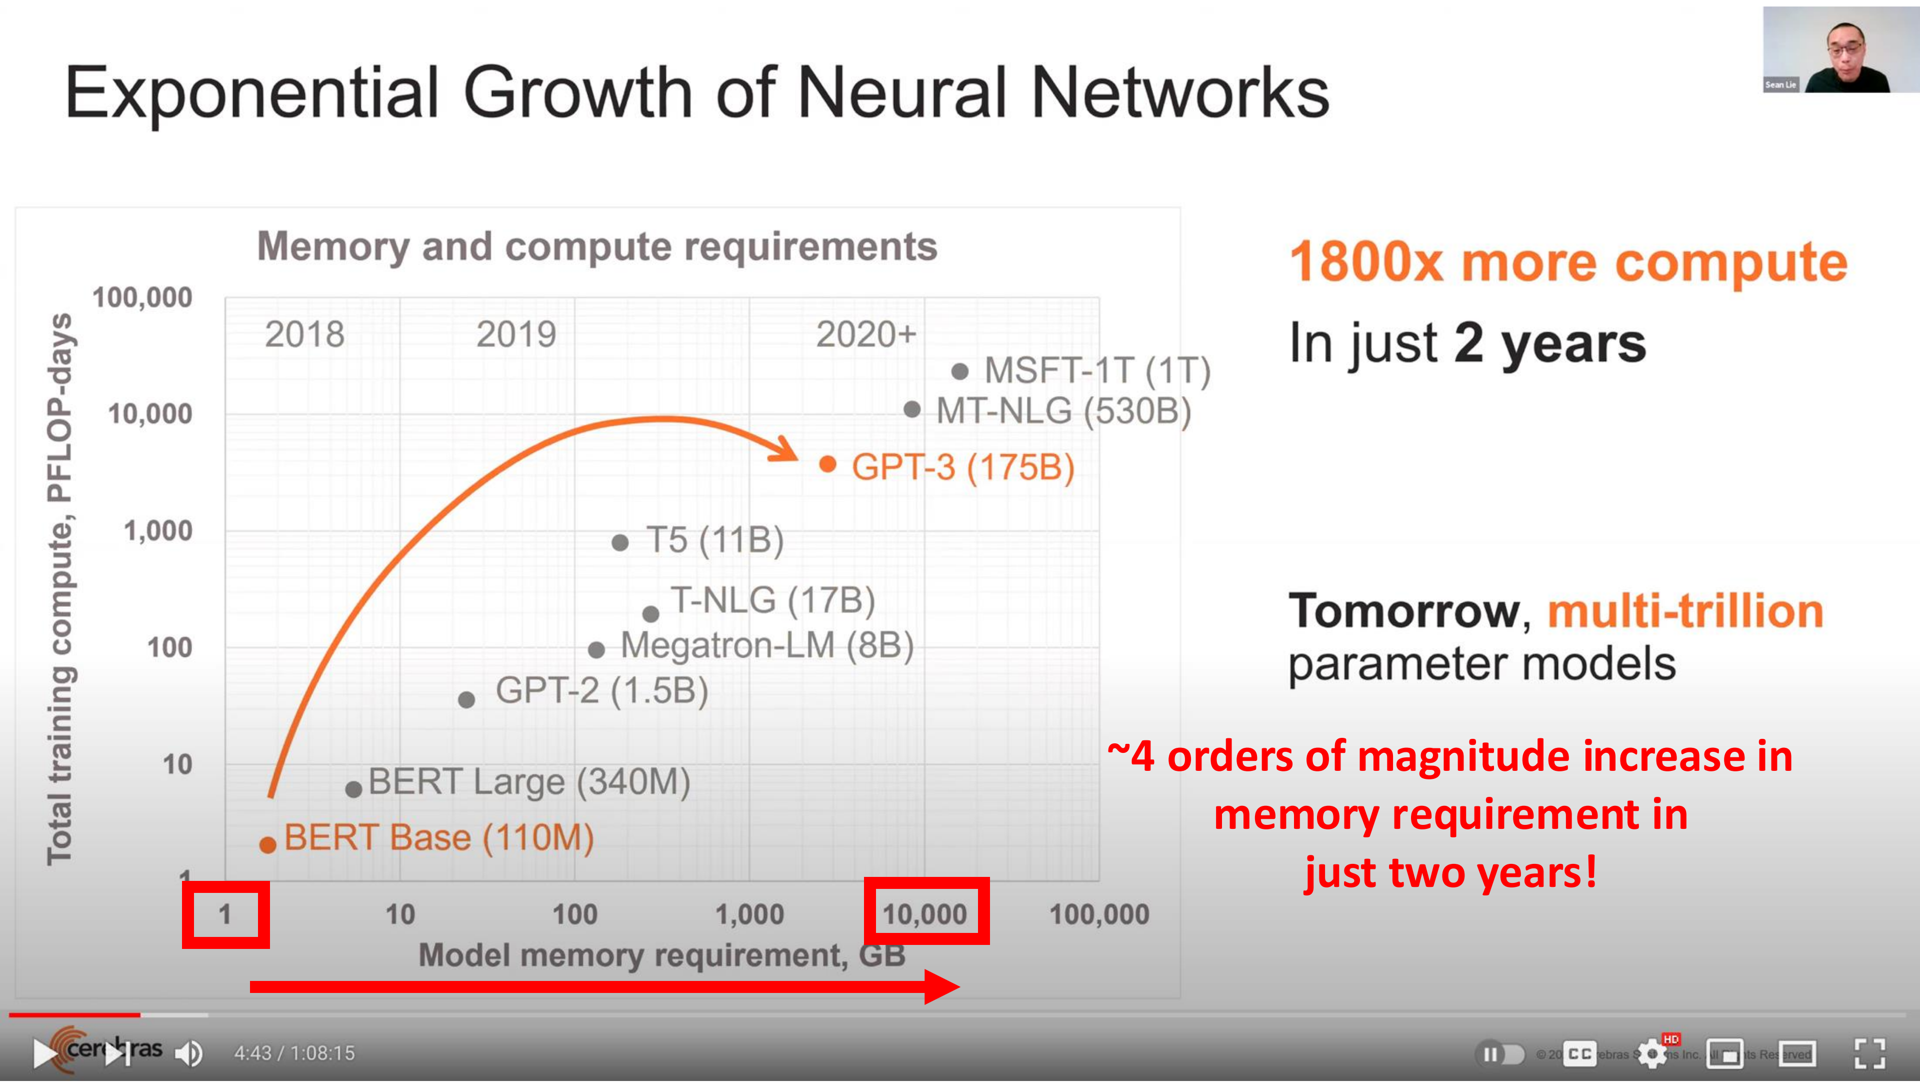

~4 orders of magnitude increase in memory requirement in just two years!

Data is Key for Future Workloads

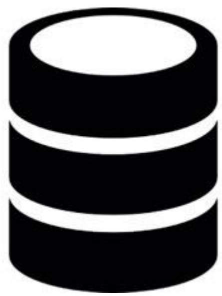

In-memory Databases
[Mao+, EuroSys'12;
Clapp+ (**Intel**), IISWC'15]

Graph/Tree Processing
[Xu+, IISWC'12; Umuroglu+, FPL'15]

In-Memory Data Analytics
[Clapp+ (**Intel**), IISWC'15;
Awan+, BDCloud'15]

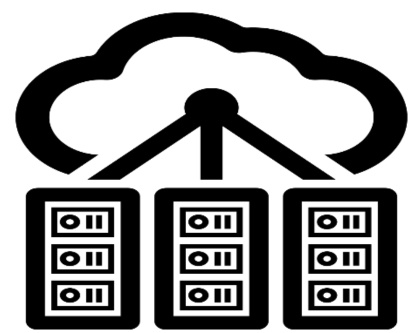

Datacenter Workloads
[Kanev+ (**Google**), ISCA'15]

Data Overwhelms Modern Machines

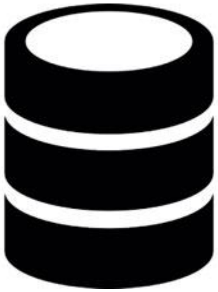

In-memory Databases

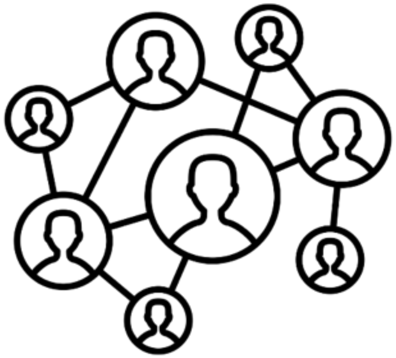

Graph/Tree Processing

Data → performance & energy bottleneck

In-Memory Data Analytics

[Clapp+ (**Intel**), IISWC'15;
Awan+, BDCloud'15]

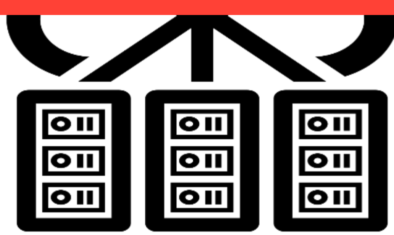

Datacenter Workloads

[Kanev+ (**Google**), ISCA'15]

Data is Key for Future Workloads

Chrome

Google's web browser

TensorFlow Mobile

Google's machine learning framework

Video Playback

Google's video codec

Video Capture

Google's video codec

Data Overwhelms Modern Machines

Chrome

TensorFlow Mobile

Data → performance & energy bottleneck

Video Playback

Google's **video codec**

Video Capture

Google's **video codec**

Data is Key for Future Workloads

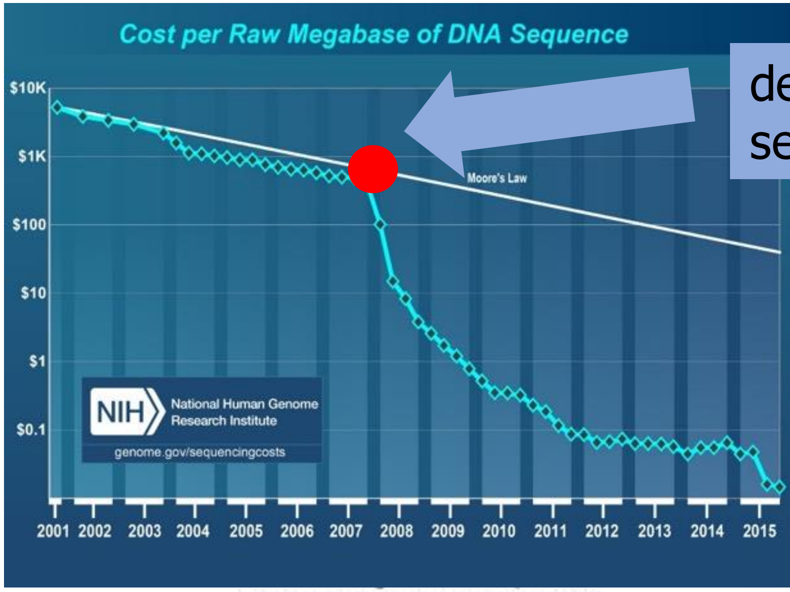

development of high-throughput sequencing (HTS) technologies

Number of Genomes Sequenced

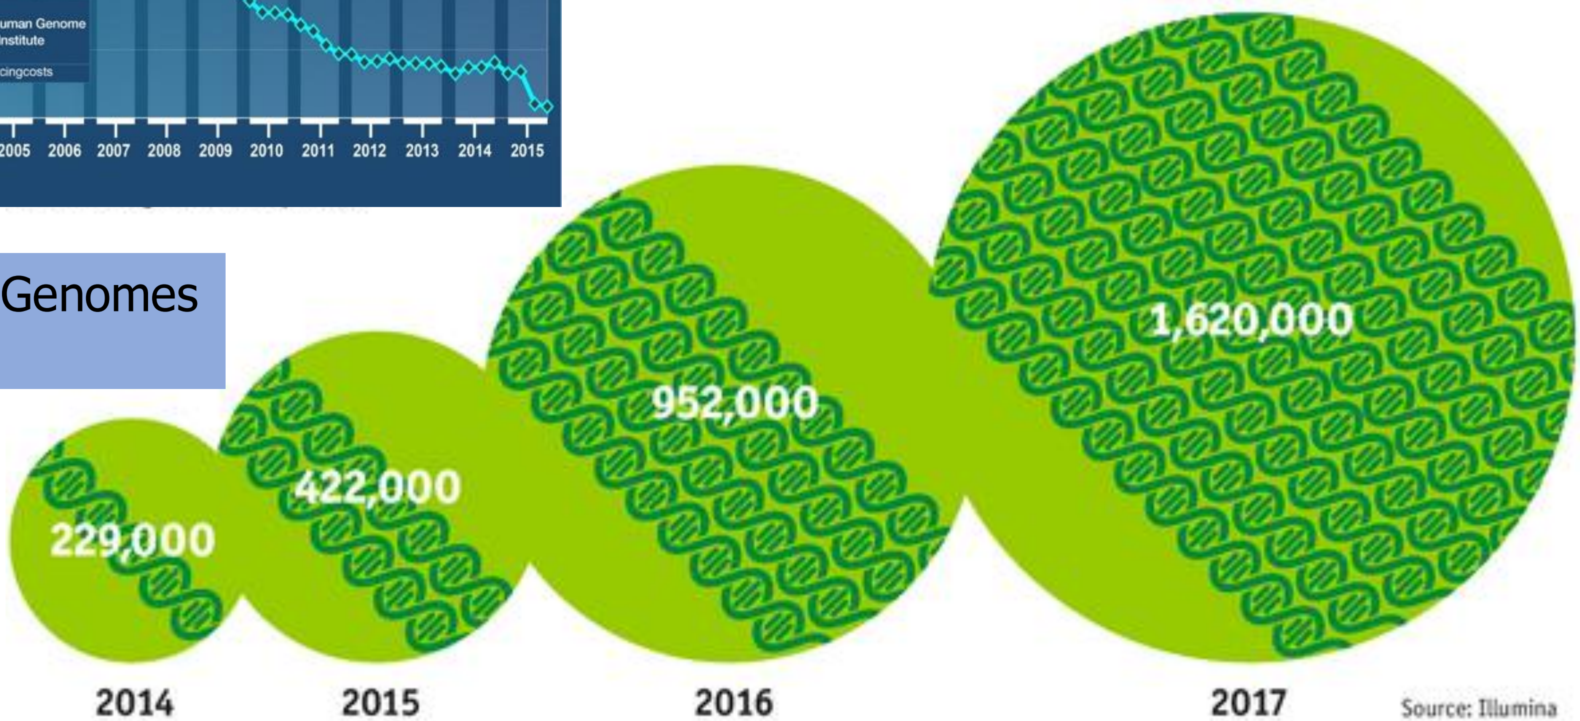

The Economist

http://www.economist.com/news/21631808-so-much-genetic-data-so-many-uses-genes-unzipped

Data is Key for Future Workloads

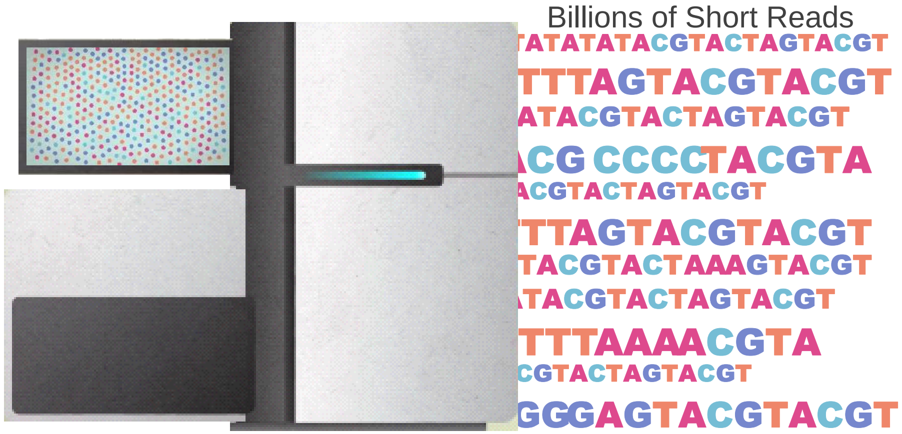

1 **Sequencing**

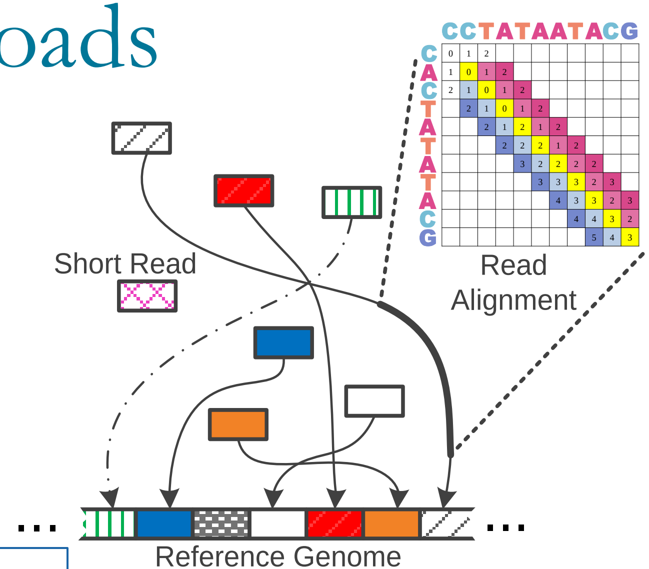

Read Mapping 2

Genome Analysis

Data → performance & energy bottleneck

```
read4:          CGCTTCCAT
read5:              CCATGACGC
read6:            TTCCATGAC
```

3 **Variant Calling**

Scientific Discovery 4

A Solution: Deeper and Larger Memory Hierarchies

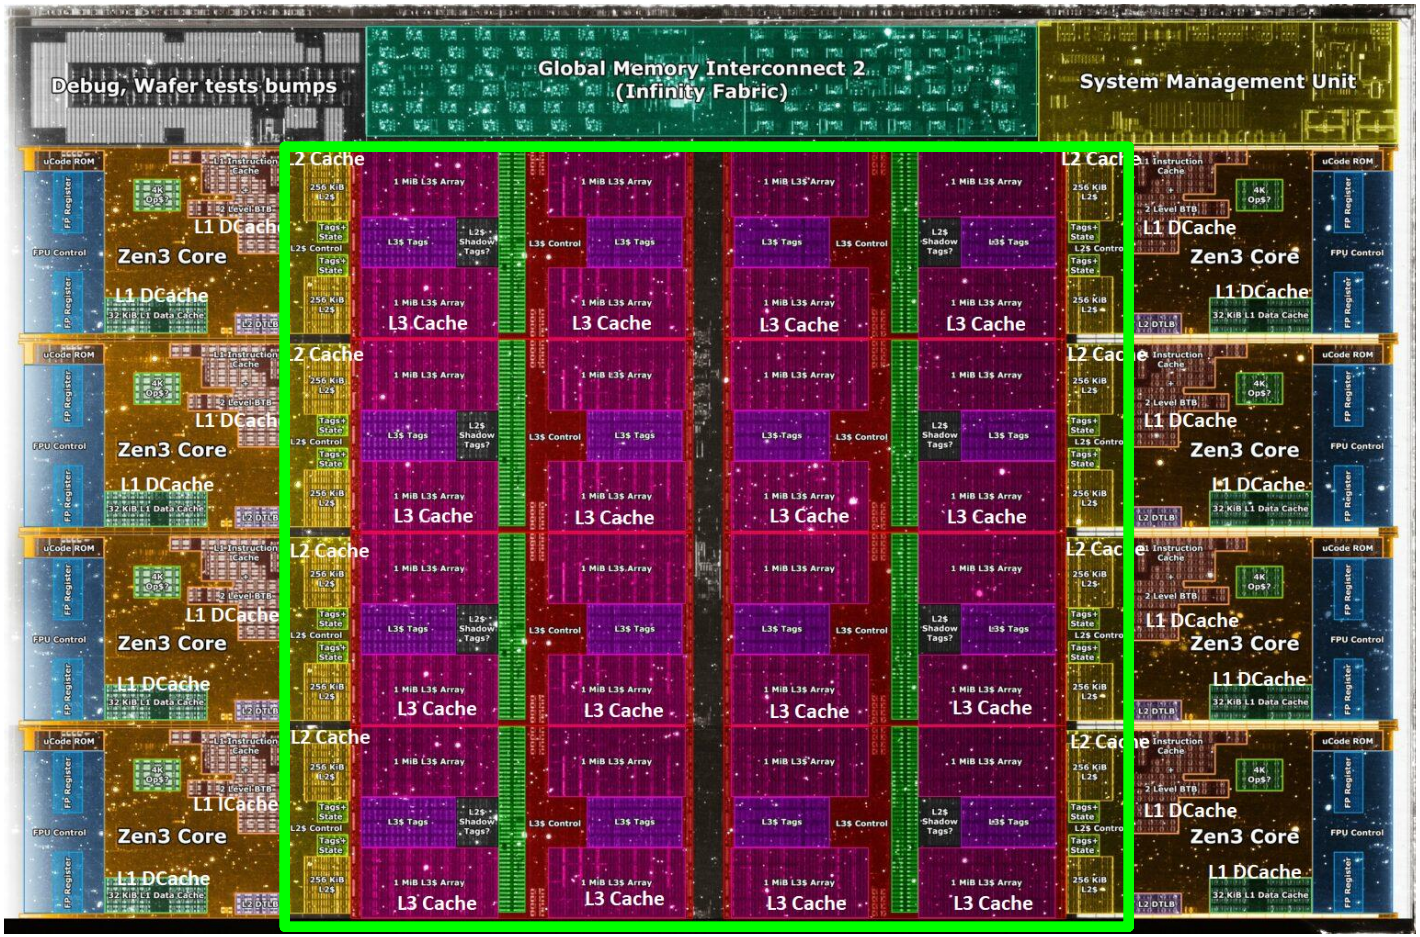

Core Count:
8 cores/16 threads

L1 Caches:
32 KB per core

L2 Caches:
512 KB per core

L3 Cache:
32 MB shared

AMD Ryzen 5000, 2020

https://wccftech.com/amd-ryzen-5000-zen-3-vermeer-undressed-high-res-die-shots-close-ups-pictured-detailed/

AMD's 3D Last Level Cache (2021)

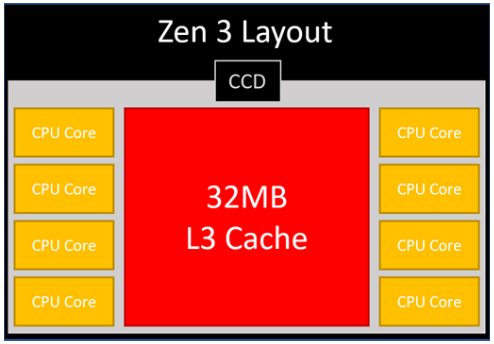

https://community.microcenter.com/discussion/5134/comparing-zen-3-to-zen-2

AMD increases the L3 size of their 8-core Zen 3 processors from 32 MB to 96 MB

Additional 64 MB L3 cache die stacked on top of the processor die
- Connected using Through Silicon Vias (TSVs)
- Total of 96 MB L3 cache

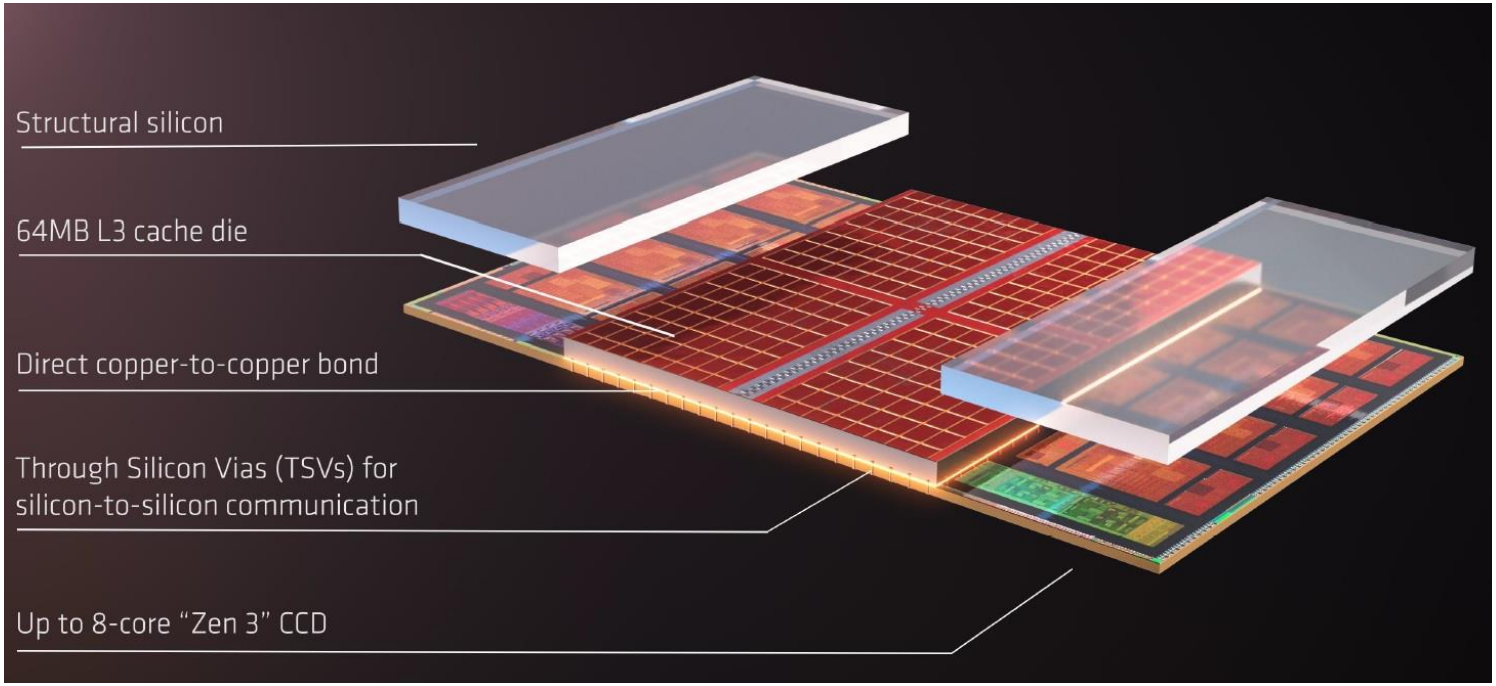

https://youtu.be/gqAYMx34euU
https://www.tech-critter.com/amd-keynote-computex-2021/

Deeper and Larger Memory Hierarchies

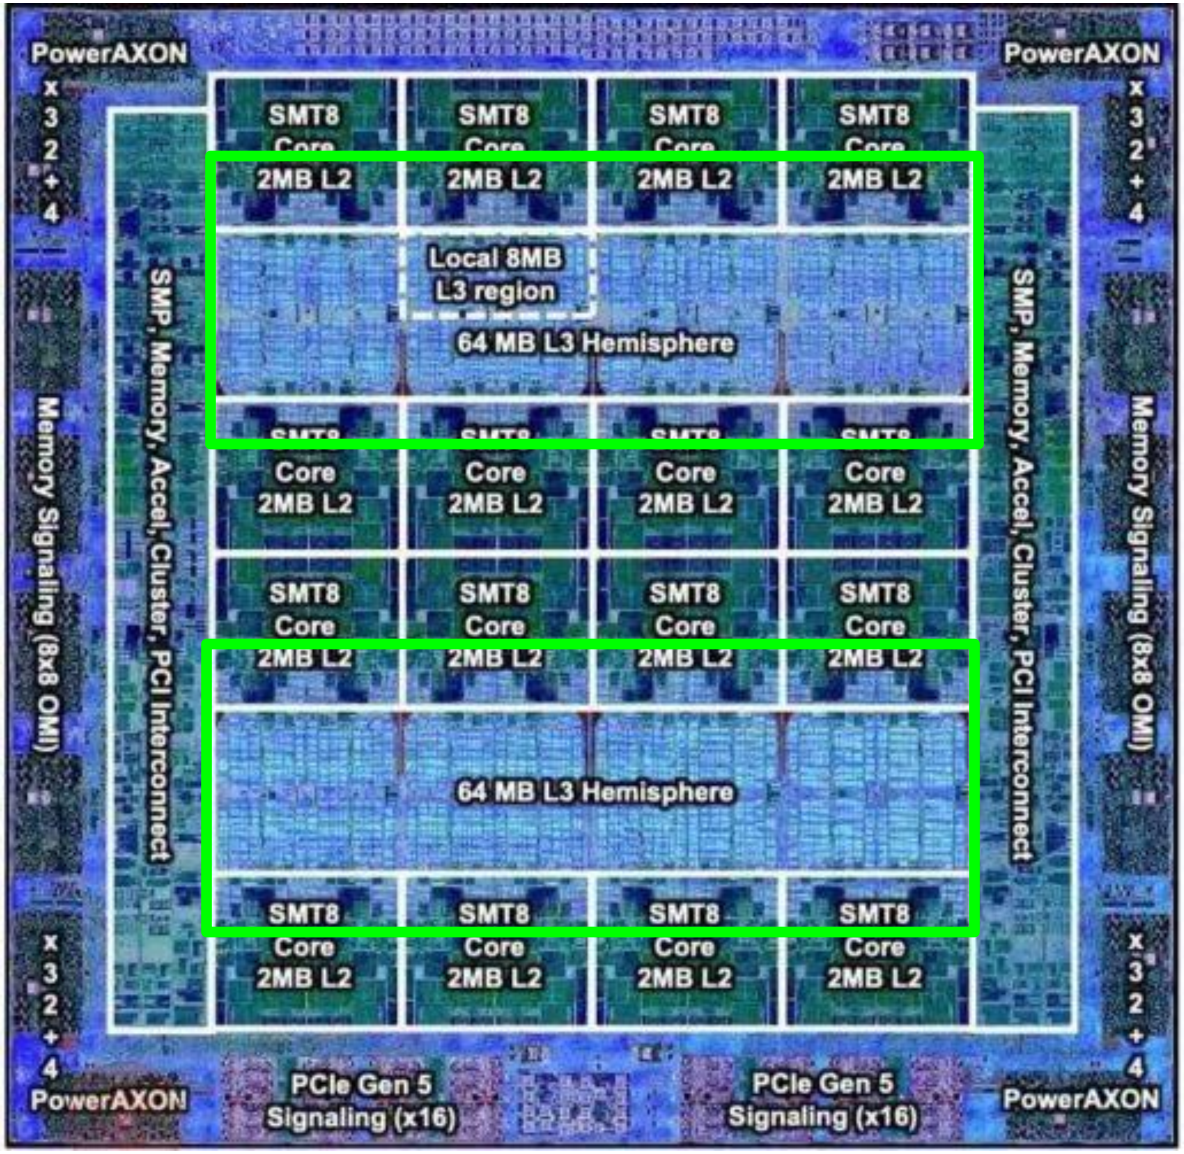

IBM POWER10, 2020

Cores:
15-16 cores,
8 threads/core

L2 Caches:
2 MB per core

L3 Cache:
120 MB shared

Deeper and Larger Memory Hierarchies

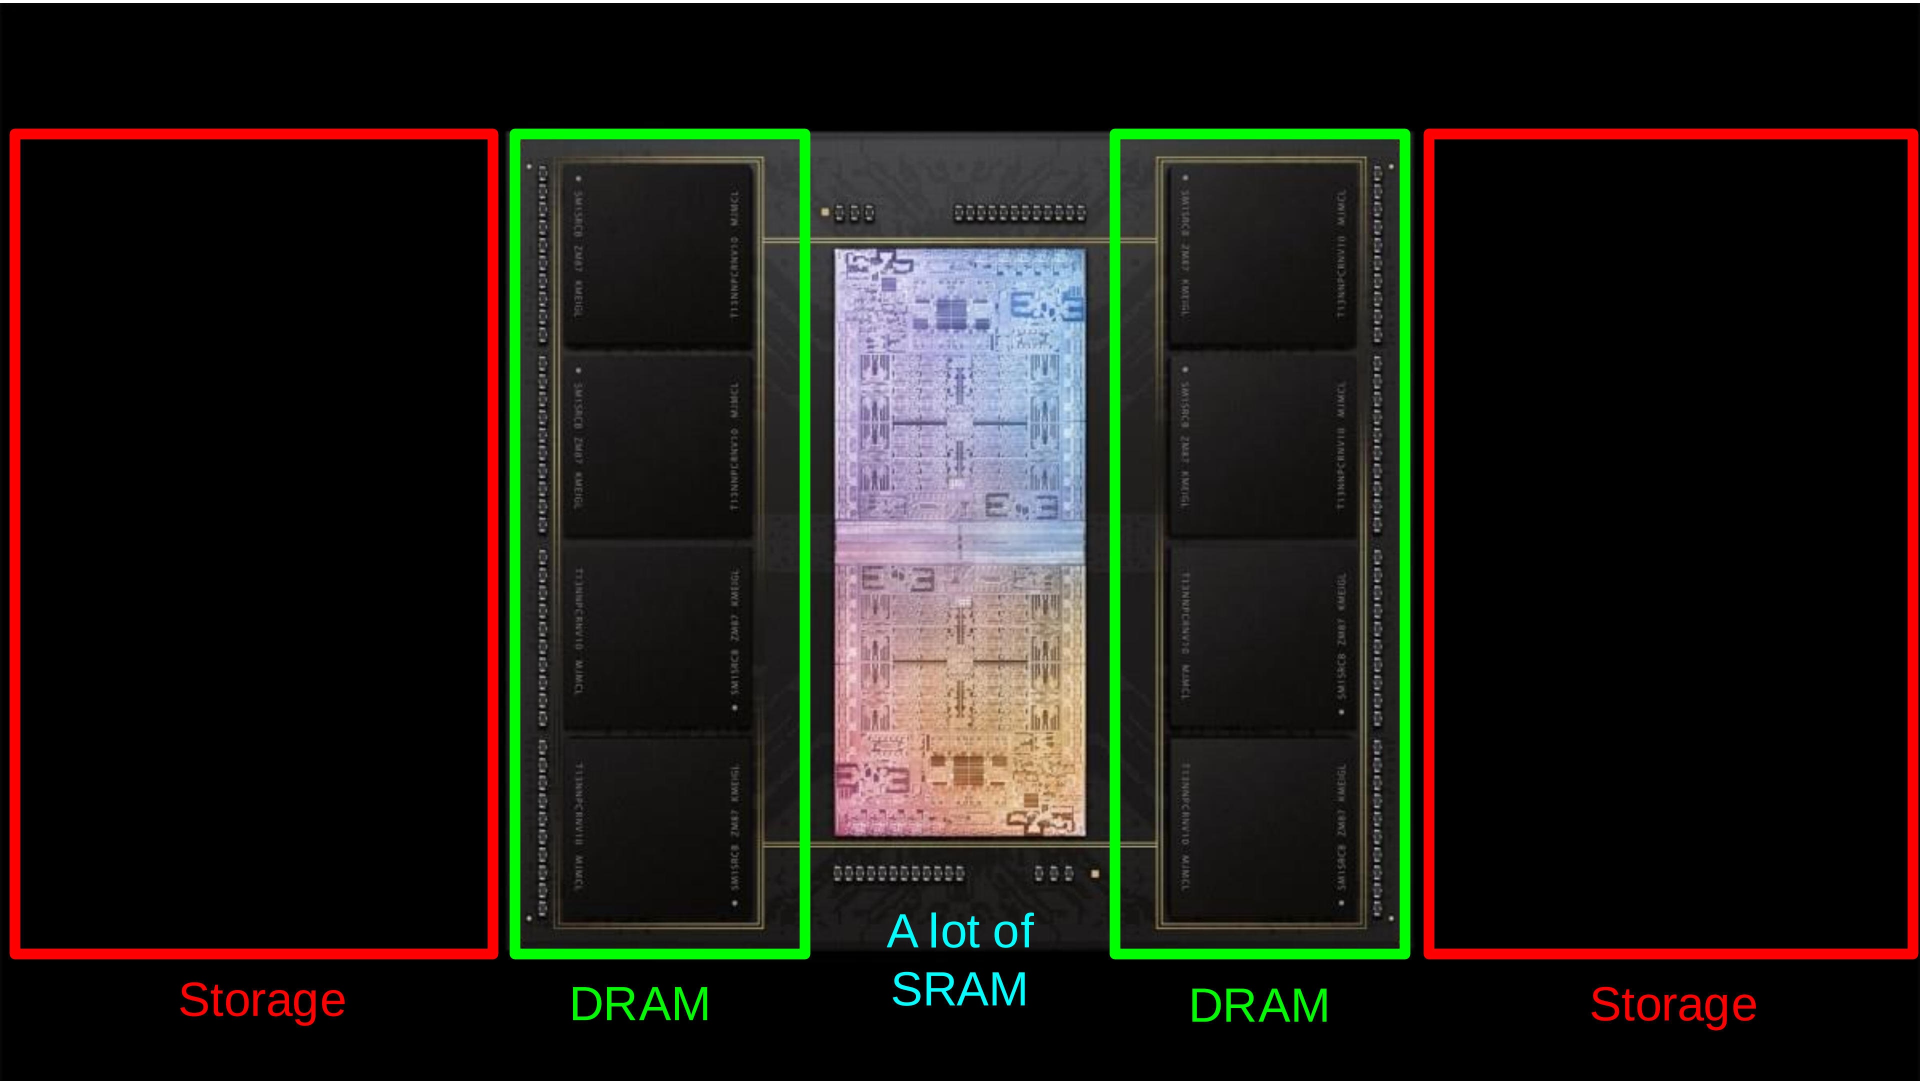

Apple M1 Ultra System (2022)

The Energy Perspective

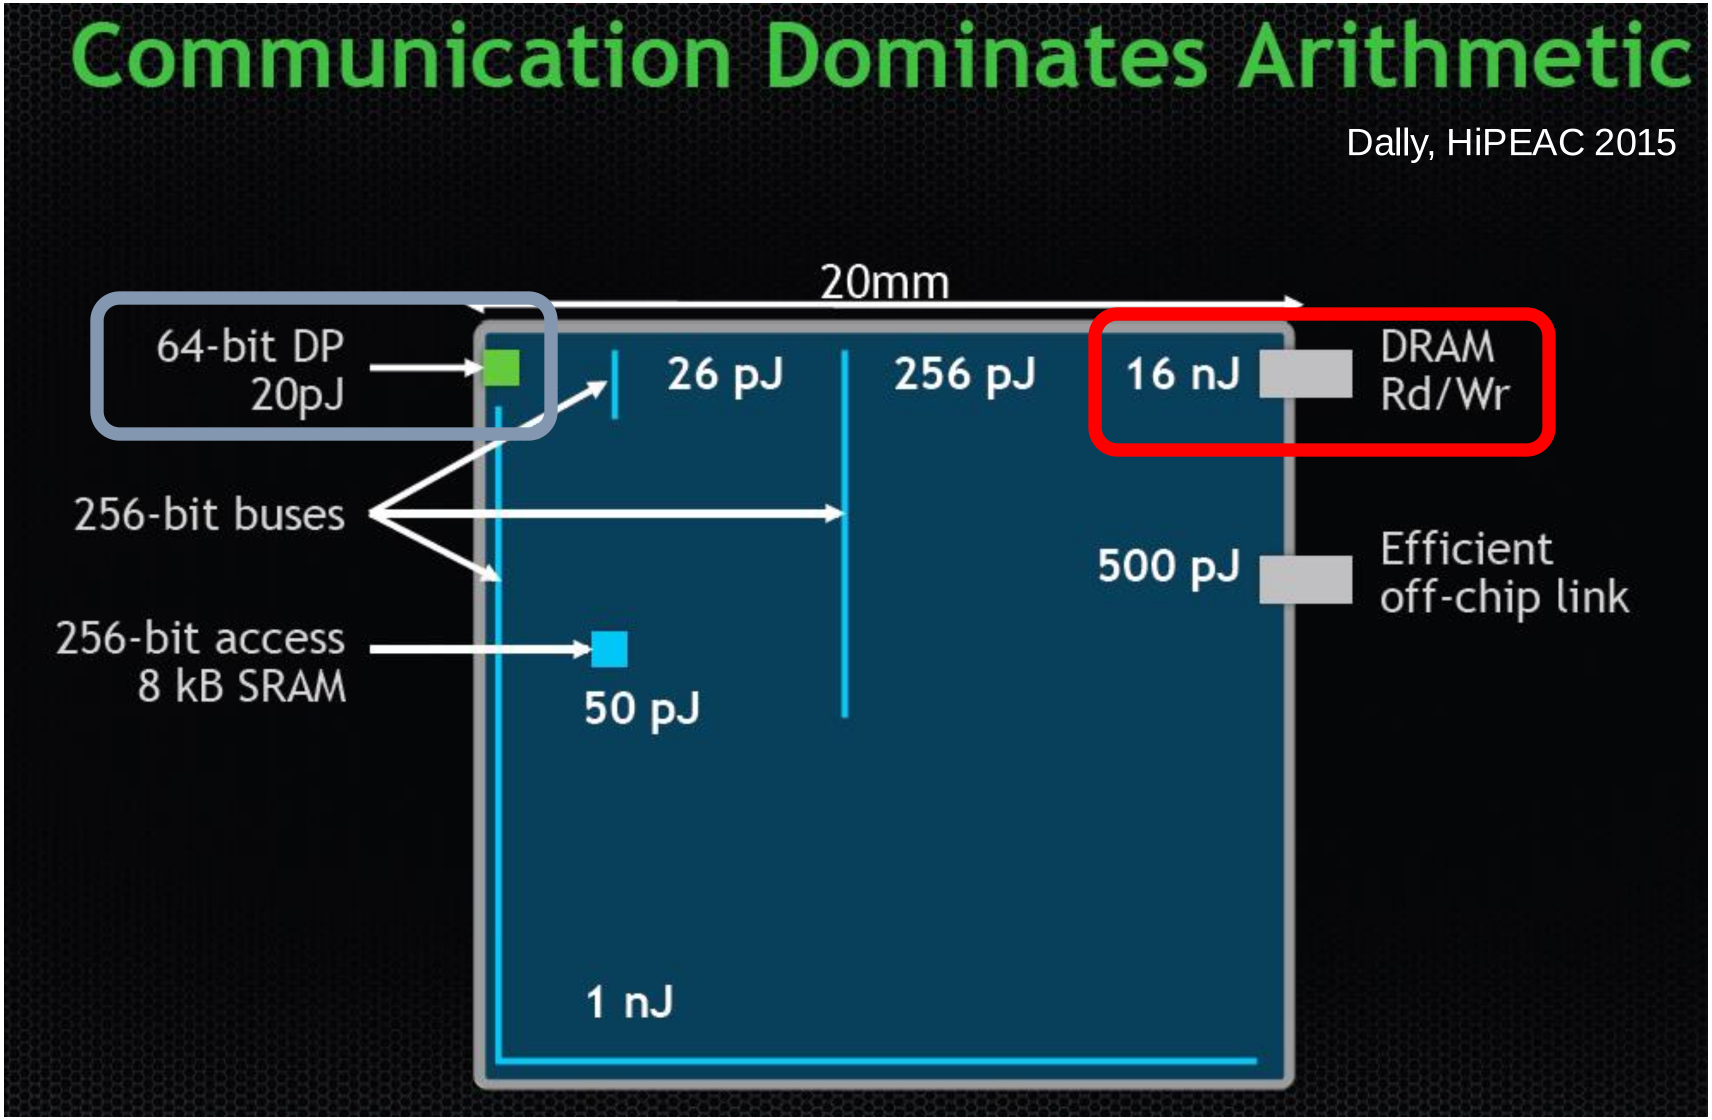

Data Movement vs Computation Energy

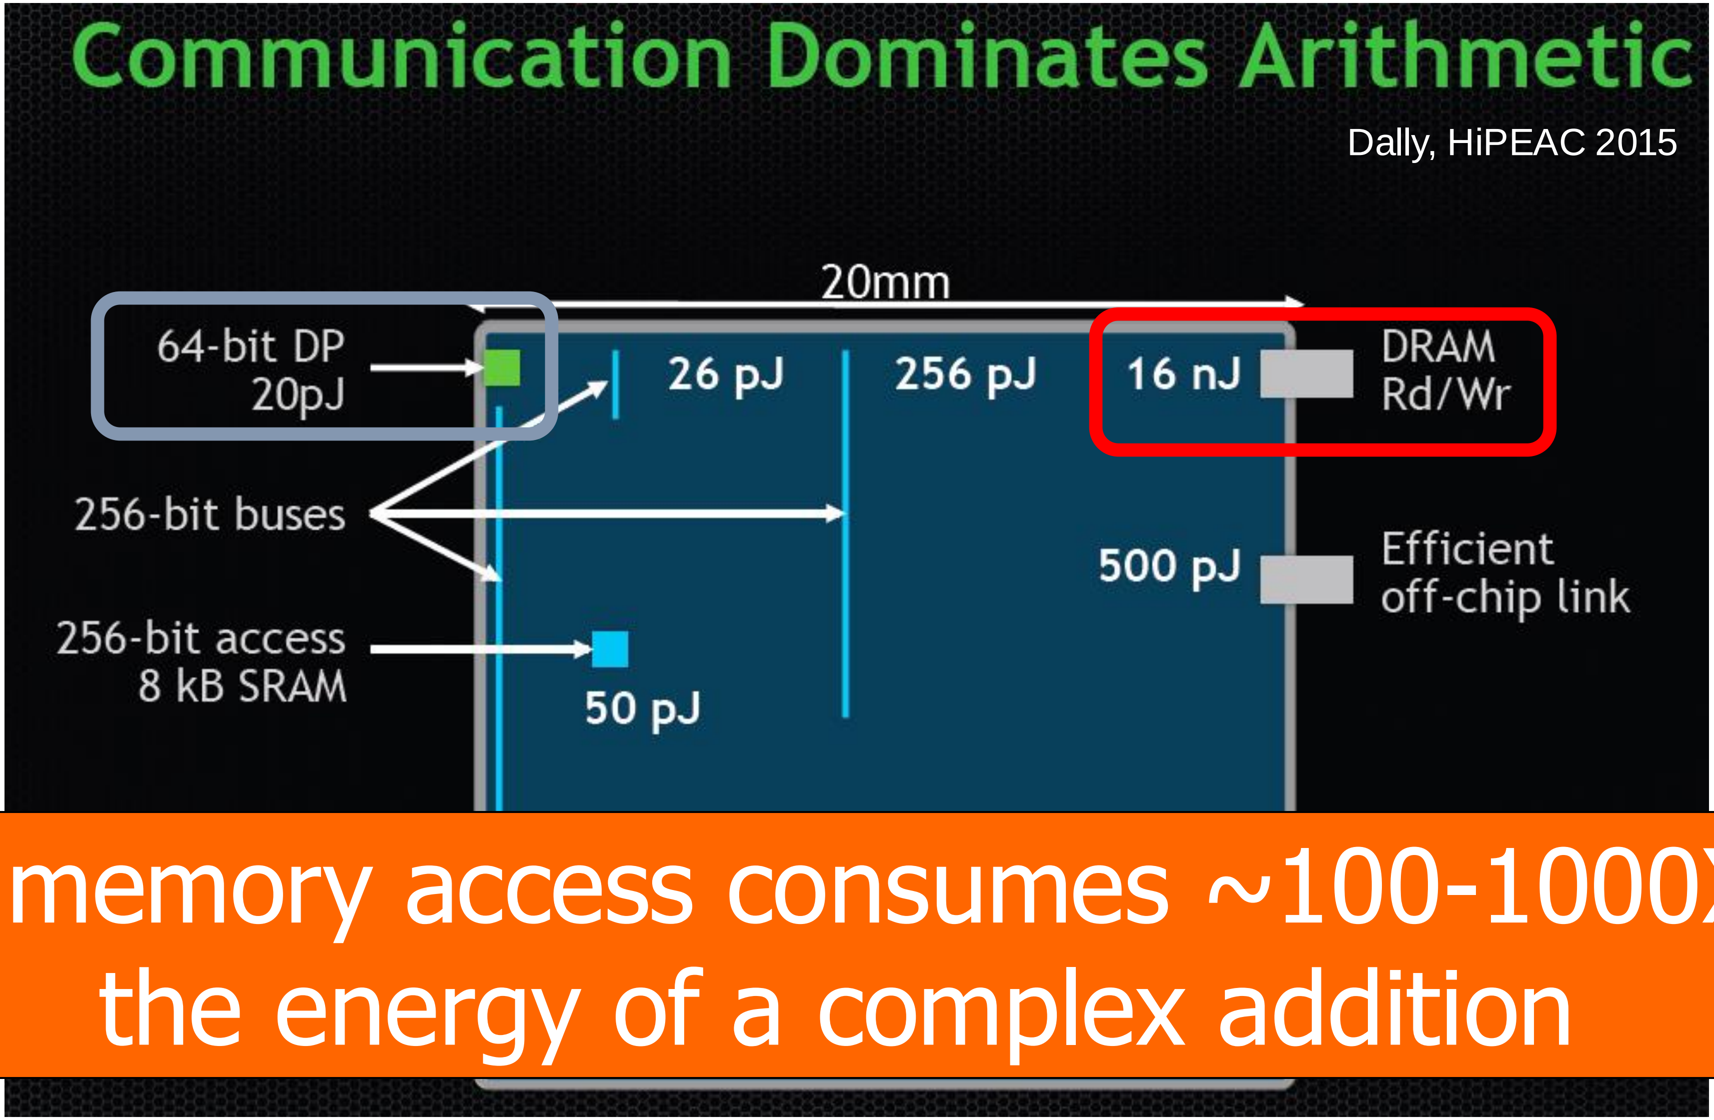

A memory access consumes ~100-1000X the energy of a complex addition

We Do Not Want to Move Data!

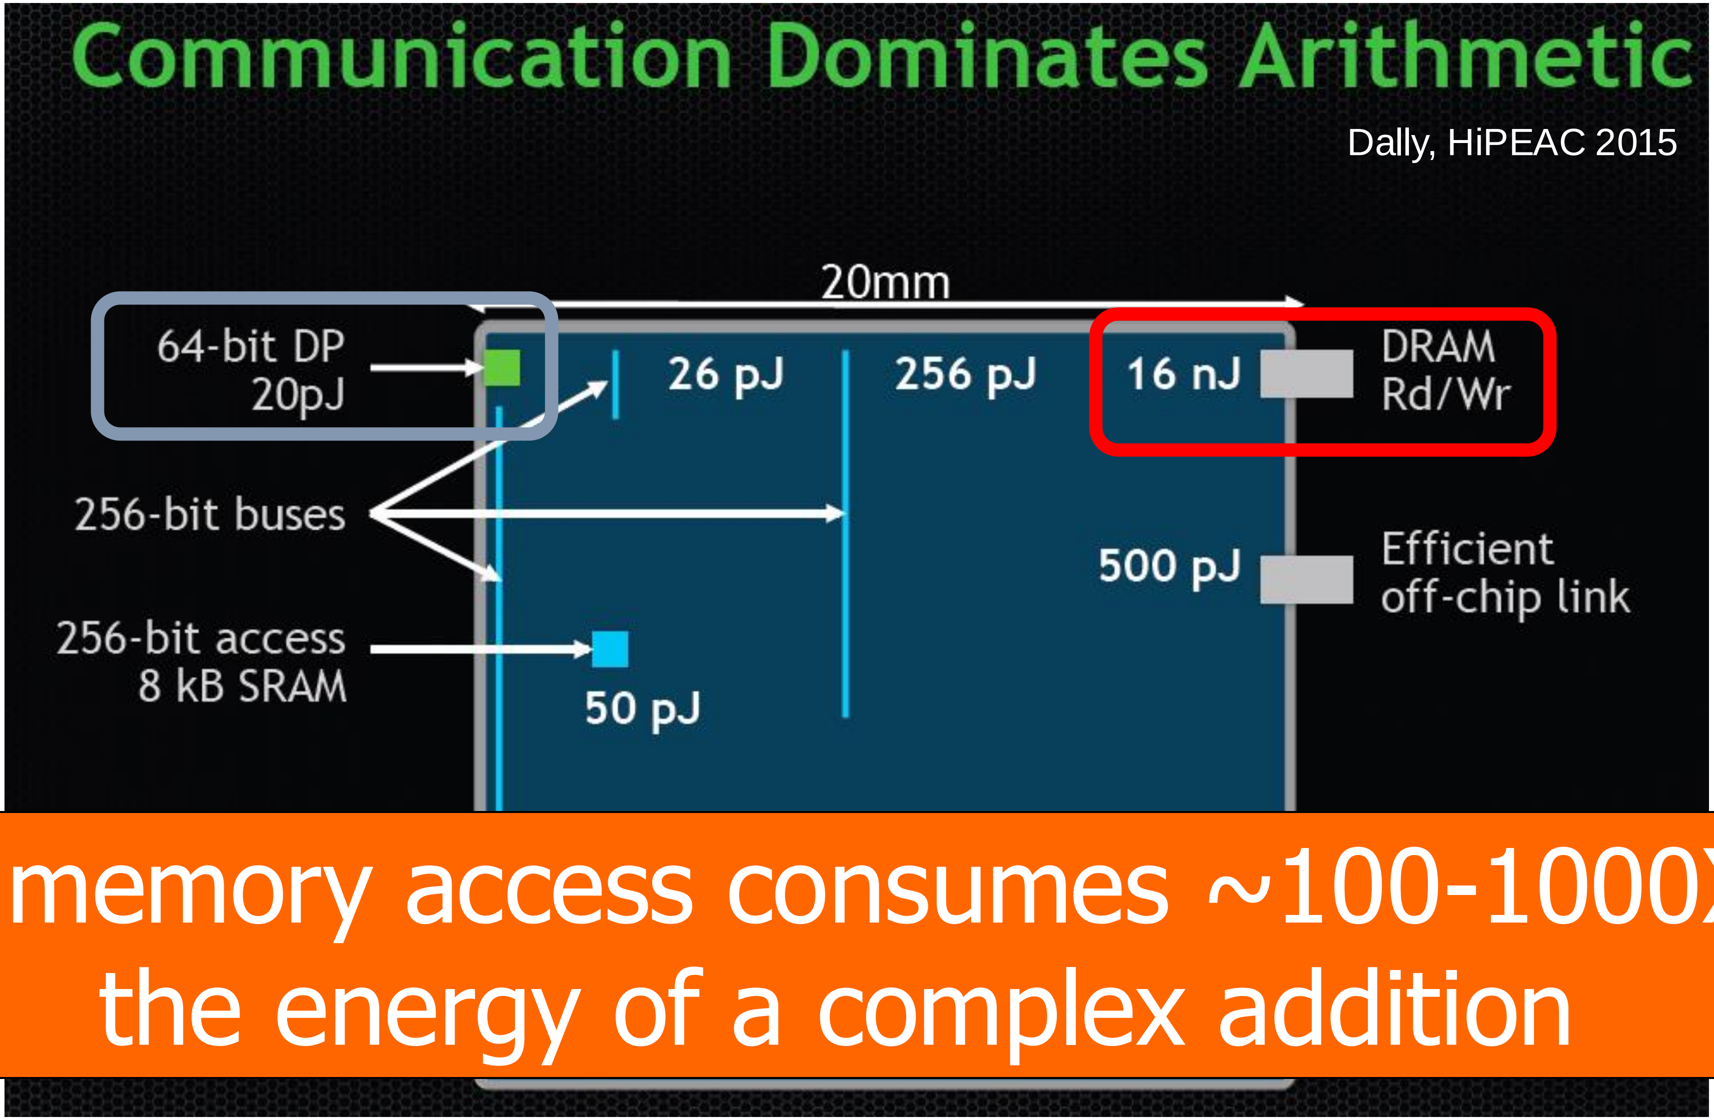

A memory access consumes ~100-1000X the energy of a complex addition

SAFARI

Data Movement vs. Computation Energy

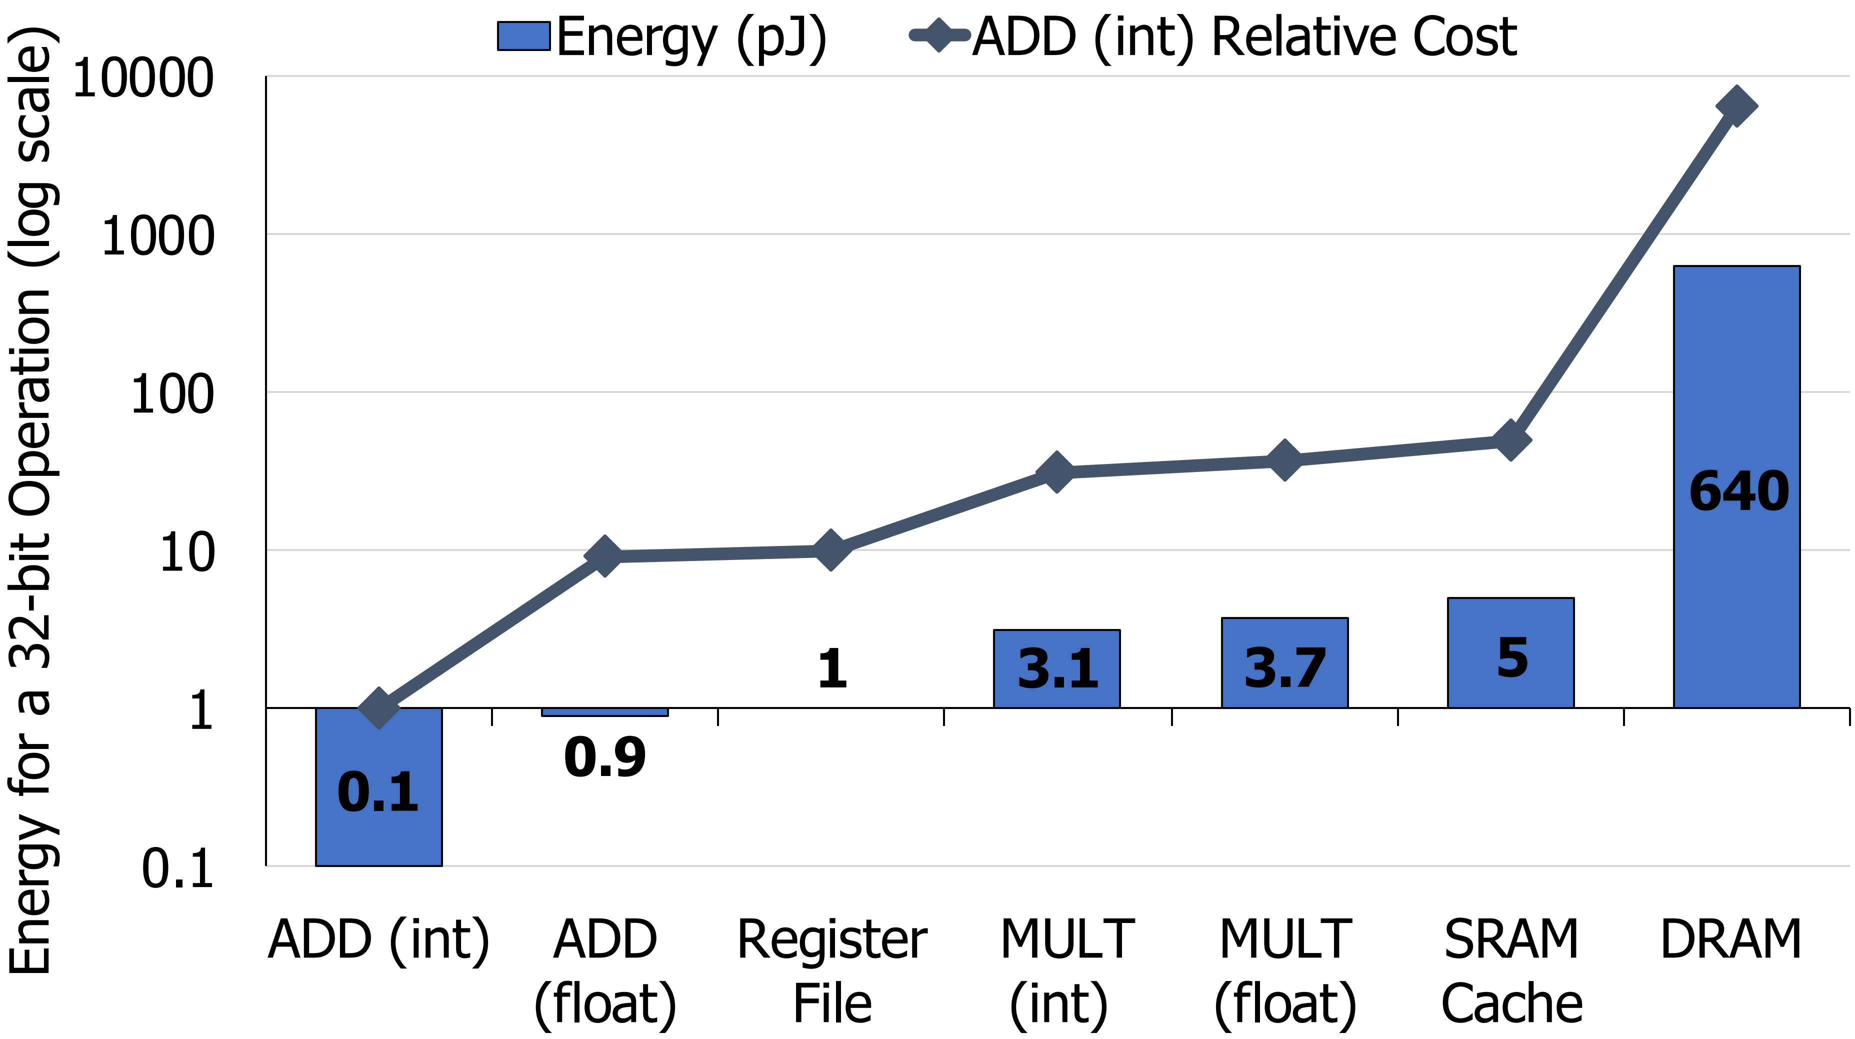

Han+, "EIE: Efficient Inference Engine on Compressed Deep Neural Network," ISCA 2016.

Data Movement vs. Computation Energy

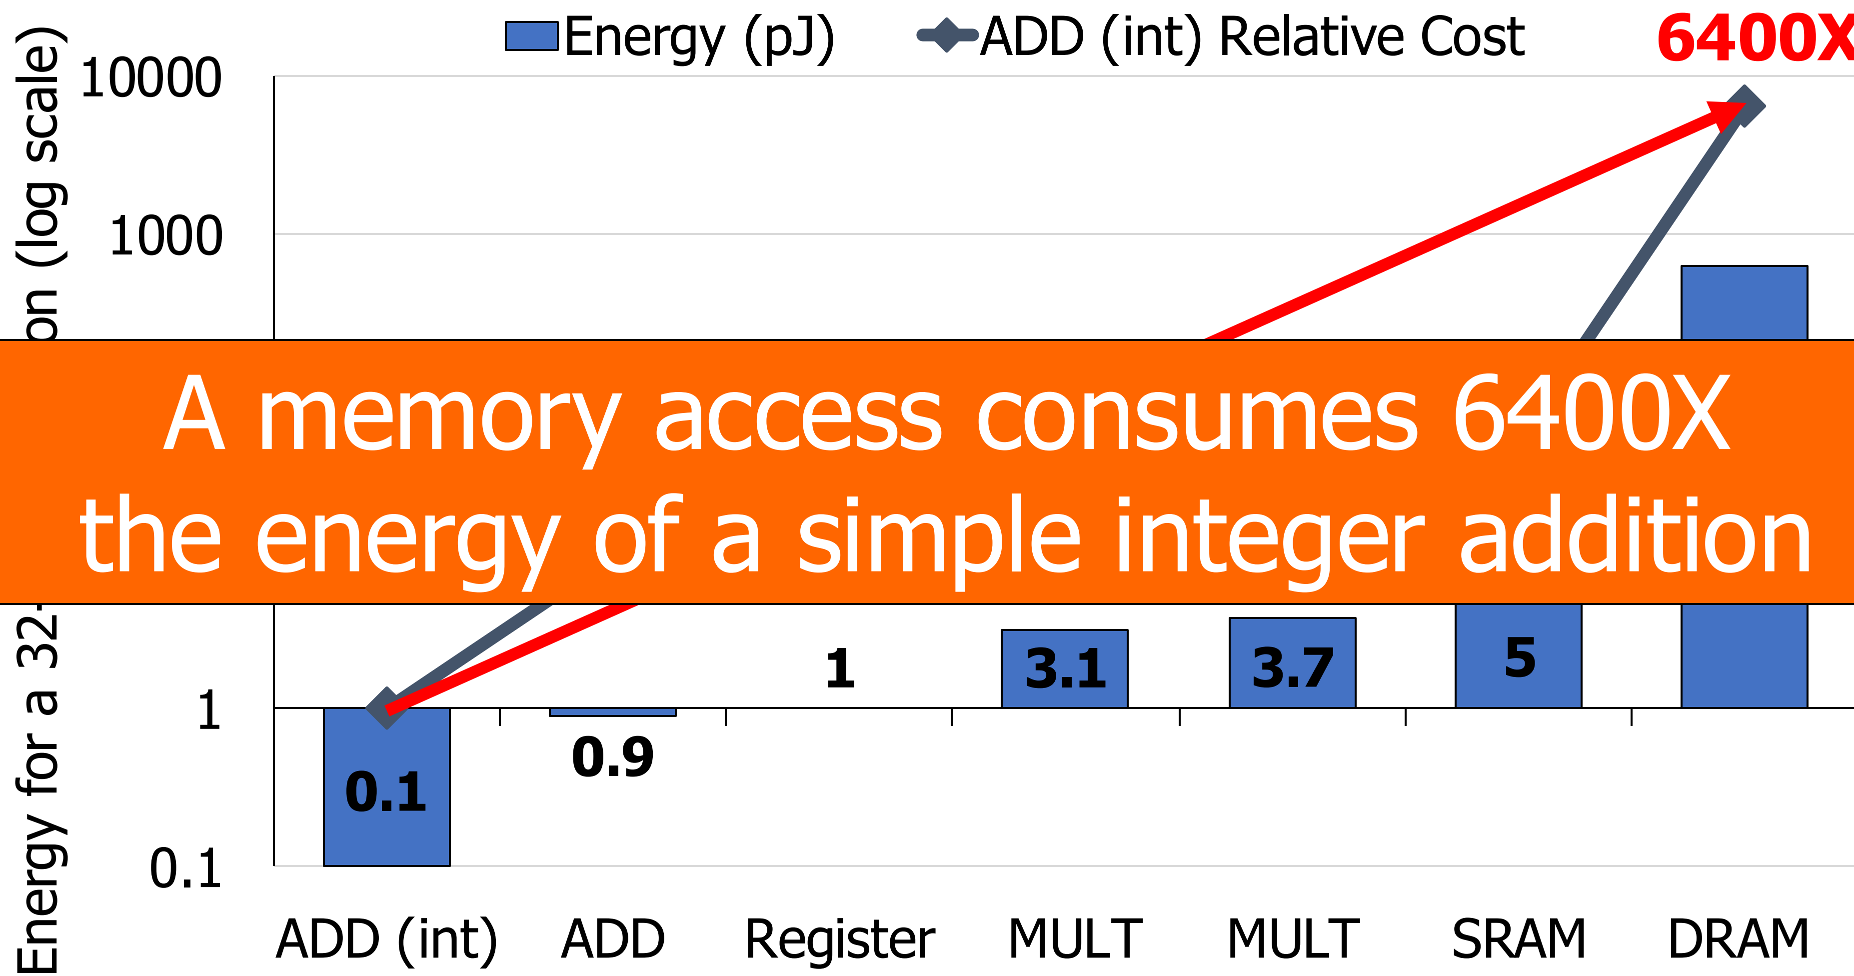

A memory access consumes 6400X the energy of a simple integer addition

Han+, "EIE: Efficient Inference Engine on Compressed Deep Neural Network," ISCA 2016.

Data Movement Overwhelms Modern Machines

- Amirali Boroumand, Saugata Ghose, Youngsok Kim, Rachata Ausavarungnirun, Eric Shiu, Rahul Thakur, Daehyun Kim, Aki Kuusela, Allan Knies, Parthasarathy Ranganathan, and Onur Mutlu, **"Google Workloads for Consumer Devices: Mitigating Data Movement Bottlenecks"** *Proceedings of the 23rd International Conference on Architectural Support for Programming Languages and Operating Systems* (***ASPLOS***), Williamsburg, VA, USA, March 2018.

62.7% of the total system energy is spent on data movement

Google Workloads for Consumer Devices: Mitigating Data Movement Bottlenecks

Amirali Boroumand[1] Saugata Ghose[1] Youngsok Kim[2]
Rachata Ausavarungnirun[1] Eric Shiu[3] Rahul Thakur[3] Daehyun Kim[4,3]
Aki Kuusela[3] Allan Knies[3] Parthasarathy Ranganathan[3] Onur Mutlu[5,1]

Data Movement Overwhelms Accelerators

- Amirali Boroumand, Saugata Ghose, Berkin Akin, Ravi Narayanaswami, Geraldo F. Oliveira, Xiaoyu Ma, Eric Shiu, and Onur Mutlu,
  "Google Neural Network Models for Edge Devices: Analyzing and Mitigating Machine Learning Inference Bottlenecks"
  Proceedings of the 30th International Conference on Parallel Architectures and Compilation Techniques (***PACT***), Virtual, September 2021.
  [Slides (pptx) (pdf)]
  [Talk Video (14 minutes)]

> 90% of the total system energy is spent on memory in large ML models

Google Neural Network Models for Edge Devices:
Analyzing and Mitigating Machine Learning Inference Bottlenecks

Amirali Boroumand[†◇] Saugata Ghose[‡] Berkin Akin[§] Ravi Narayanaswami[§]
Geraldo F. Oliveira[⋆] Xiaoyu Ma[§] Eric Shiu[§] Onur Mutlu[⋆†]

[†]*Carnegie Mellon Univ.* [◇]*Stanford Univ.* [‡]*Univ. of Illinois Urbana-Champaign* [§]*Google* [⋆]*ETH Zürich*

Yet …

I expect that over the coming decade memory subsystem design will be the *only* important design issue for microprocessors.

- "**It's the Memory, Stupid!**" (Richard Sites, MPR, 1996)

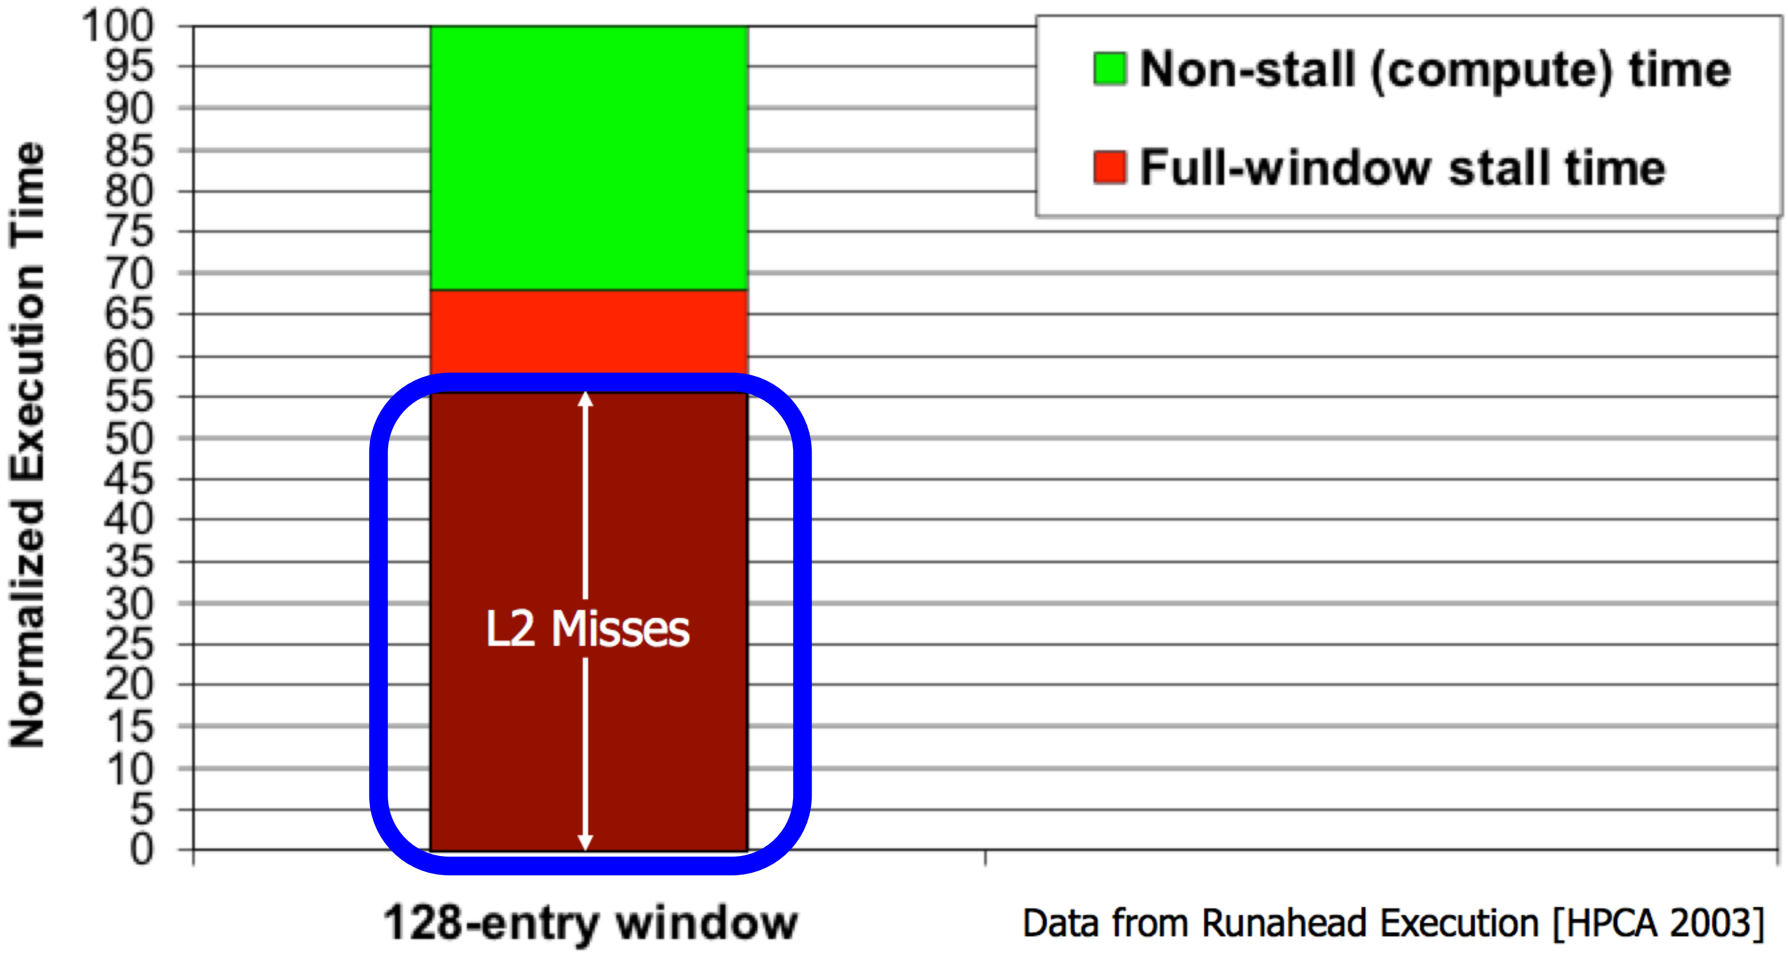

Data from Runahead Execution [HPCA 2003]

Mutlu+, "Runahead Execution: An Alternative to Very Large Instruction Windows for Out-of-Order Processors," HPCA 2003.

Performance Perspective (Today)

- All of Google's Data Center Workloads (2015):

Kanev+, "Profiling a Warehouse-Scale Computer," ISCA 2015.

Performance Perspective (Today)

- All of Google's Data Center Workloads (2015):

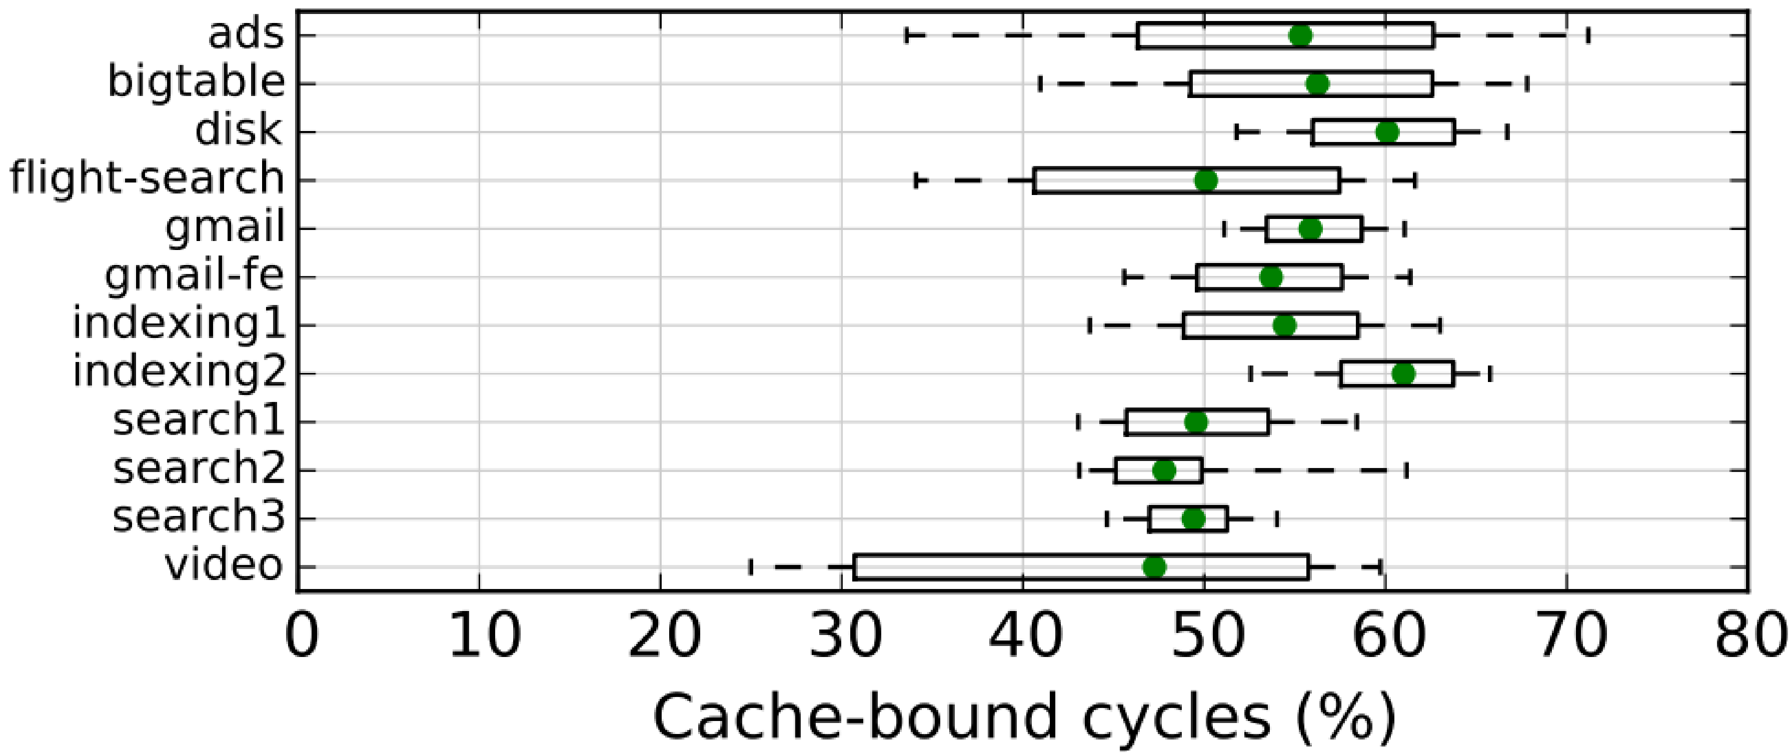

Figure 11: Half of cycles are spent stalled on caches.

Kanev+, "Profiling a Warehouse-Scale Computer," ISCA 2015.

Problem

Data access is the major performance and energy bottleneck

Our current
design principles cause
great energy waste
and
performance loss

Problem

Processing of data
is performed
far away from the data

Promising Solutions: Processing Near Memory (PnM)

- UPMEM, founded in January 2015, announces the first real-world PIM architecture in 2016
- UPMEM's PIM-enabled DIMMs start getting commercialized in 2019

- In early 2021, Samsung announces FIMDRAM at ISSCC conference
- Samsung's LP-DDR5 and DIMM-based PIM announced a few months later

- In early 2022, SK Hynix announces AiM and Alibaba announces HB-PNM at ISSCC conference

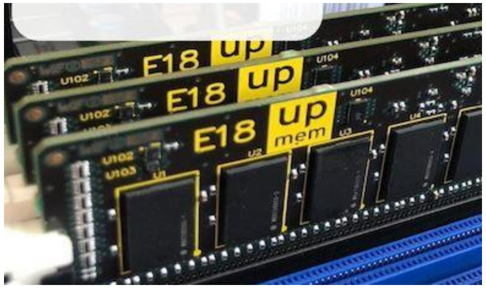

Startup plans to embed processors in DRAM

October 13, 2016 // By Peter Clarke

Email print Share Share reddit

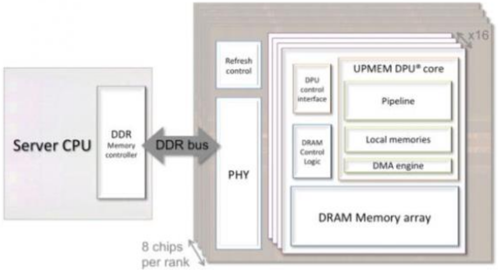

Fabless chip company Upmem SAS (Grenoble, France), founded in January 2015, is developing a microprocessor for use in data-intensive applications in the datacenter that will sit embedded in DRAM to be close to the data.

Promising Solutions: Processing Near Memory (PnM)

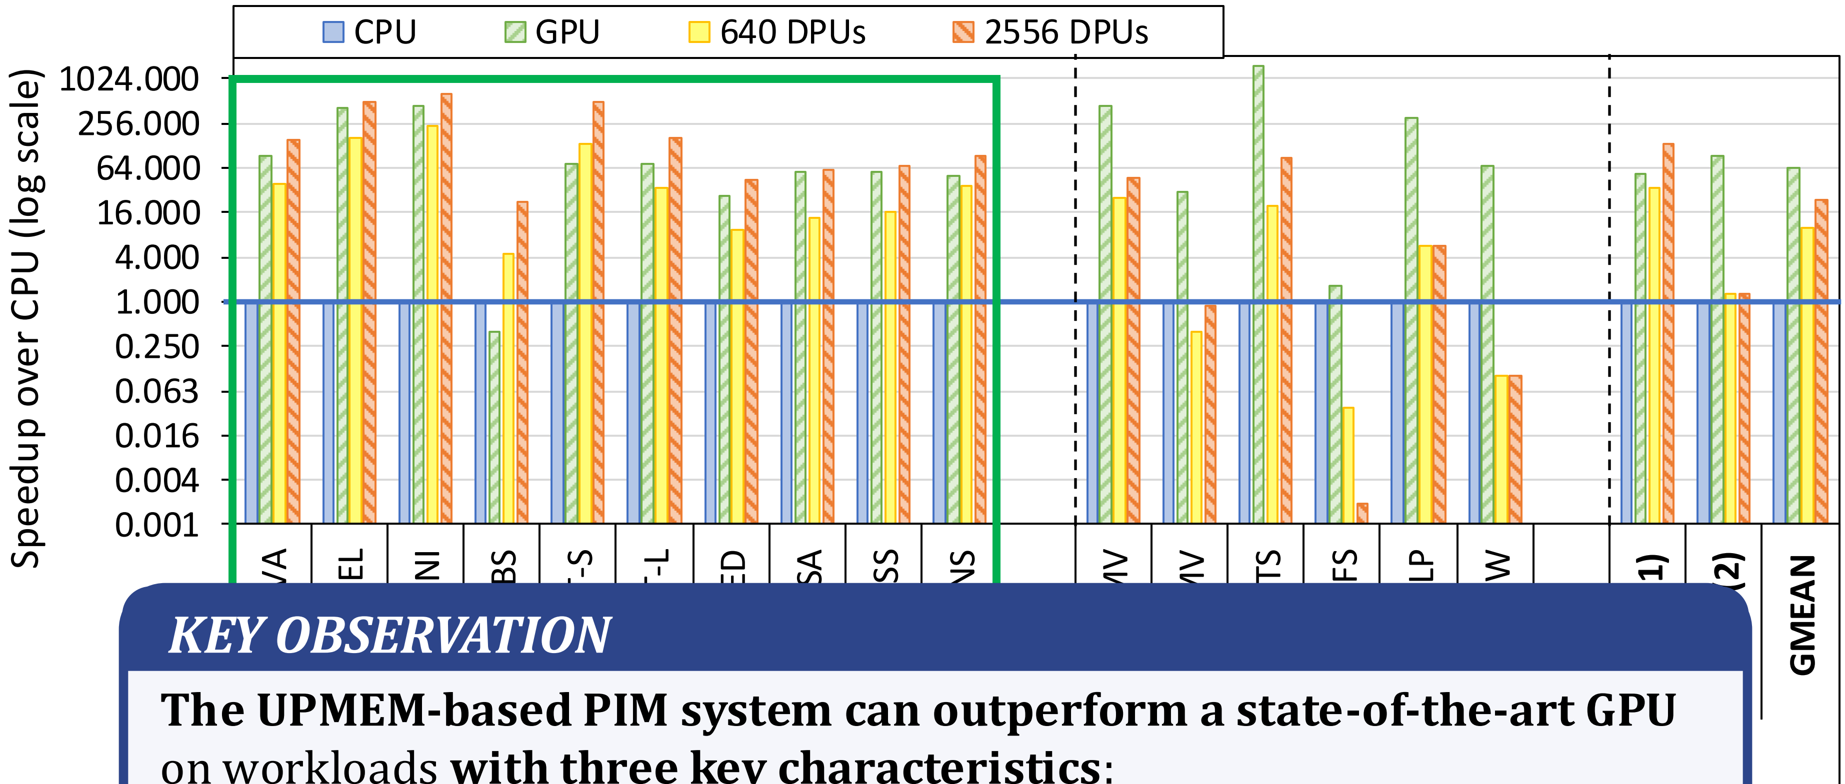

KEY OBSERVATION

The UPMEM-based PIM system can outperform a state-of-the-art GPU on workloads **with three key characteristics**:

1. Streaming memory accesses
2. No or little inter-DPU synchronization
3. No or little use of integer multiplication, integer division, or floating point operations

These three key characteristics make a **workload potentially suitable to the UPMEM PIM architecture**.

Promising Solutions: Processing Using Memory (PuM) Implementing Data Movement

Idea: Two consecutive ACTivates

Negligible HW cost

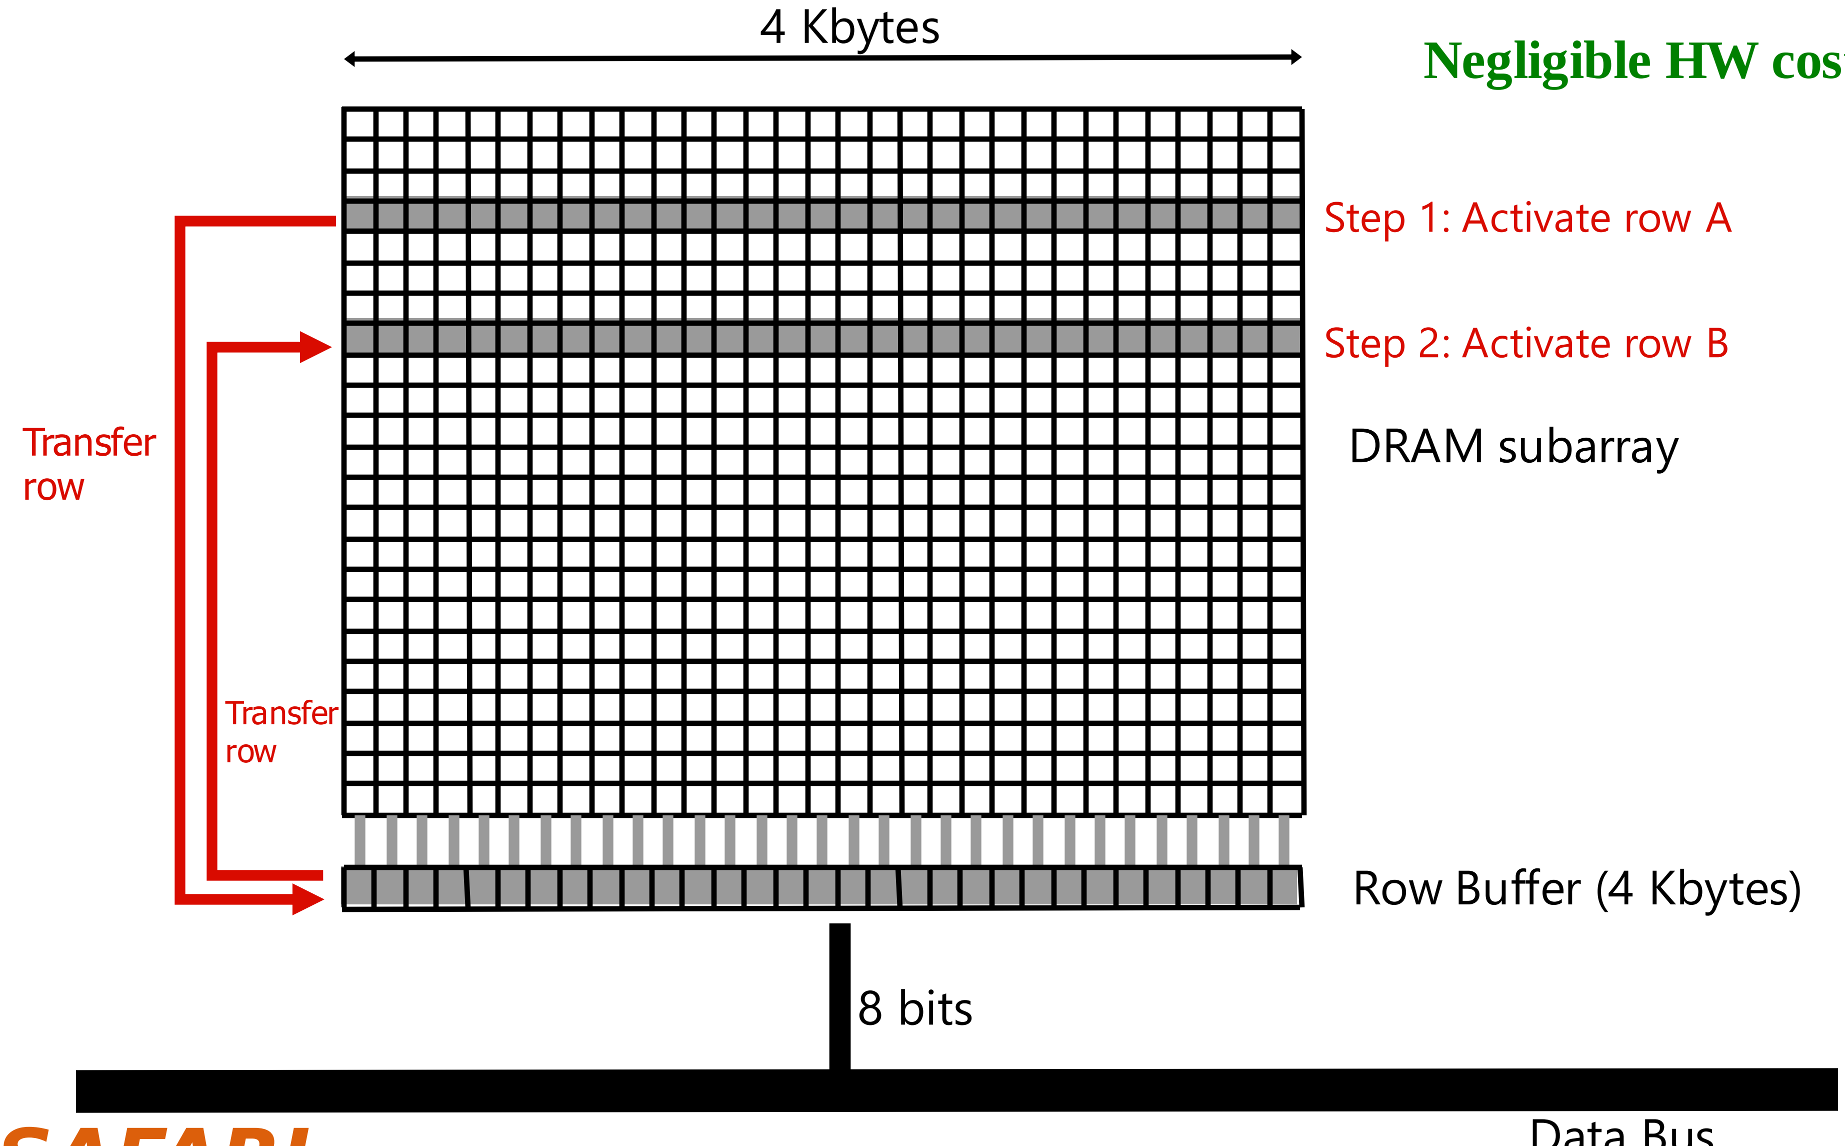

Promising Solutions: Processing Using Memory (PuM)
Implementing Data Movement

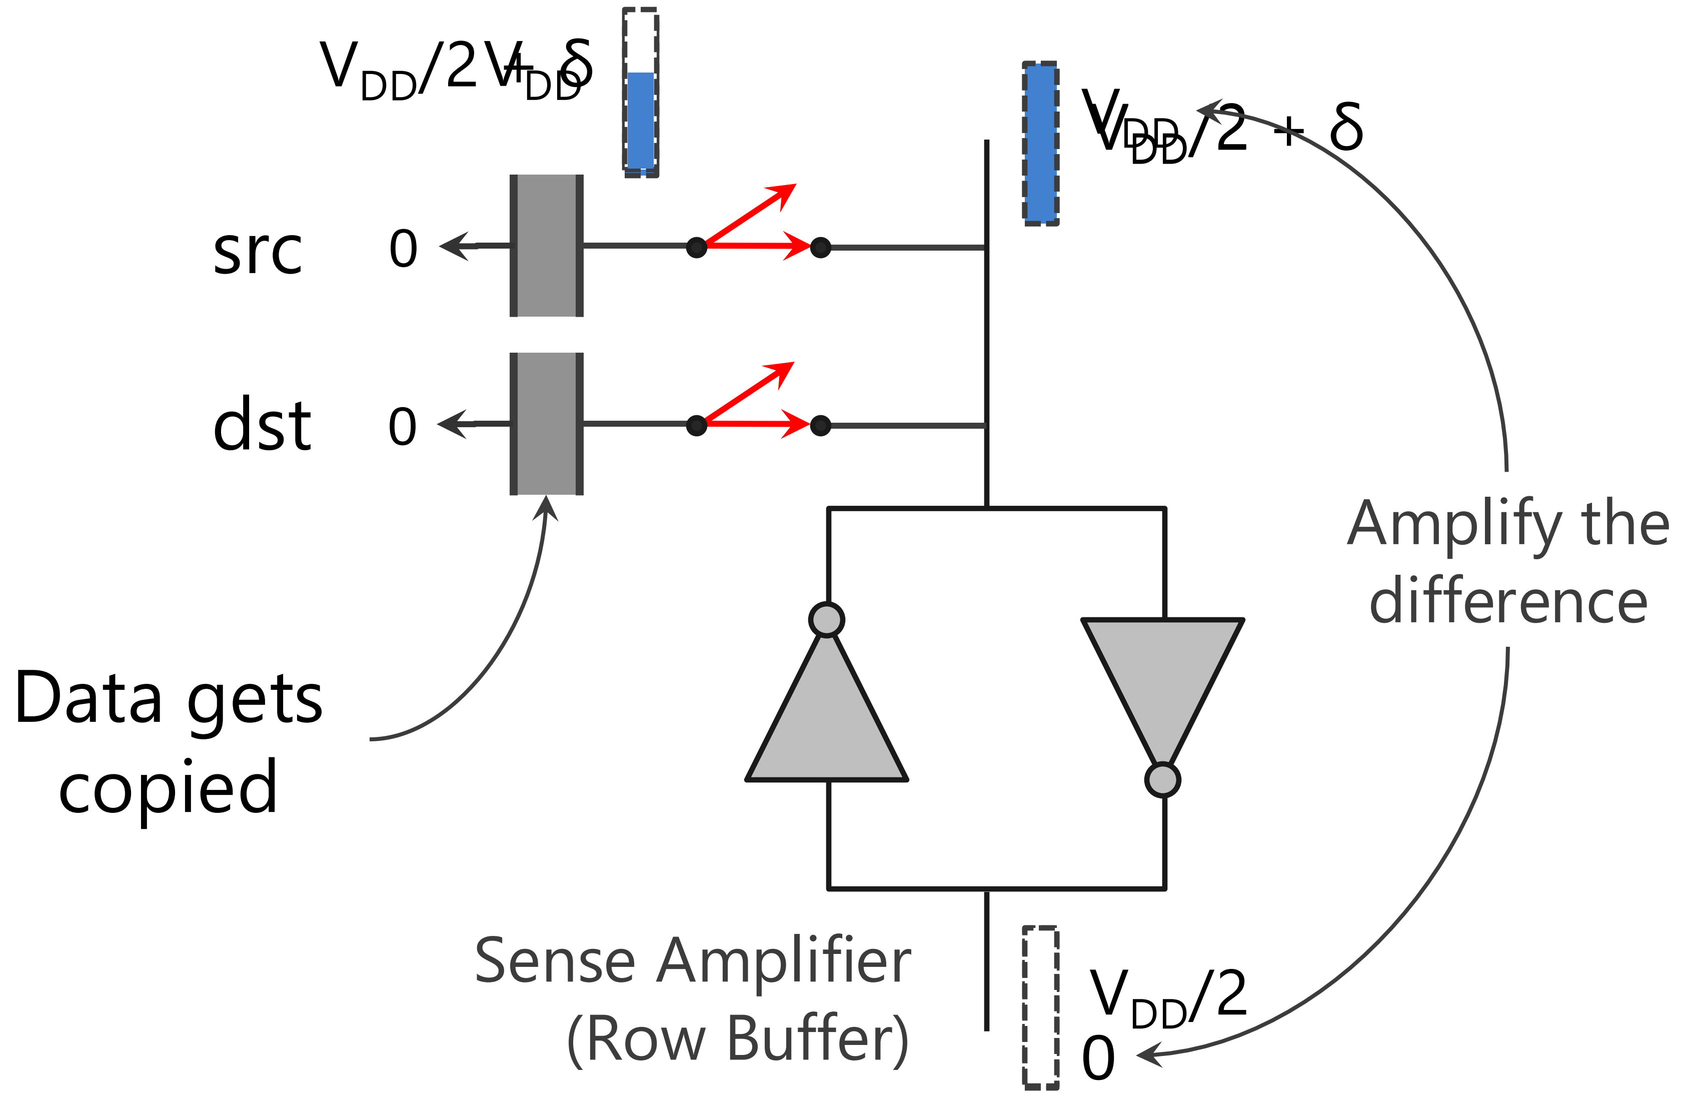

Promising Solutions: Processing Using Memory (PuM) Implementing Data Movement

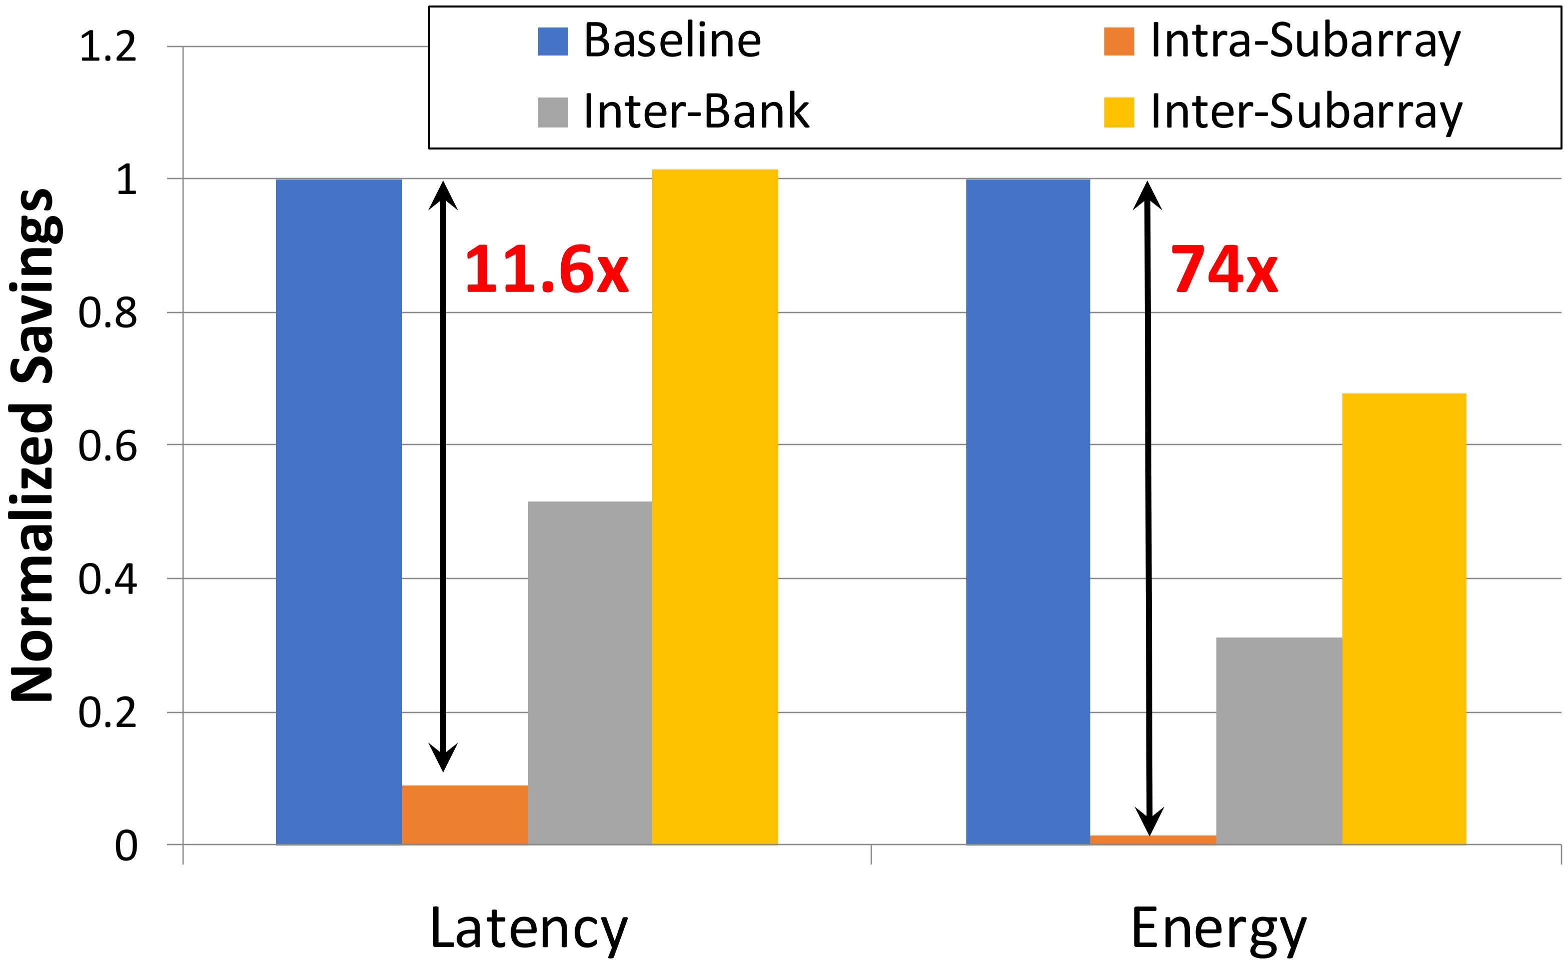

Seshadri et al., "RowClone: Fast and Efficient In-DRAM Copy and Initialization of Bulk Data," MICRO 2013.

Promising Solutions: Processing Using Memory (PuM)
Implementing Logic Operations

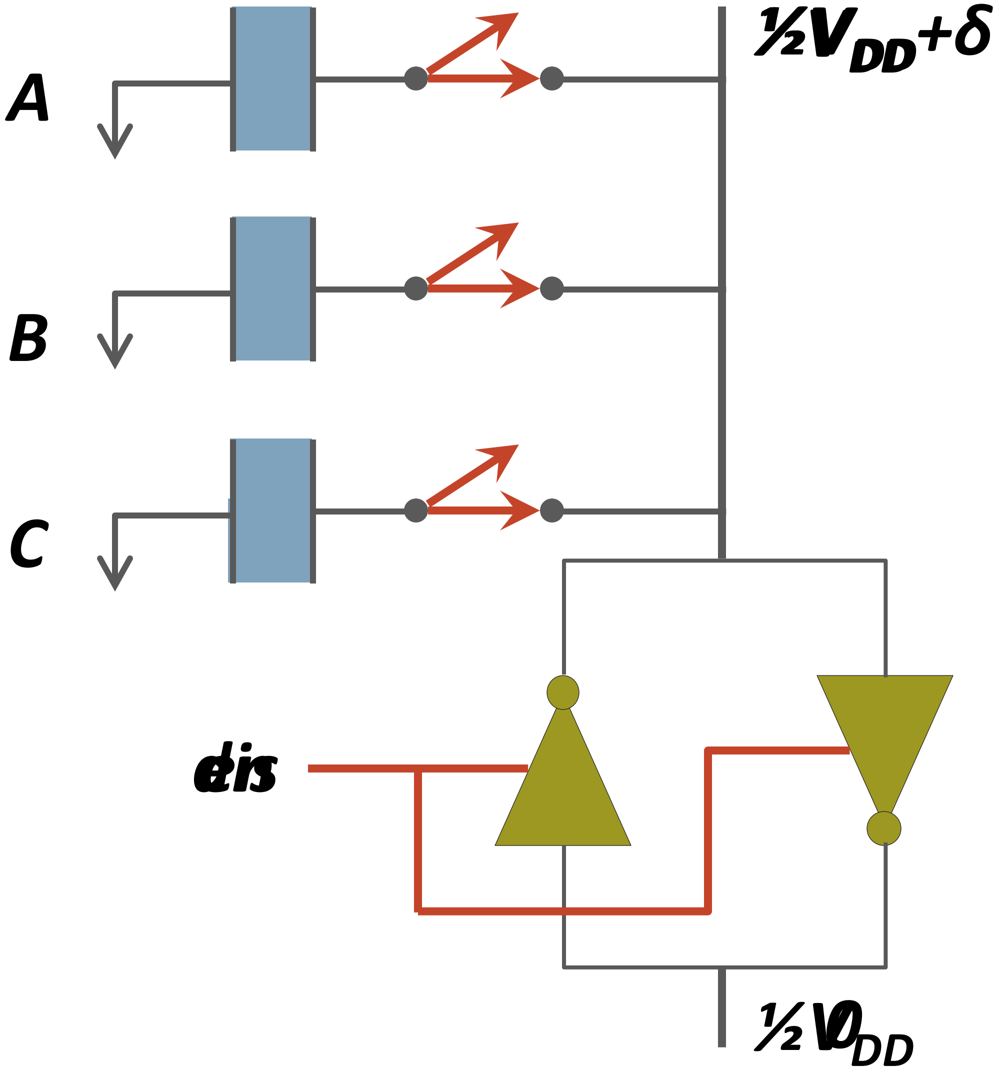

Final State
AB + BC + AC

C(A + B) +
~C(AB)

Seshadri+, "Fast Bulk Bitwise AND and OR in DRAM", IEEE CAL 2015.

Promising Solutions: Processing Using Memory (PuM) Implementing Logic Operations

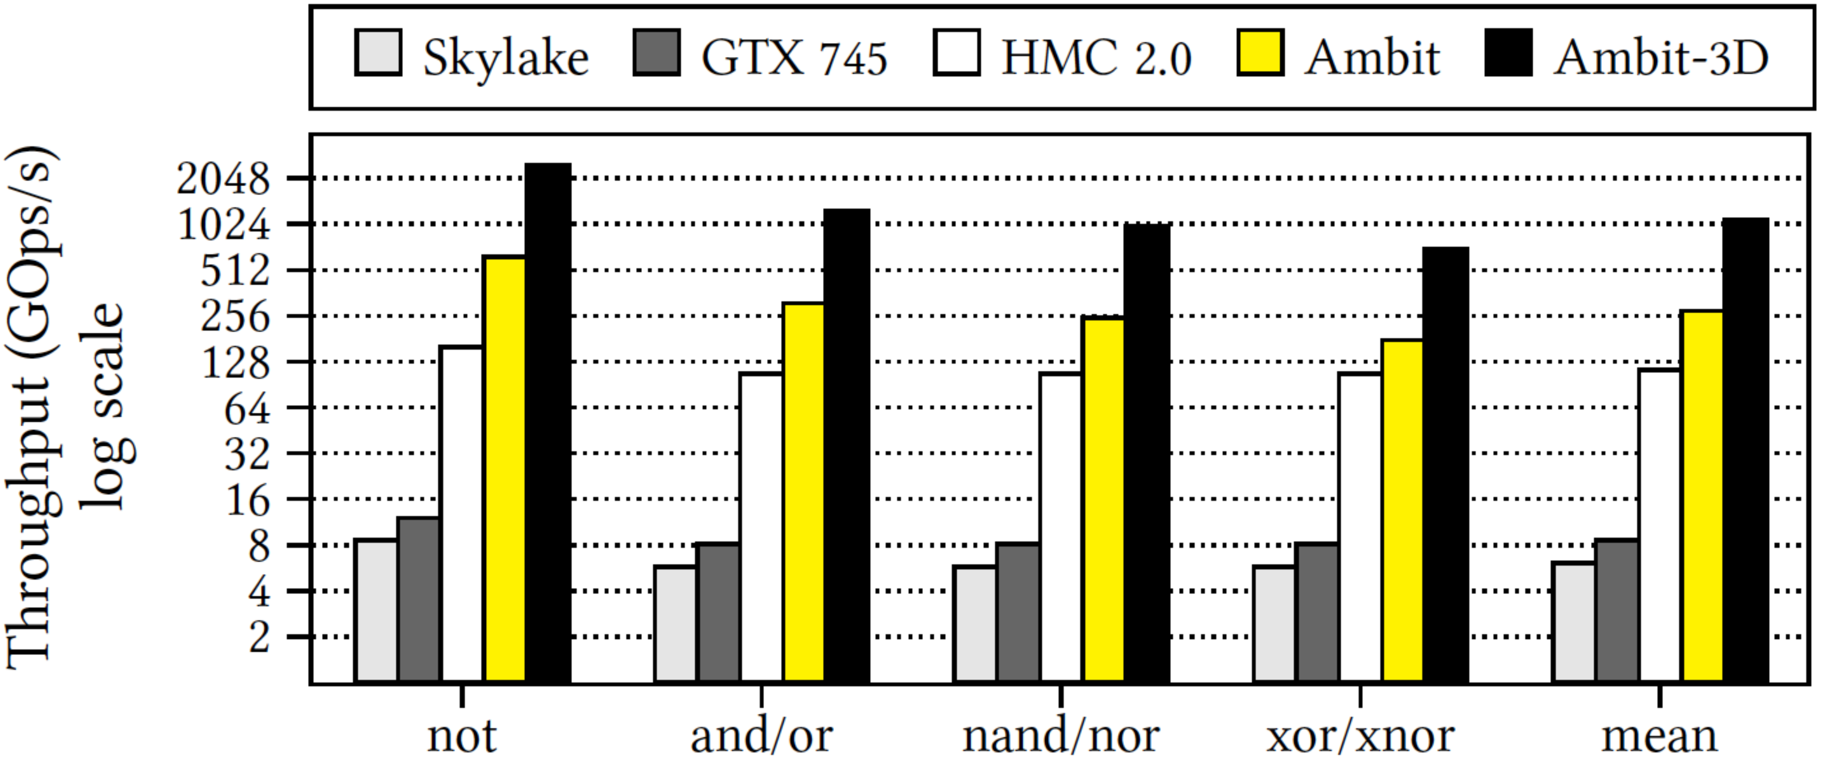

Figure 9: Throughput of bitwise operations on various systems.

Promising Solutions: Processing Using Memory (PuM) Implementing Logic Operations

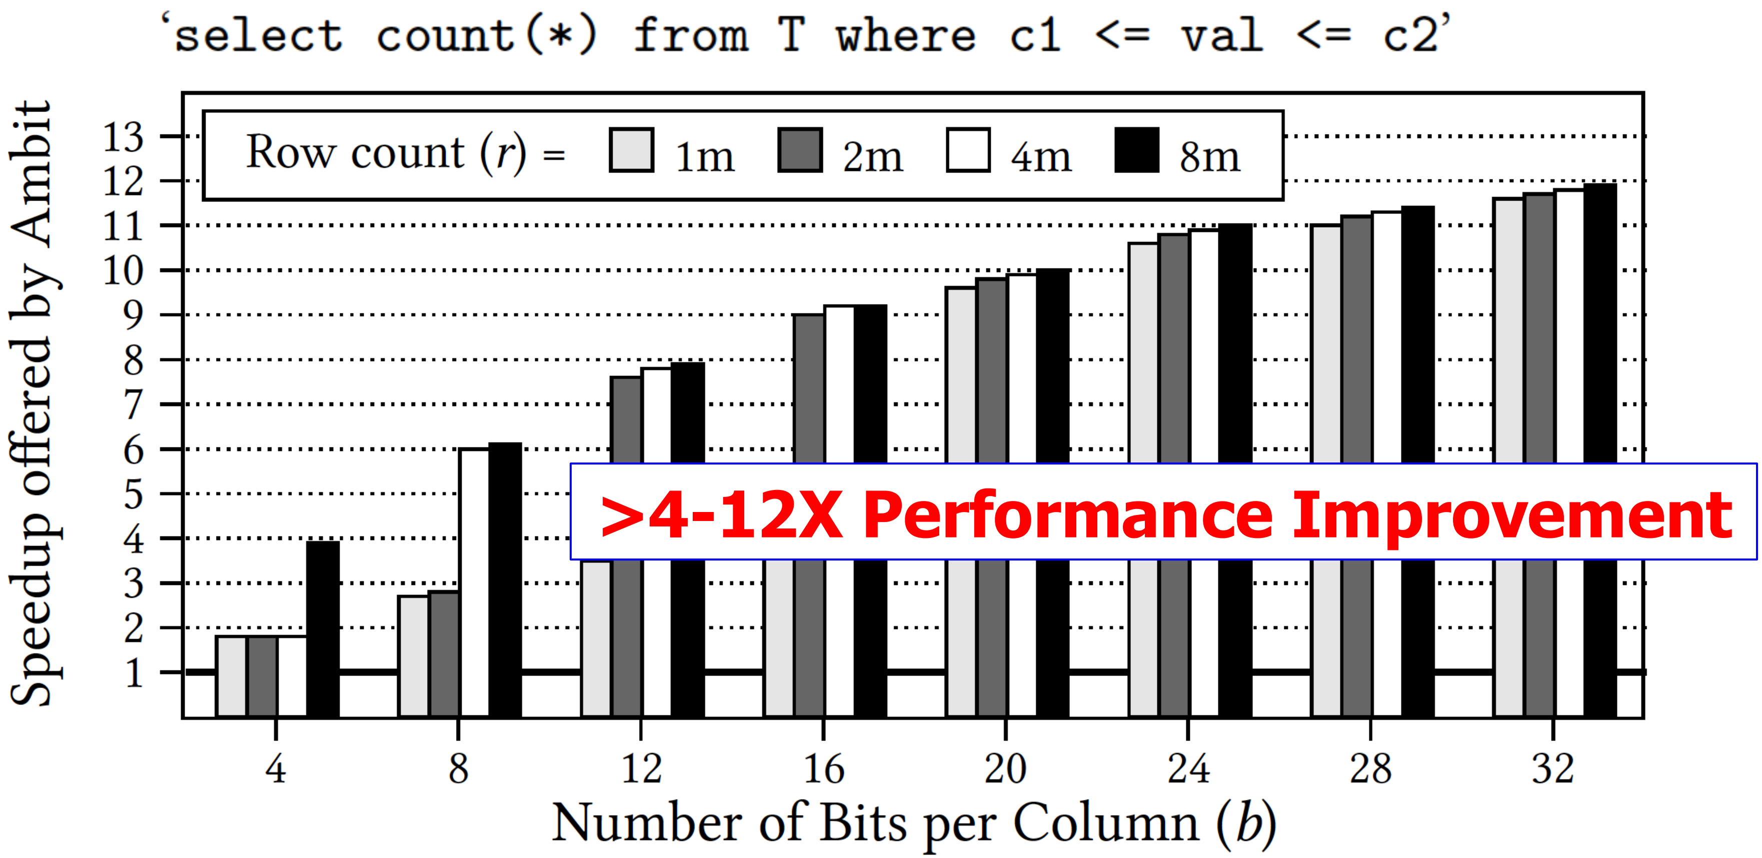

Figure 11: Speedup offered by Ambit over baseline CPU with SIMD for BitWeaving

Seshadri+, "Ambit: In-Memory Accelerator for Bulk Bitwise Operations using Commodity DRAM Technology," MICRO 2017.

Motivation

Processing near/using memory
is a promising computation paradigm

Distributing execution across
CPU, GPU, PnM, and PuM
should NOT hurt security and privacy

Research Scope (Going Forward)

- **Scope**: Confidentiality, integrity, and availability in large-scale & vastly-shared processing in memory and storage
- **Confidentiality**:
  - Physical isolation
  - Architecture support for encryption in memory/storage
- **Integrity**:
  - In-field patchable and dynamically tunable maintenance mechanisms
  - Integrity check and efficient roll-back mechanisms
- **Availability**:
  - Intelligent scheduling and fairness mechanisms
  - Adaptive maintenance supported with online profiling

Enabling Efficient and Scalable DRAM Read Disturbance Mitigation via New Experimental Insights into Modern DRAM Chips

Abdullah Giray Yağlıkçı

Research Summary and Future Directions

July 9, 2024

SAFARI

ETH zürich

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary
- Interfaces should be **more flexible**

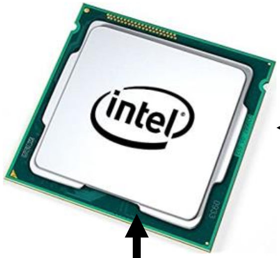

Memory controller decides what should be done when

DRAM chip has read disturbance solution inside (tracking+prevention)

- The memory controller should provide the DRAM chip with **necessary time window** to perform **preventive actions (e.g., refreshing rows)**
- The memory controller **does not have** the tracking information
- Communicating is **not straightforward** due to strict communication protocol

A more flexible interface is necessary

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary
- Interfaces should be **more flexible**

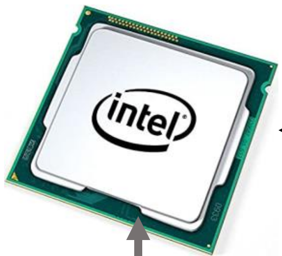

A new DDR5 extension
(JESD79-5C, 2024)

ACT
ACT
ACT

alert_n
alert_n
alert_n

Memory controller decides what should be done when

DRAM chip has read disturbance solution inside (tracking+prevention)

- **Two Key Issues:**
- A naïve implementation can lead to **very high command bus occupancy**
  - Very high bandwidth consumption
  - Poor parallelism

- DRAM **self-manages preventive refreshes**. Why not **other maintenance routines**?

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary
- Interfaces should be **more flexible**
- Memory controllers should be **more intelligent** in detecting **malicious activity**

- DRAM chips become **more and more vulnerable** to RowHammer **and RowPress**
- **Key Insight:** Activation counts of **benign applications** get close to **unsafe levels**
- **Problem:** DRAM read disturbance solutions are getting **prohibitively expensive**
- **Research Question:** How to identify malicious threads/processes/users?
- More **intelligent detection** mechanisms are needed
- The **memory controller** observes all memory accesses

More intelligent memory controllers can help

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**
  Preliminary data on aging via 68-day of continuous hammering

Aging can lead to read disturbance bitflips at **smaller** hammer counts

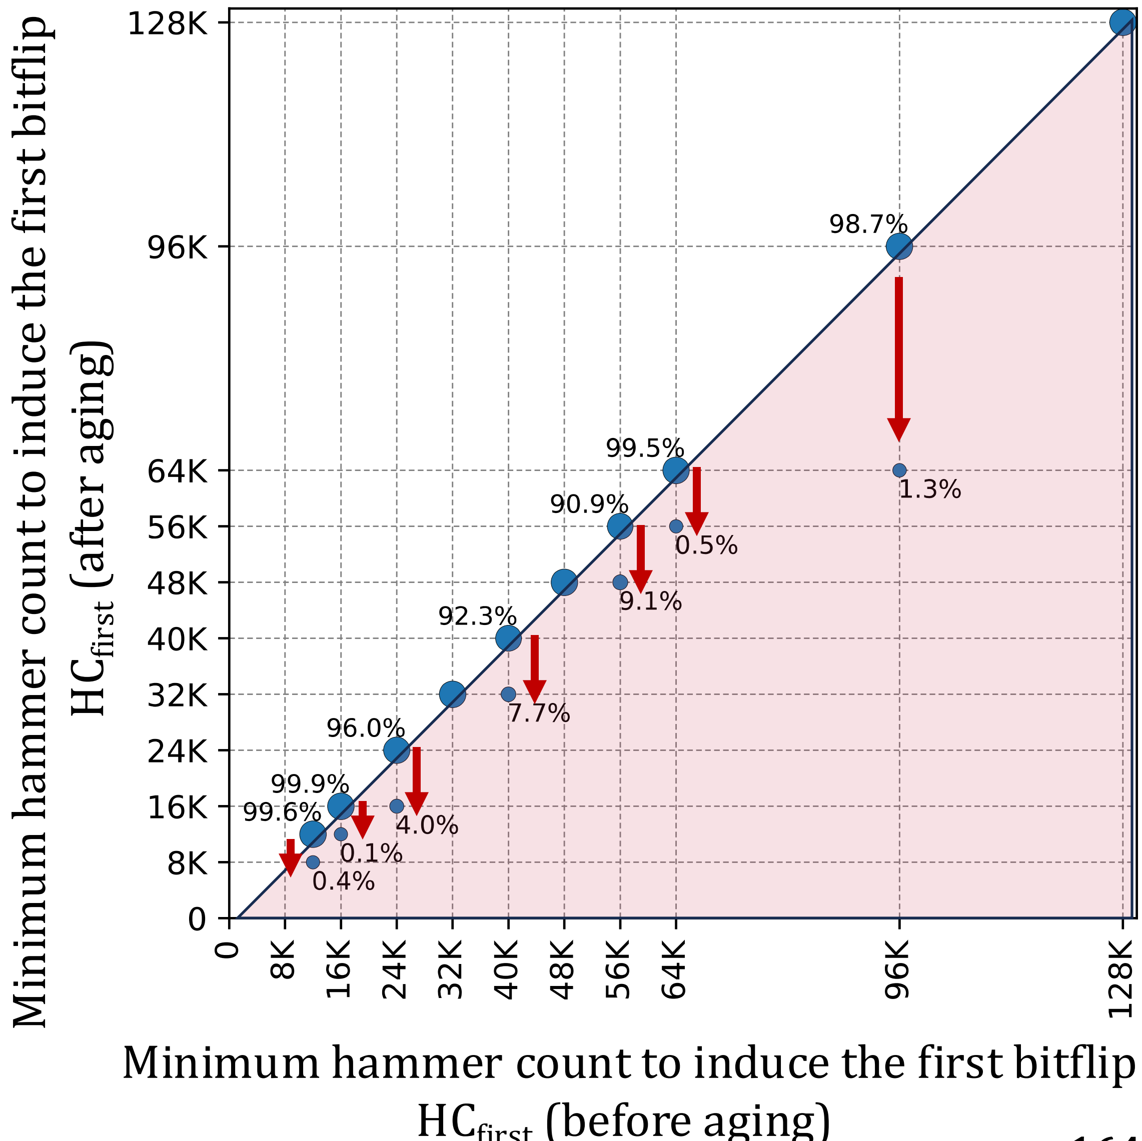

SAFARI

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**
  Preliminary data on aging via 68-day of continuous hammering

Aging can lead to read disturbance bitflips at **smaller** hammer counts

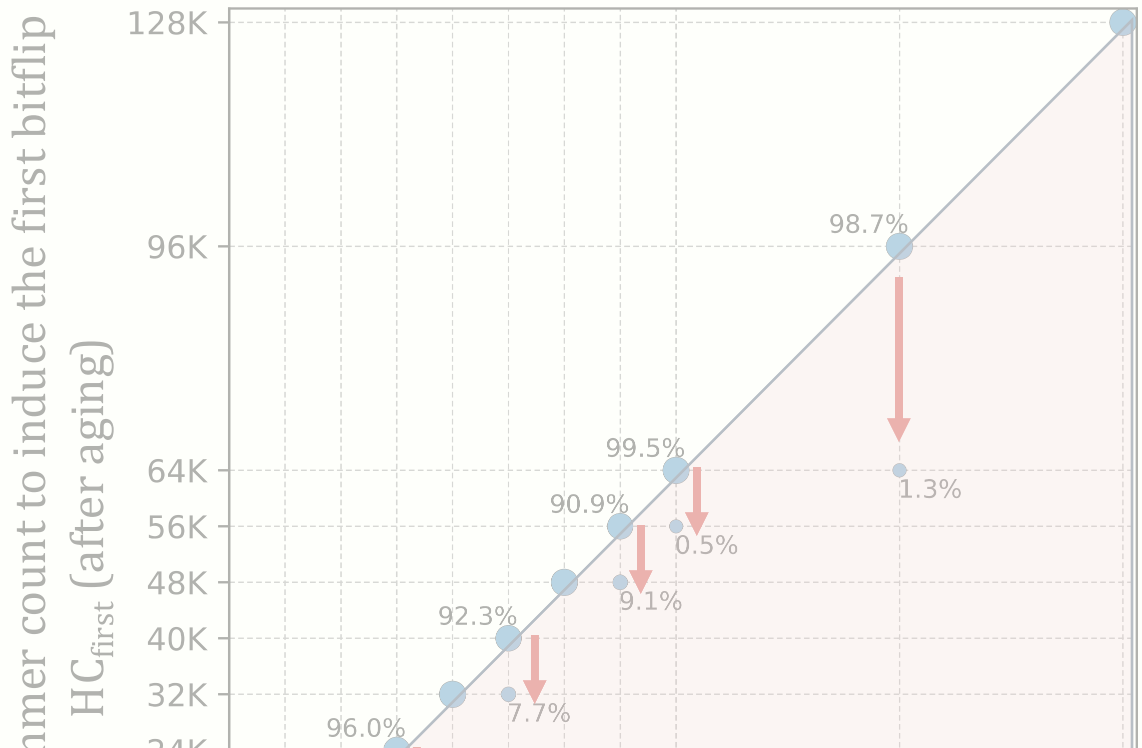

Future work:
rigorous aging characterization
and **online profiling of read disturbance vulnerability**

Minimum hammer count to induce the first bitflip
HC_{first} (before aging)

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**

- **Interactions** across different error mechanisms
  - RowHammer
  - RowPress
  - Data retention time errors
  - Variable retention time

  ...

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**

- **Interactions** across different error mechanisms

- What is **the worst-case**?
  - Temperature
  - Data pattern
  - Memory access pattern
  - Spatial variation
  - Voltage

What is **the worst-case** considering all **these sensitivities**?

What is **the minimum hammer count** to induce a read disturbance bitflip?

Deeper Understanding of Physics and Vulnerabilities

- The effect of **aging**

- **Interactions** across different error mechanisms

- What is **the worst-case**?

How reliable are our DRAM chips?

How reliable will our DRAM chips be tomorrow?

We **do not** know! This is an **open research problem**

Future Research for Better Memory Systems

Deeper Understanding of
Physics and Vulnerabilities

Flexible and Intelligent Memory
Chips, Interfaces, Controllers

Cross-Layer
Communication

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary

DRAM read disturbance solution is here

Configuration

DRAM chip is here

DRAM module can be replaced

Reconfiguration is needed

+ add more hardware counters
+ increase metadata cache capacity

Aging might change the chip's characteristics

Deployed solutions should be patchable in field

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary
- Interfaces should be **more flexible**

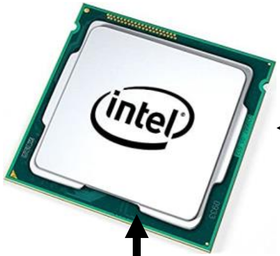

Memory controller decides what should be done when

DRAM chip has read disturbance solution inside (tracking+prevention)

- The memory controller should provide the DRAM chip with **necessary time window** to perform **preventive actions (e.g., refreshing rows)**
- The memory controller **does not have** the tracking information
- Communicating is **not straightforward** due to strict communication protocol

A more flexible interface is necessary

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary
- Interfaces should be **more flexible**

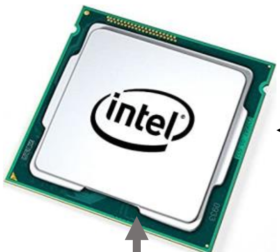

A new DDR5 extension
(JESD79-5C, 2024)

ACT
ACT
ACT

alert_n
alert_n
alert_n

Memory controller decides what should be done when

DRAM chip has read disturbance solution inside (tracking+prevention)

- **Two Key Issues**:
- A naïve implementation can lead to **very high command bus occupancy**
  - Very high bandwidth consumption
  - Poor parallelism

- DRAM **self-manages preventive refreshes**. Why not **other maintenance routines**?

Self-Managing DRAM

The three major DRAM maintenance operations:

- ❖ Refresh
- ❖ RowHammer Protection
- ❖ Memory Scrubbing

Implementing new **maintenance mechanisms** often requires difficult-to-realize changes

Self-Managing DRAM (SMD)

Enables implementing

- new **in-DRAM** maintenance mechanisms
- with **no further changes** in the *DRAM interface* and the *memory controller*

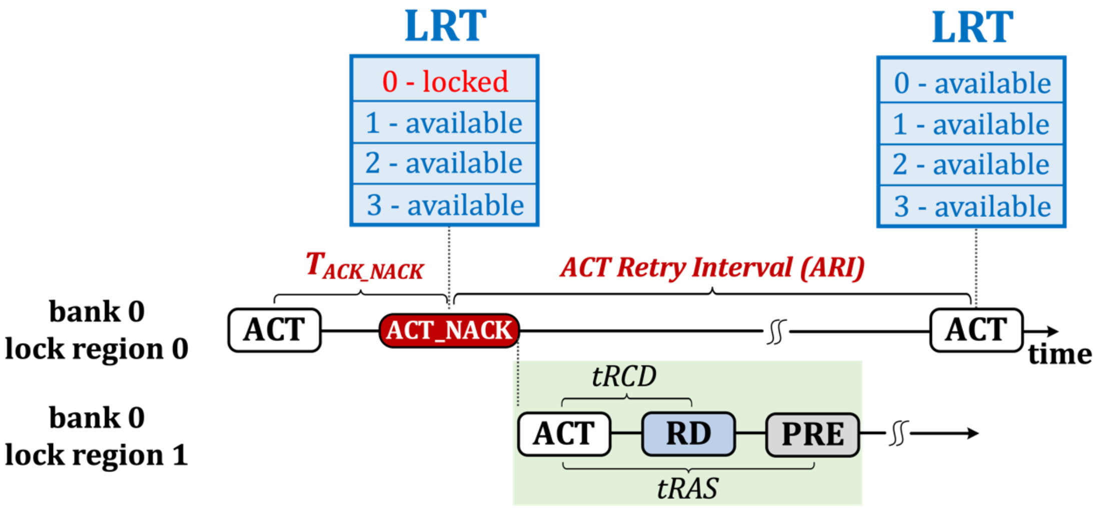

SMD-based *maintenance operations* provide significant **performance** and **energy** benefits across many system configurations and workloads

Hasan Hassan, Ataberk Olgun, A. Giray Yaglikci, Haocong Luo, and Onur Mutlu, ``Self-Managing DRAM: A Low-Cost Framework for Enabling Autonomous and Efficient in-DRAM Operations,'' cs.AR:2207.13358, 2022 (last update in 2024)

Flexible and Intelligent Chips, Interfaces, Controllers

- **In-field patching** is necessary
- Interfaces should be **more flexible**
- Memory controllers should be **more intelligent** in detecting **malicious activity**

- DRAM chips become **more and more vulnerable** to RowHammer **and RowPress**
- **Key Insight:** Activation counts of **benign applications** get close to **unsafe levels**
- **Problem:** DRAM read disturbance solutions are getting **prohibitively expensive**
- **Research Question:** How to identify malicious threads/processes/users?
- More **intelligent detection** mechanisms are needed
- The **memory controller** observes all memory accesses

More intelligent memory controllers can help

BreakHammer

- **Key Observation**: Mitigating DRAM read disturbance causes delays in memory accesses
- **Our Exploit**: Denial of memory service is possible via triggering mitigation mechanisms
- **Key Idea**: Throttling memory accesses of threads that trigger mitigation mechanisms repeatedly
- **BreakHammer**:
  - Detects the threads that repeatedly trigger the mitigation mechanisms
  - Limits their on-the-fly memory request counts and MSHRs
  - Near-zero area overhead and no additional memory access latency
- **Evaluation**:
  - Improves **system performance** by **48.7%** on average (**105.5%** max)
  - Reduces the **maximum slowdown** by **14.6%** on average

Oğuzhan Canpolat, A. Giray Yağlıkçı, Ataberk Olgun, İsmail Emir Yüksel, Yahya Can Tuğrul, Konstantinos Kanellopoulos, Oğuz Ergin, Onur Mutlu, "Leveraging Adversarial Detection to Enable Scalable and Low Overhead RowHammer Mitigations," [cs.CR] 2404.13477, 2024.

Future Research for Better Memory Systems

? Deeper Understanding of Physics and Vulnerabilities

? Flexible and Intelligent Memory Chips, Interfaces, Controllers

? Cross-Layer Communication

Cross-Layer Communication

		Detection	Mitigation
Software 	• Memory allocations • Memory access patterns • Control flow patterns • Time / power measurements	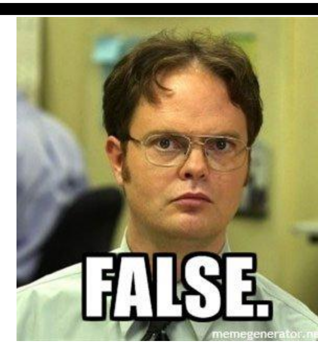	++
uArch 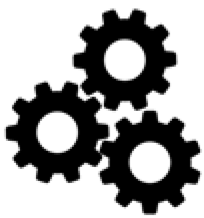	• Memory request scheduling • Speculative execution • Prefetching, branch prediction • Power management	+	+
Device 	• Bitflips occur • Memory isolation is broken	++	~

How Large is 1000 Activations?

- Bitflips occur **at ~1000 activations**
- Mitigation mechanisms trigger **preventive actions** (e.g., preventive refresh) **at ~500 activations**
- Is 500 **a distinctive activation count**?
- Benign workloads activate **hundreds of rows** more than **500 times** in a refresh window

Memory intensive workloads
from SPEC’06, SPEC’17, TPC, YCSB, and MediaBench

Workload	MPKI	ACT-64+	ACT-128+	ACT-512+
429.mcf	68.27	2564	2564	2564
470.lbm	28.09	7089	6596	664
462.libquantum	25.95	1	0	0
549.fotonik3d	25.28	10065	88	0
459.GemsFDTD	24.93	10572	218	0
519.lbm	24.37	5824	5455	2482
434.zeusmp	22.24	11085	4825	292
510.parest	17.79	803	185	94
433.milc	17.22	321	92	0
437.leslie3d	15.82	4678	631	7
483.xalancbmk	13.67	4354	776	113
482.sphinx3	12.59	1385	762	304
505.mcf	11.35	1582	1384	732
471.omnetpp	10.72	1015	419	122
tpch2	9.09	875	307	88
520.omnetpp	9.00	1185	84	32
tpch17	7.43	1196	158	26
473.astar	5.18	5957	22	0
436.cactusADM	4.94	6151	2354	1134
jp2_encode	4.18	0	0	0

Benign workloads **might not be so benign**

Cross-Layer Communication

		Detection	Mitigation
Software	• Memory allocations • Memory access patterns • Control flow patterns • Time / power measurements	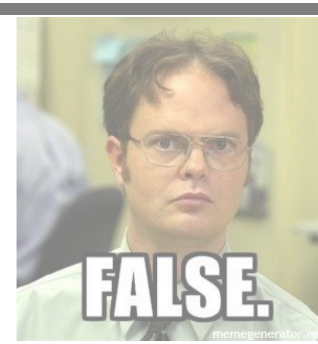	++
Device	• Bitflips occur • Memory isolation is broken	++	~

Cross-layer communication is crucial going forward

Enabling Efficient and Scalable DRAM Read Disturbance Mitigation

via New Experimental Insights into Modern DRAM Chips

Backup Slides

Research Summary and Future Directions

July 9, 2024

SAFARI

ETH zürich

Awards and Honorable Mentions

<u>BlockHammer</u>: One of four finalists
Intel Hardware Security Academic Awards in 2022

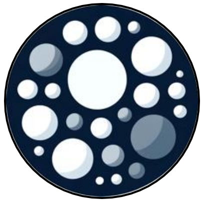

<u>Svärd</u>: First place
ACM SRC (PACT) in 2023 (grand final is ongoing)

<u>Thesis</u>: One of five finalists (ongoing)
HOST Ph.D. Dissertation Competition in 2024

IChannels

Exploiting Current Management Mechanisms to Create Covert Channels in Modern Processors

Jawad Haj-Yahya

Jeremie S. Kim *A. Giray Yağlıkçı* *Ivan Puddu* *Lois Orosa*

Juan Gómez Luna *Mohammed Alser* *Onur Mutlu*

SAFARI ETH zürich

Executive Summary

Problem: Current management mechanisms throttle instruction execution and adjust voltage/frequency to accommodate power-hungry instructions (PHIs).
These mechanisms may compromise a system's confidentiality guarantees

Goal:
1. Understand the throttling side-effects of current management mechanisms
2. Build high-capacity covert channels between otherwise isolated execution contexts
3. Practically and effectively mitigate each covert channel

Characterization: Variable execution times and frequency changes due to running PHIs
We observe five different levels of throttling in real Intel systems

IChannels: New covert channels that exploit side-effects of current management mechanisms
- On the same hardware thread
- Across co-located Simultaneous Multi-Threading (SMT) threads
- Across different physical cores

Evaluation: On three generations of Intel processors, IChannels provides a channel capacity
- 2× that of PHIs' variable latency-based covert channels
- 24× that of power management-based covert channels

ABACuS: All-Bank Activation Counters

- Ataberk Olgun, Yahya Can Tugrul, Nisa Bostanci, Ismail Emir Yuksel, Haocong Luo, Steve Rhyner, Abdullah Giray Yaglikci, Geraldo F. Oliveira, and Onur Mutlu, **"ABACuS: All-Bank Activation Counters for Scalable and Low Overhead RowHammer Mitigation"**
  *To appear in Proceedings of the 33rd USENIX Security Symposium (**USENIX Security**)*, Philadelphia, PA, USA, August 2024.
  [arXiv version]
  [ABACuS Source Code]

ABACuS: All-Bank Activation Counters for Scalable and Low Overhead RowHammer Mitigation

Ataberk Olgun Yahya Can Tugrul Nisa Bostanci Ismail Emir Yuksel
Haocong Luo Steve Rhyner Abdullah Giray Yaglikci Geraldo F. Oliveira Onur Mutlu

ETH Zurich

ABACuS: All-Bank Activation Counters

Goal: Prevent RowHammer bitflips at low performance, energy, and area cost

Key Observation: Workloads tend to access the same row in all DRAM banks at around the same time

Key Idea: Use one hardware counter to keep track of activation counts of the same row across all banks

- Make high-performance, area-hungry counter-based mechanisms practical

Key Results:

Faster than the lowest-area-cost counter-based defense mechanism
Smaller than the lowest-performance-overhead counter-based defense mechanism

0.59% avg. performance overhead (single-core) at a RowHammer threshold (1K)

- Only 9.79 KiB on-chip storage per DRAM rank (0.02% of a Xeon processor)

1.52% avg. performance overhead (single-core) at an ultra-low threshold (125)

- 75.70 KiB on-chip storage per DRAM rank (0.11% of the Xeon processor)

https://github.com/CMU-SAFARI/ABACuS

CoMeT: Count-Min-Sketch-based Row Tracking

- Nisa Bostanci, Ismail Emir Yuksel, Ataberk Olgun, Konstantinos Kanellopoulos, Yahya Can Tugrul, A. Giray Yaglikci, Mohammad Sadrosadati, Onur Mutlu
  "CoMeT: Count-Min-Sketch-based Row Tracking to Mitigate RowHammer at Low Cost,"
  *in Proceedings the 30th International Symposium on High-Performance Computer Architecture (**HPCA**), Edinburgh, March 2024.*
  [arXiv version]
  [CoMeT Source Code]

CoMeT: Count-Min-Sketch-based Row Tracking to Mitigate RowHammer at Low Cost

F. Nisa Bostancı İsmail Emir Yüksel Ataberk Olgun Konstantinos Kanellopoulos
Yahya Can Tuğrul A. Giray Yağlıkçı Mohammad Sadrosadati Onur Mutlu

ETH Zürich

Executive Summary

Goal: Prevent RowHammer bitflips with low area, performance, and energy overheads in highly RowHammer-vulnerable DRAM-based systems

Key Idea: Use low-cost and scalable hash-based counters to accurately track DRAM rows

CoMeT:
- tracks most DRAM rows with scalable hash-based counters by employing the Count-Min-Sketch technique to achieve a low area cost
- tracks only a small set of DRAM rows that are activated many times with highly accurate per-DRAM-row activation counters to reduce performance penalties

Evaluation: CoMeT achieves a good trade-off between area, performance and energy costs
- incurs significantly less area overhead (74.2×) compared to the state-of-the-art technique
- outperforms the state-of-the-art technique (by up to 39.1%)

https://github.com/CMU-SAFARI/CoMeT